

FAUST
January 19, 2009

Ole Weidner (oweidner@cct.lsu.edu)

FAUST

A Framework for Adaptive Ubiquitous Scalable Tasks

Abstract

Contents

| | | |
|----------|-------------------------------------|-----------|
| 1 | Random Thoughts | 3 |
| 1.1 | Modeling Job Dependencies | 3 |
| 1.1.1 | Types of Dependencies | 3 |
| 1.1.2 | Describing Dependencies | 4 |
| 2 | Implementation | 7 |
| 2.1 | Overview | 7 |
| 2.2 | Interface | 7 |
| 2.3 | Scheduling Engine | 7 |
| 2.4 | System Agents | 7 |
| 2.4.1 | Mode of Operation | 8 |
| 3 | Appendix | 9 |
| 4 | References | 10 |

1 Random Thoughts

1.1 Modeling Job Dependencies

The overall goal of the FAUST framework is to schedule a given set of jobs on a number of distributed resources as effective as possible. Effectiveness in our case means minimum *makespan*¹ scheduling or time to completion. In case of an application which consists of a set of independent jobs (embarrassingly parallel EP), scheduling is rather trivial: execute as many jobs as possible at the same time on all available resources. An example for such an application would be a parameter sweep which generates and executes a set of independent model instances with different input parameters.

However, lots of distributed applications do not fall into the category of EP applications. Jobs often require communication with other jobs or they may rely on data that has to be generated by other jobs. Message-passing (e.g. MPI) as well as distributed workflows are good examples for these types of applications. Unfortunately, scheduling becomes way more complex in this case, since it has to take not only the availability of resources but also things like interconnect bandwidth, shared filesystems, etc. into account to minimize the overhead exposed by job dependencies.

In this section, we try to identify different types of job dependencies, describe how to model them on application level and discuss the implications for possible minimum makespan scheduling algorithms.

1.1.1 Types of Dependencies

So far, we identified two types of dependencies in distributed applications that are relevant for job scheduling and placement. We distinguish between dependencies that rely on data availability and dependencies that rely on communication:

Data Dependencies occur whenever a job requires data that is generated by another job or set of jobs. Imagine for example an image processing application (Figure 1) that splits up an image into several regions and applies a filter to each of the regions in parallel. Another job takes the processed fragments and puts them back together. This job depends on the output generated by the filter instances.

Communication Dependencies occur whenever two or more jobs need to exchange information while they are running. Imagine a 2D heat-transfer appli-

¹The makespan of a schedule is its total execution time.




Figure 1: Example of a dependency graph (r) for an image processing application where job M depends on the data generated by jobs $A...N$. The grey vertical bar in the scheduling scheme (l) represents the time overhead generated by data transfer.

cation (Figure 2) that splits up the problem space in 4 regions and maps them to 4 different jobs (domain decomposition). Communication has to occur whenever the heat transfers across domain boundaries. This concept is also known as ghost-zone exchange and a very well known concept in MPI.




Figure 2: Example of a dependency graph (r) for set of communicating jobs in a domain decomposition application. The grey vertical bars within the jobs in the scheduling scheme (l) represent the time overhead generated by communication.

1.1.2 Describing Dependencies

Describing just the dependencies and type of dependencies (data or communication) between jobs enables a scheduler to execute the jobs in a way that satisfies the dependencies. However, without additional information about the jobs and the logical dependencies, a scheduling algorithm won't be able to *place* the jobs efficiently on the distributed resources. These attributes usually can't be extracted from the application automatically. They have to be described explicitly on application level. We identify a minimum set of these attributes and show how they can be described using the FAUST API.

Data Dependencies expose a potential data transfer overhead. A scheduling

algorithm has to decide whether it should either move the data to the computation or the computation to the data (place the dependent job as close² to the data as possible). To be able to make this decision, the following information has to be provided on application level:

- **Expected runtime** of the jobs that are part of the dependency.

```
faust::attribute::walltime
```

- **Expected amount of data** generated by a job.

```
faust::attribute::data_volume
```

The FAUST framework provides an interface to describe data dependencies in applications through the *job submission* interface. In case of the example image processing application described above, this could look like the following (simplified) code fragment:

Describing Data Dependencies

```
1  job::description filter_jd;
2  filter_jd.set_attribute(walltime, "10.0");
3  filter_jd.set_attribute(data_volume, "0.5GB");
4
5
6  std::vector<std::string> filter_desc;
7  for(int i=1; i<10; ++i)
8      filter_desc.push_back(filter_jd); // create 10 filter instances
9
10 job::service s;
11 job::group filters = s.create_job_group(filter_desc);
12
13 // create the composer job which has a DATA dependency with the
14 // filter job group.
15 job::description composer_jd;
16 job::job composer = s.create_job(composer_jd, filters, type::DATA);
17
18 s.schedule();
19
```

Communication Dependencies expose a potential communication overhead. A scheduling algorithm has to decide whether it should place the jobs on resource which are connected via high-bandwidth interconnects or if processing power for the individual jobs is more important than communication. The following application attributes can be specified on application level to help the scheduler to make the right decisions:

² *Close* in this context is defined as the interconnect bandwidth between two locations.

- **Computation/communication ratio**

`faust::attribute::comm_comp_ratio`

- **Communication pattern** (regular, irregular, ...)

`faust::attribute::comm_pattern`

FIXME: I'm still struggling if these are the right attributes and if it's feasible to describe them on application level. Reading some papers about this might help...

2 Implementation

2.1 Overview

2.2 Interface

2.3 Scheduling Engine

2.4 System Agents

The System Agents are used to provide the Scheduling Engine as well as the programmer (through the API) with informations about all participating systems that are relevant for scheduling descissions and job execution. *FAUST* System Agents are realized as independent command line applications that run on the execution systems and report vital system and status informations like queue status, network load, etc. to back to the *FAUST* instance.




Figure 3: Example of a *FAUST* application instance running agents on three execution hosts. To ensure application persitency, the communication between the application and the agents flows through a proxy database.

Besides reporting system informations to the application, agents can optionally act as job submission endpoints. In this scenario, an application sends jobs directly to the agents for execution and not to the local job service (like Globus, GridSAM, PBS, etc). This will usually happen, if the application scheduler (or the programmer) decides to apply 'hijacking' strategies like *Glide-In* which require circumvention of local queueing systems.

2.4.1 Mode of Operation

The System Agents are transparently deployed and started by a *FAUST* Framework application as soon as a new `faust::service` instance gets instantiated. Each `service` instance is defined by a set of resource `descriptions` which represent targets for possible job submission. Each `resource` has its own System Agent for management, monitoring and information retrieval.

Unfortunately, different distributed infrastructures require different techniques to extract the informations that are required for effective scheduling. Although there are several more or less well defined interfaces to query these informations like *NWS* or *BQP*, it appears that in practice their existence and proper operation cannot be assumed. For this reason, agents require a relatively rich set of initial system informations that are used to generate subsequent runtime informations and performance predictions. Once deployed and started, System Agents loop through the following execution procedure:

3 Appendix

4 References