



CENTER FOR COMPUTATION
& TECHNOLOGY


Grid Interoperability at the Application Level using SAGA

Shantenu Jha^{12*}, H Kaiser¹, A Merzky¹, O Weidner¹

¹Center for Computation and Technology, LSU

² Department of Computer Science, LSU

*Also with NeSC, Edinburgh and UC-London





Outline

- Interoperability: An Application's Perspective
 - Application vs Service
 - Grid Aware vs Unaware
- Simple API for Grid Applications(SAGA): Introduction
- SAGA: Interoperability at different levels
- SAGA in relation to GIN
- Application Example(s)



Motivation

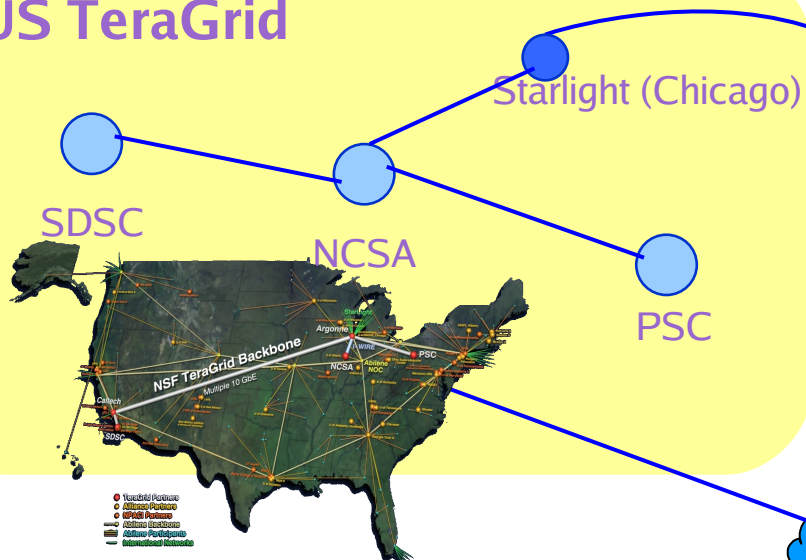
- GIN:
 - Interoperation at the service level
 - Bottom-up approach
- Specific Application Projects:
 - “Federating Grids”: SPICE, Cactus BH simulation
 - not scalable
 - Top-down approach
- e-Science Applications must be able to utilise infrastructure
 - Handle current heterogeneity and future developments
- “Real” eScience Applications: Interoperability(!)




"Global" Grid Infrastructure

CENTER FOR COMPUTATION
& TECHNOLOGY

US TeraGrid



NGS



Netherlight
(Amsterdam)

UKLight

App Sci

All sites connected by
production network

DEISA

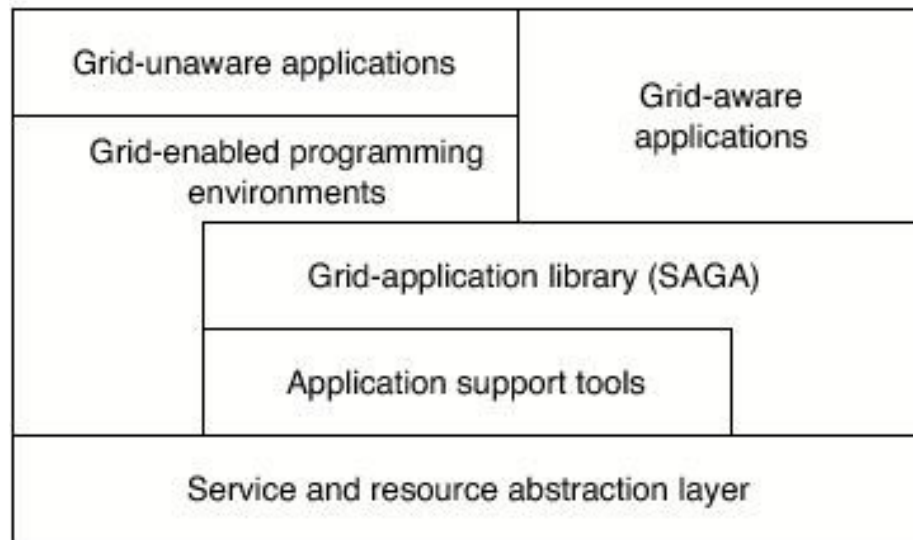
Distributed
European
Infrastructure for
Supercomputing
Applications

● Computation ● Network PoP ● Visualization



Application Level Interoperability (1)

Rough Taxonomy of Applications



- Some applications are Grid-unaware and want to remain so
 - Use tools/environments (e.g, NanoHub, GridChem)
 - May run on Grid-aware/Grid-enabled environments (e.g. Condor) or programming environment (e.g, MPICH-G2)
- Some applications are explicitly Grid-aware
 - Control, Interact & Exploit distributed systems at the *application level*



Application Level Interoperability (2)

CENTER FOR COMPUTATION
& TECHNOLOGY

- Some features (not definition!) of ALI:
 - Beyond compiling, no further changes for platform
 - Automated, scalable and extensible solutions to use new **resources**
 - Not via bilateral or bespoke arrangements
 - Semantics of any services are consistent (eg error handling)
 - ALI often requires QoS, co-scheduling
 - eg. remote visualisation, multiple-resources
 - In general, GU applications require mostly SLI
 - establishing interoperability for GU is easier than GA applications



SAGA: In a Nutshell

- A lack of:
 - Programming interface that provides common grid functionality with the correct level of abstractions?
 - Ability to hide underlying complexities, varying semantics, heterogenities and changes from application program(er)
- Simple, integrated, stable, uniform, high-level interface
- Simplicity: Restricted in scope, 80/20
- Measure(s) of success:
 - Does SAGA enable quick development of “new” distributed applications?
 - Does it enable greater functionality using less code?



Copy a File: Globus GASS

CENTER FOR COMPUTATION

```
& TECHNOLOGY
int globus_gass_copy_file (char const* source, char const* target)
{
    globus_url_t                source_url;
    globus_io_handle_t          dest_io_handle;
    globus_ftp_client_operationattr_t source_ftp_attr;
    globus_result_t              result;
    globus_gass_transfer_requestattr_t source_gass_attr;
    globus_gass_copy_attr_t      source_gass_copy_attr;
    globus_gass_copy_handle_t    gass_copy_handle;
    globus_gass_copy_handleattr_t gass_copy_handleattr;
    globus_ftp_client_handleattr_t ftp_handleattr;
    globus_io_attr_t             io_attr;
    int                           output_file = -1;

    if ( globus_url_parse (source_URL, &source_url) != GLOBUS_SUCCESS ) {
        printf ("can not parse source_URL \"%s\"\n", source_URL);
        return (-1);
    }

    if ( source_url.scheme_type != GLOBUS_URL_SCHEME_GSIFTP &&
        source_url.scheme_type != GLOBUS_URL_SCHEME_FTP &&
        source_url.scheme_type != GLOBUS_URL_SCHEME_HTTP &&
        source_url.scheme_type != GLOBUS_URL_SCHEME_HTTPS ) {
        printf ("can not copy from %s - wrong prot\n", source_URL);
        return (-1);
    }

    globus_gass_copy_handleattr_init (&gass_copy_handleattr);
    globus_gass_copy_attr_init (&source_gass_copy_attr);

    globus_ftp_client_handleattr_init (&ftp_handleattr);
    globus_io_fileattr_init (&io_attr);

    globus_gass_copy_attr_set_io (&source_gass_copy_attr, &io_attr);
    globus_gass_copy_handleattr_set_ftp_attr
        (&gass_copy_handleattr,
         &ftp_handleattr);
    globus_gass_copy_handle_init (&gass_copy_handle,
                                   &gass_copy_handleattr);
```

```
    if (source_url.scheme_type == GLOBUS_URL_SCHEME_GSIFTP ||
        source_url.scheme_type == GLOBUS_URL_SCHEME_FTP ) {
        globus_ftp_client_operationattr_init (&source_ftp_attr);
        globus_gass_copy_attr_set_ftp (&source_gass_copy_attr,
                                         &source_ftp_attr);
    }
    else {
        globus_gass_transfer_requestattr_init (&source_gass_attr,
                                                source_url.scheme);
        globus_gass_copy_attr_set_gass(&source_gass_copy_attr,
                                         &source_gass_attr);
    }

    output_file = globus_libc_open ((char*) target,
                                    O_WRONLY | O_TRUNC | O_CREAT,
                                    S_IRUSR | S_IWUSR | S_IRGRP |
                                    S_IWGRP);

    if ( output_file == -1 ) {
        printf ("could not open the file \"%s\"\n", target);
        return (-1);
    }
    /* convert stdout to be a globus_io_handle */
    if ( globus_io_file_posix_convert (output_file, 0,
                                       &dest_io_handle)

        != GLOBUS_SUCCESS) {
        printf ("Error converting the file handle\n");
        return (-1);
    }

    result = globus_gass_copy_register_url_to_handle (
        &gass_copy_handle, (char*)source_URL,
        &source_gass_copy_attr, &dest_io_handle,
        my_callback, NULL);
    if ( result != GLOBUS_SUCCESS ) {
        printf ("error: %s\n", globus_object_printable_to_string
            (globus_error_get (result)));
        return (-1);
    }
    globus_url_destroy (&source_url);
    return (0);
}
```




SAGA Example: Copy a File

High-level, uniform

```
#include <string>
#include <saga/saga.hpp>

void copy_file(std::string source_url, std::string target_url)
{
    try {
        saga::file f(source_url);
        f.copy(target_url);
    }
    catch (saga::exception const &e) {
        std::cerr << e.what() << std::endl;
    }
}
```

- Provides the high level abstraction, that application programmers need; will work across different systems
- Shields gory details of lower-level middle-ware and system issues
- Like MapReduce – leave details of distribution *etc.* out

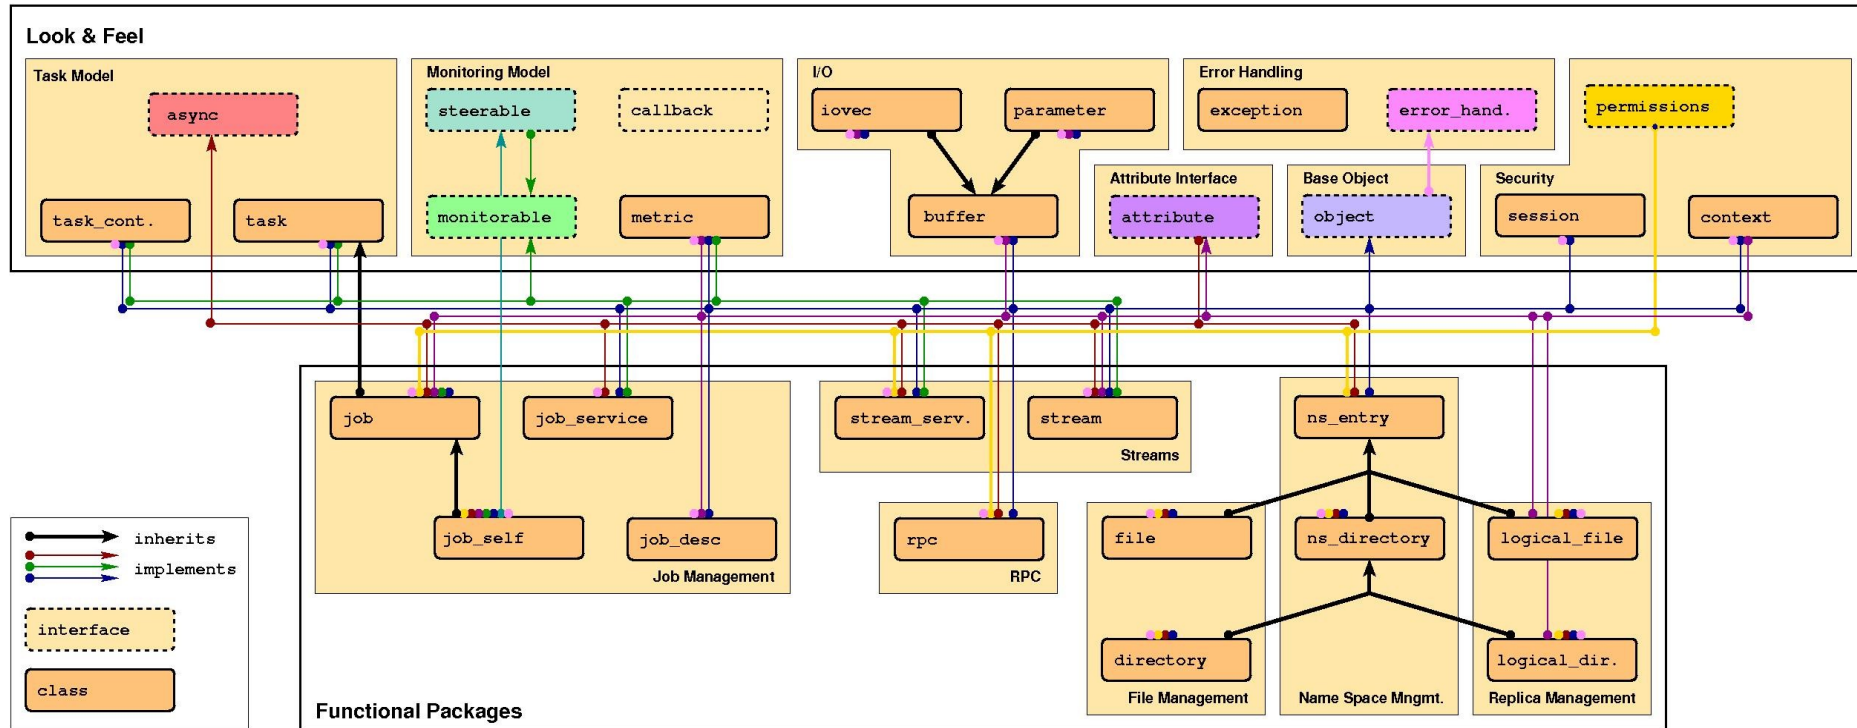


SAGA: Scope

- Is:
 - Simple API for Grid-Aware Applications
 - Deal with distributed infrastructure explicitly
 - High-level (= application-level) abstraction
 - An uniform interface to different middleware(s)
 - Client-side software
- Is NOT:
 - Middleware
 - A service management interface!
 - Does not hide the resources - remote files, job (but the *details*)



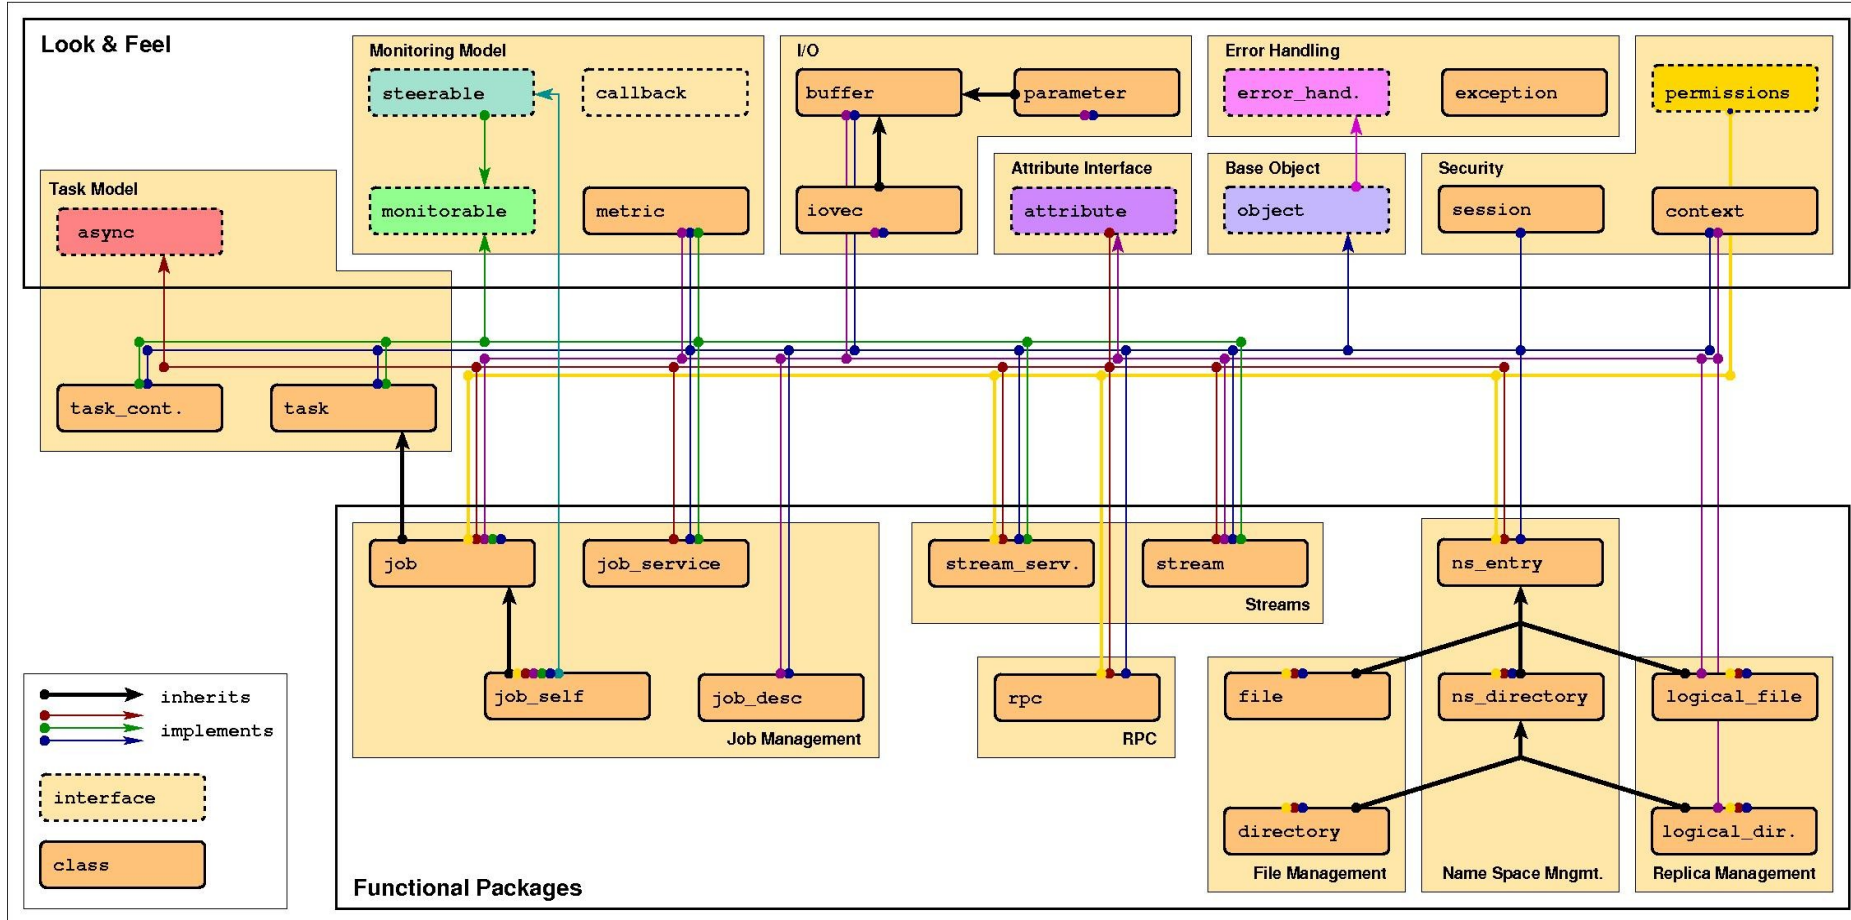
SAGA: Class Diagram





SAGA: Class Diagram (2)

CENTER FOR COMPUTATION
& TECHNOLOGY





CENTER FOR COMPUTATION
& TECHNOLOGY

SAGA API: Towards a Standard

Standards help Interoperability

- The need for a standard programming interface
 - “Go it alone” versus “Community” model
 - Reinventing the wheel again, yet again, and again
 - MPI as a useful analogy of community standard
 - OGF the natural choice; establish SAGA-RG
- “Tedium” of the standardisation process?
 - Not all technology needs to be standardised upfront
 - Standardisation not a guarantee to success
- Requirements Document
 - Design and requirements derived from 23 Use Cases
 - Different projects, applications and functionality



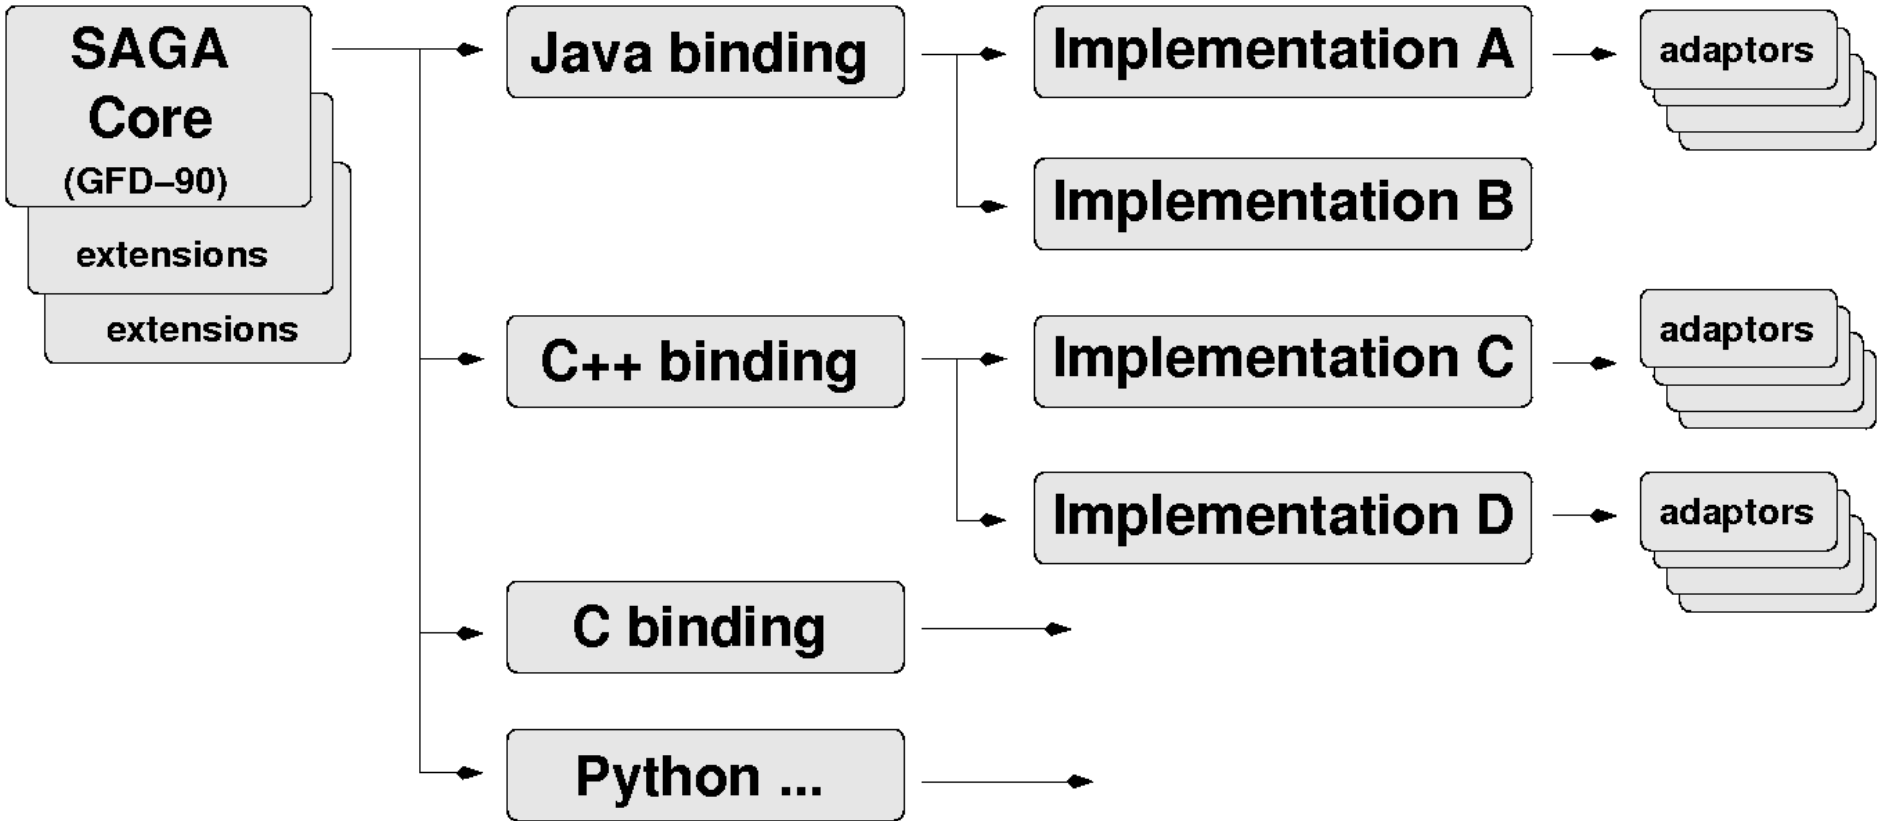
SAGA API: Design & Specification

Process enabling Interoperability

- Open Grid Forum: Design and requirements derived from 23 Use Cases (some “correlated”)
 - Different projects, applications and functionality
 - Biological, Coastal-modelling, visualization ..
 - Functional Areas: Job Mgmt, Resource Mgmt, Data Mgmt, Logical Files, Streams....
 - Non-functional Areas: Asynchronous, QoS, Bulk
- Interface is language independent, object-oriented and each sub-system is independent
- Specified using Scientific Interface Description Language (SIDL)
 - extensible and easily implementable



The SAGA Landscape





SAGA Interface: Providing Interoperability Job Submission API

```
01: // Submitting a simple job and wait for completion
02: //
03: saga::job_description jobdef;
04: jobdef.set_attribute ("Executable", "job.sh");
05:
06: saga::job_service js;
07: saga::job job = js.create_job ("remote.host.net", jobdef);
08:
09: job.run();
10:
11: while( job.get_state() == saga::job::Running )
12: {
13:     std::cout << "Job running with ID: "
14:               << job.get_attribute("JobID") << std::endl;
15:     sleep(1);
16: }
```




SAGA Implementation

How it supports Interoperability

- Requirements
- Overall Architecture
 - Horizontal, Extensibility, Vertical
- Generic Call Routing
 - Route API calls to appropriate MW adaptors; underlying technology is Abstract Objects
- Lessons Learnt
 - Uniform implementations for different languages; C, C++, Python and Java
 - Flexible adaptor selection (more than one)



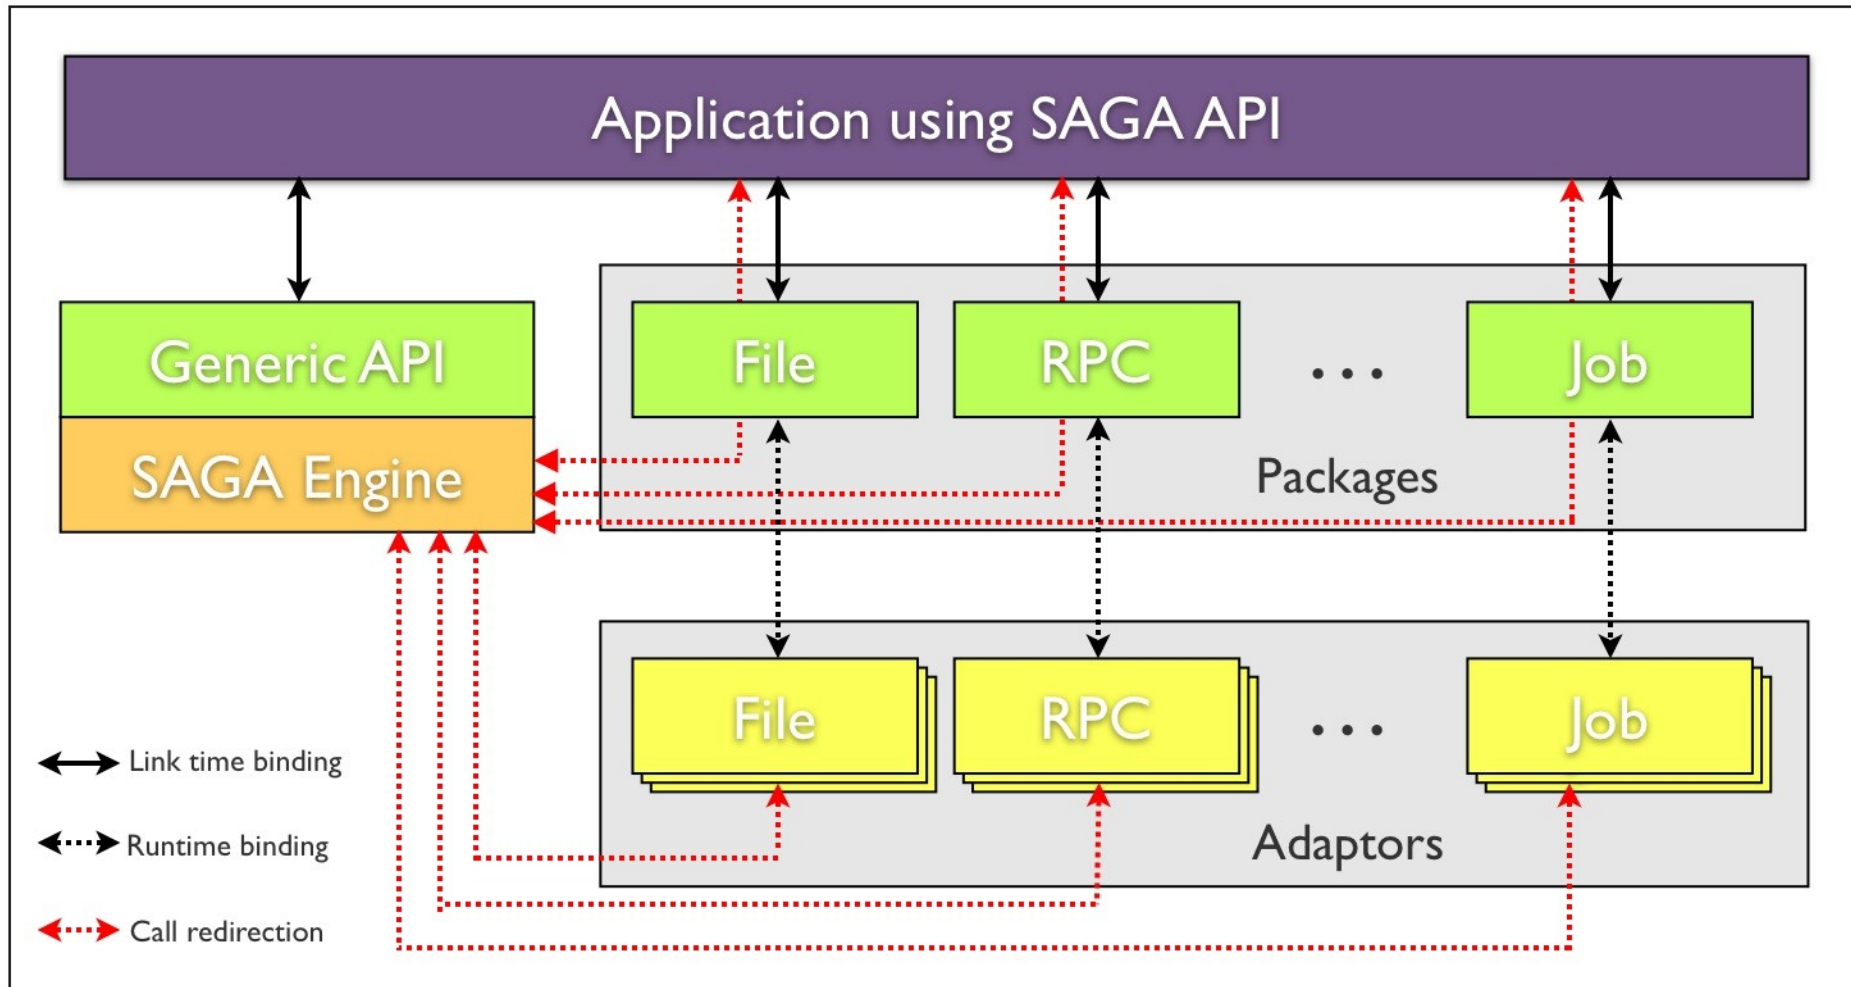
SAGA Implementation (2)

Some requirements to support interoperability

- Non-trivial set of requirements:
 - Allow heterogenous middleware to co-exist
 - Cope with evolving grid environments; dyn resources
 - Future SAGA API extensions
 - Portable, syntactically and semantically platform independent; permit latency hiding mechanisms
 - Ease of deployment, configuration, multiple-language support, documentation etc.
 - Provide synchronous, asynchronous & task versions
- Portability, modularity, flexibility, adaptability, extensibility



SAGA C++ (LSU)





SAGA Implementation

How the Architecture Supports Interoperability

- Horizontal Extensibility – API Packages
 - Current packages:
 - file management, job management, remote procedure calls, replica management, data streaming
 - Steering, information services, checkpoint in pipeline
- Vertical Extensibility – Middleware Bindings
 - Different adaptors for different middleware
 - Set of ‘local’ adaptors
- Extensibility for Optimization and Features
 - Optimization: e.g. Bulk optimization, modular design
 - Features: e.g latency hiding, late binding etc.
 - PIMPL = Private (API..) + Implementation (Engine)



SAGA Implementation

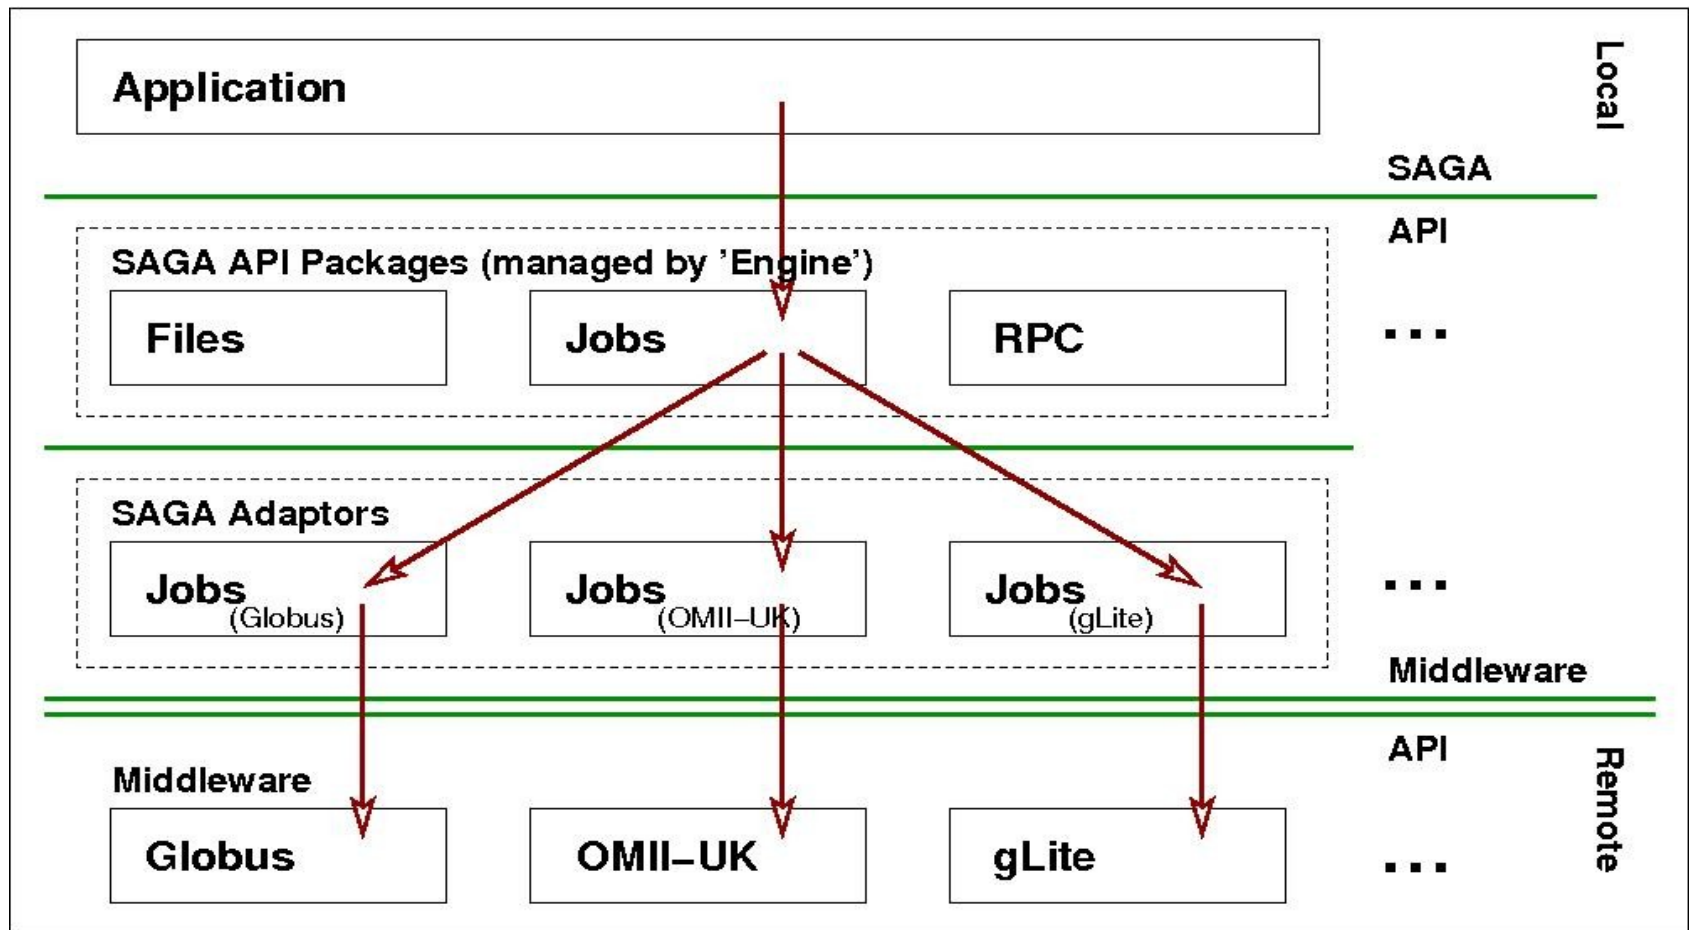
How it supports Interoperability

- Requirements
- Overall Architecture
 - Horizontal, Extensibility, Vertical
- Generic Call Routing
 - Route API calls to appropriate MW adaptors;
underlying technology is Abstract Objects
- Lessons Learnt
 - Uniform implementations for different languages;
C, C++, Python and Java
 - Flexible adaptor selection (more than one)



SAGA: Interoperable Job Submission

Role of Adaptors (middleware binding)





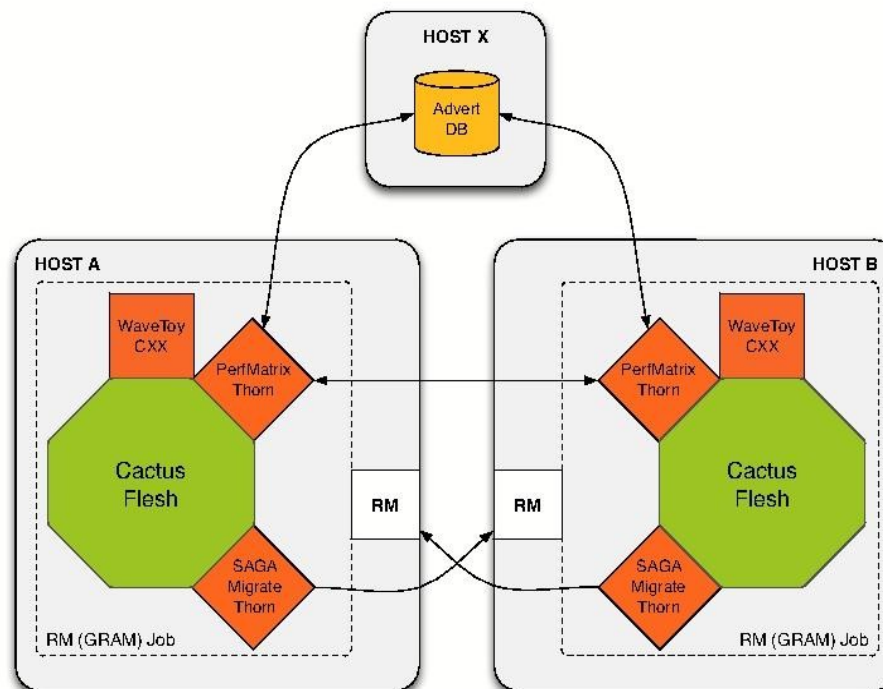
SAGA and (the five areas of) GIN

- SAGA and ...
 - Job Submission & Management
 - Provides same functionality as OGSA-BES, but at a higher level of abstraction, e.g., job id w/o contacting job service
 - Same job state model as BES (not a random coincidence)
 - Data Management and Movement
 - Data movement protocol agnostic
 - Supports pseudo-schema (`any://remote.host.net/dir/file.typ`) and allows implementation to determine
 - Authorization and Identity Management
 - Essentially not exposed at application development level
 - Cross Grid Applications
 - Information Services and Modelling




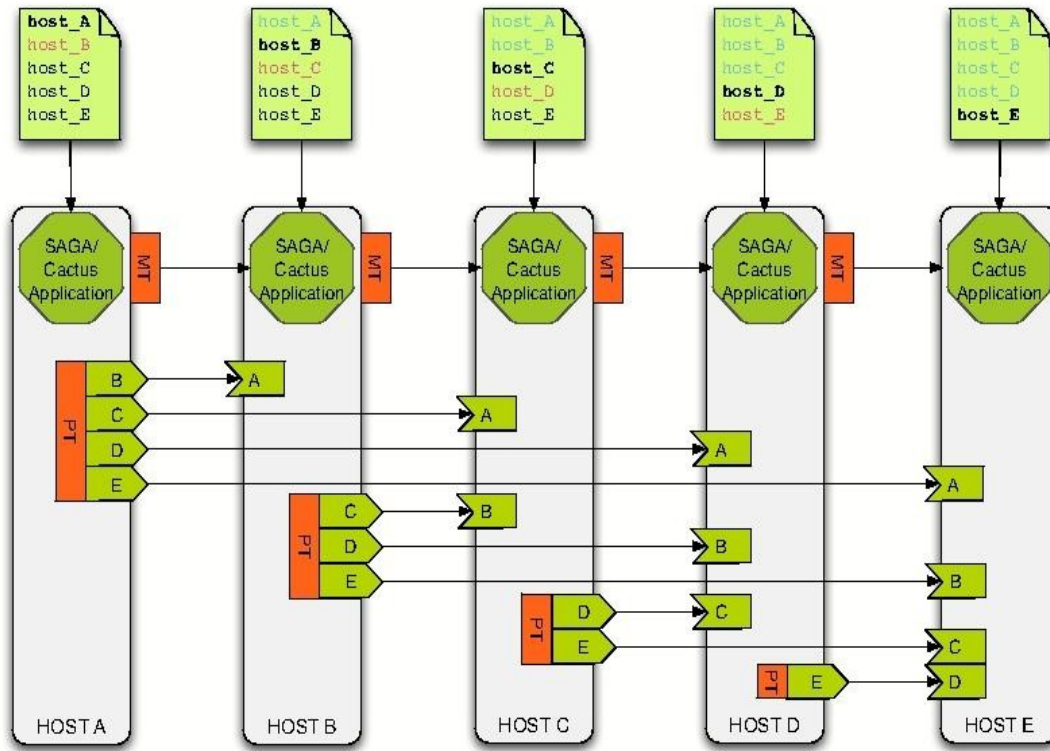
Network Performance Aware Application (1)

- Capable of acquiring **application-specific** network characteristic data and determining **ideal** migration target across **heterogeneous** grids **without changes** in the application





Network Performance Aware Application (2)





Future Work

- Many applications under development...
- Middleware bindings (adaptors): Key to “realising” interoperability
- Use SAGA based applications in the next generation of GIN demonstrations(!)
 - Suggestions for applications?
 - US TeraGrid and EU Deisa?



Acknowledgments

- Funding Agencies:
 - UK EPSRC:
 - Through OMII-UK “OMII SAGA” project
 - Through NeSC, Edinburgh Distributed Programming Abstractions
 - US NSF
 - Cybertools projects
 - CCT Internal Funds



CENTER FOR COMPUTATION
& TECHNOLOGY



Hartmut Kaiser



Andre Merzky



Ole Weidner

+ Many other students



Thilo Kielmann



Cerial Jacobs



Kees Verstop



Upcoming API Extensions

- MessageBus
 - Structured data transfer, also many-to-many
- Service Discovery
 - Based on GLUE schema
- Adverts
 - Persistent storage of application-level data
- Checkpointing/Recovery
 - Based on GridCPR