Adaptive Distributed Replica-Exchange Simulations

BY ANDRE LUCKOW[1], SHANTENU JHA[2,3,4], JOOHYUN KIM[2], ANDRE MERZKY[2]
AND BETTINA SCHNOR[1]

[1]*Institute of Computer Science, Potsdam University, Germany*
[2]*Center for Computation & Technology, Louisiana State University, USA*
[3]*Department of Computer Science, Louisiana State University, USA*
[4]*e-Science Institute, Edinburgh, UK*

Due to the loose-coupling between replicas, the Replica-Exchange class of algorithms should be able to benefit greatly from utilising as many resources as available. However, the ability to effectively utilise multiple distributed resources to reduce the time-to-completion remains a challenge at many levels. Additionally, an implementation of a *pleasingly-distributed* algorithm such as Replica-Exchange, which is independent of infrastructural details does not exist. This paper proposes an extensible and scalable framework based on SAGA that provides a general-purpose, opportunistic mechanism to effectively utilise multiple resources in an infrastructure independent way. By analysing the requirements of the Replica-Exchange algorithm and the challenges of implementing it on real production systems, we propose a new abstraction (BigJob), which forms the basis of the adaptive redistribution and effective scheduling of replicas.

1. Introduction

Several classes of applications are well suited for distributed environments. Probably the best known and most powerful examples are those that involve an ensemble of decoupled tasks, such as simple parameter sweep applications (Casanova et al. 2000). A slightly more complicated and challenging class of distributed applications are those that have a degree of coupling between individual sub-tasks. An interesting example of such applications are those based upon the *Replica-Exchange (RE)* (Hansmann 1997, Sugita and Okamoto 1999) algorithm. RE simulations are used to understand physical phenomena – ranging from protein folding dynamics to binding affinity calculations.

The RE method involves the concurrent execution of multiple similar simulations – the *replicas*. The coupling between the replicas occurs via periodic exchange attempts between paired replicas. The exchange is typically infrequent compared to the run-time of each replica, and is small in terms of communication bandwidth requirements. Thus, RE is *prima facie* a perfect algorithm to exploit distributed resources. We label such a class of algorithms as *pleasingly-distributed*.

Most RE implementations are either infrastructure specific (Woods, C. et al. 2005), or if using multiple distributed resources, require prior co-scheduling (Manos et al. 2008). Ref. (Manos et al. 2008) is an important example of a first-generation Grid application, wherein the effectiveness of coupling multiple distributed resources for scientific problems has been demonstrated. The real power of distributed systems however, arises from

adaptive algorithms and implementations that provide applications with an agile execution model, and thus the ability to utilise resources dynamically, as opposed to a static execution model inherited from parallel and cluster computing. Unfortunately, the barrier to the development of such adaptive applications is high and the infrastructure support is poor. Specifically, there is no implementation of an adaptive RE algorithm, which is both able to effectively and reliably utilise multiple distributed resources without prior scheduling as well as being independent of any specific infrastructure.

In this paper we address some of the challenges and performance bottlenecks encountered when performing RE simulations over multiple distributed resources, such as the overall slowdown due to synchronisation arising from the light-coupling and the lack of co-scheduled resources. The unique contribution of this paper is the implementation of a RE framework that overcomes the described limitations by being able to adapt at runtime to a change in the availability of resources and application resource requirements. The framework builds upon preliminary work of integrating SAGA and Migol to provide fault-tolerance. While SAGA represents a well-defined, standardized interface for writing Grid applications, Migol provides the underlying middleware services to guarantee the correct and reliable execution of applications even in the presence of failures (Luckow et al. 2008).

We provide evidence that, as more resources become available, our framework can opportunistically utilise these resources, leading to a reduction in the time-to-completion of the scientific problem. The remainder of the paper is structured as follows: In the next section we provide the basic ideas and advantages of using RE simulations to understand physical properties of a RNA system. In section 3 we describe the details of the REMD framework implementation. Section 4 discusses the new BigJob abstraction and the SAGA Glide-In framework. In section 5 we describe the deployment and results of different experiments using the SAGA based RE framework on the TeraGrid, and in section 6 we present data establishing the advantages of REMD for the physical system under study.

2. Hepatatis-C Virus (RNA) Using Replica-Exchange

In Molecular Dynamics (MD) approaches, sufficient sampling of configurations is an important requirement for connecting atomistic results to macroscopic or thermodynamic quantities available from experiments. This provides an important motivation for researching ways to accelerate sampling and to enhance the "effective" time-scales studied. Generalized ensemble approaches – of which REMD (Sugita and Okamoto 1999) is a prominent example – represent a promising attempt to overcome the general limitations of insufficient time-scales, as well as specific limitations of inadequate conformational sampling arising from kinetic trappings. The fact that one single long-running simulation can be substituted for an ensemble of loosely-coupled shorter-running simulations, make these ideal candidates for distributed environments.

RE simulations consist of two distinct components: the underlying simulation engine used for each replica, and the coupling-mechanism between the individual replicas. The degree and frequency of coupling and exchange can be either regular (Sugita and Okamoto 1999), or irregular (Shirts and Pande 2001). An example of the latter – parallel replica dynamics as implemented in Folding@home (Fol 2008) involves coordination between replicas only when an "event" occurs. In contrast, for regular RE applications, attempts to exchange states between certain pairs occur at fixed intervals.

The hepatitis C virus (HCV) internal ribosome entry site (IRES) is recognized specifically by the small ribosomal subunit and eukaryotic initiation factor 3 (eIF3) before viral

translation initiation. This makes it a good candidate for new drugs targeting Hepatitis-C virus. Our aim is to use REMD to enhance the sampling of the conformational flexibility of the internal loop referred to as *HCV IRES IIIb CA variant* (Collier et al. 2002) as well as the equilibrium energetics. The model of the physical system under investigation in this work is comprised of a RNA system of nucleotides; the total number of atoms in the simulating box is 21887 – including the RNA system, explicit water molecules, and ions for neutralization of the system. The initial conformation of the RNA is taken from the NMR structure (PDB ID: 1PK7).

3. Implementing Distributed Replica-Exchange Using SAGA/Migol

RE-Manager Architecture: The framework comprises of three components, the RE-Manager, the Replica-Agent and the Migol infrastructure. The *RE-Manager*, also referred to as task manager, is deployed on the user's desktop and provides the interface to the overall RE run. It orchestrates all replicas, which involves file staging, job spawning and the conduction of the exchange attempts, using the SAGA APIs.

The second element is the task agent, the *Replica-Agent*, that resides on the machines where the RE replicas are executed. The Replica-Agent is launched using SAGA CPR and Migol. NAMD (Phillips, J. et al. 2005), a highly scalable, parallel MD code, is used to carry out the MD simulation corresponding to each replica run. It is important to mention that any other MD or Monte Carlo code could be used just as simply and effectively. Finally, Migol handles the reliable execution of the Replica-Agent and the replicas, i. e. the submission, the monitoring and, if required, the recovery of replicas or the application itself.

Replica-Exchange Logic: RE simulations involve the running of multiple replica jobs. Each replica is assigned a different temperature. Depending on the number of processes n, the RE-Manager creates $\frac{n}{2}$ pairs of replicas. Before launching a job the RE-Manager ensures that all required input files are transferred to the respective resource. For this purpose, the SAGA File API and the GridFTP adaptor are used. The replica jobs are then submitted to the resource using the SAGA CPR API and the Migol/GRAM middleware.

When all replicas reach a pre-determined state (e. g. the NAMD job finishes after a fixed number of steps), the decision as to whether to pairwise exchange temperatures between neighbouring replicas is determined by the Metropolis scheme. The run of an ensemble of replicas in parallel and the subsequent pairwise exchange attempt is referred to as *generation*. No two replicas can belong to different generations. If the exchange attempt is successful, parameters such as the temperature, are swapped. Both jobs are then relaunched. Often the Metropolis scheme returns a negative result, and an exchange is not carried out; thus it is difficult to respond to a possible exchange speculatively.

Deploying on Production Environments: The RE framework has been successfully deployed on LONI and TeraGrid production resources (Luckow et al. 2008). In these environments a significant slowdown was observable, in particular when running a larger number of replicas. A major reason for this slowdown was the fact that the re-started replicas are required to queue again at the local scheduler. In pathological cases, the complete system came to a halt solely due to a single crowded or slow resource.

To avoid such bottlenecks, the multiple sub-tasks that comprise distributed applications need to avoid re-queuing at the system batch-queue level. Additionally, distributed applications that are decomposable into sub-tasks should be able to respond to the dynamic availability of resources. Unfortunately, current infrastructures do not support such dynamic scheduling directly. To provide this capability to applications, we need (i) ab-
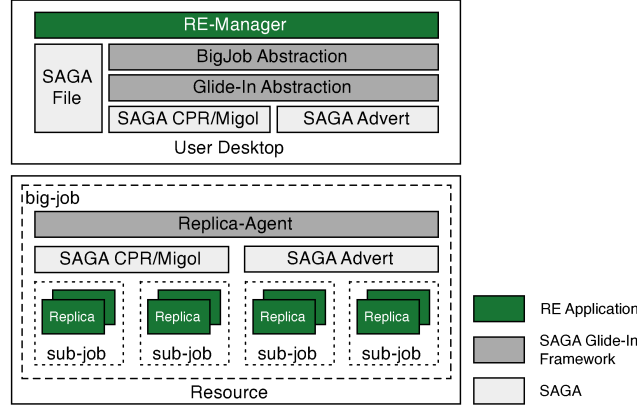
Figure 1: **RE-Manager Abstractions: The BigJob abstraction provides the capability to cluster sub-jobs into a larger big-job, and is implemented on top of the Glide-In abstraction.**

stractions that enable agile execution models via application-level allocation of resources, and (ii) different adaptivity strategies that determine how resources are efficiently utilised. *The following section describes the extensions to the simple RE framework that enables efficient scheduling of sub-tasks and supports adaptive applications.*

4. Adaptive Replica-Exchange: Abstractions and Implementation

As motivated before, the use of multiple *simple* Grid jobs to execute many replicas has a severe limitation: all simple jobs must queue at the resource management system, i. e. a single delayed job can cause an overall slowdown. We overcome this issue by using an efficient dispatching scheme, which builds upon the ability to cluster replicas using the novel BigJob abstraction before submission. Based on this abstraction, we propose different strategies that address the dynamic conditions of distributed environments.

Abstractions: A common approach to avoid queuing delays is the use of place-holder jobs, which are able to dispatch several sub-jobs without each sub-job needing to queue at the local scheduler. A specific mechanism to support this pattern is the *Glide-In* abstraction, in reference to the Condor Glide-In system (Frey et al. 2002), which pioneered this idea. A Glide-In job requests a sufficiently large chunk of resources; smaller sub-jobs can then rapidly be executed through the Glide-In job. By avoiding the high initial costs for queueing each individual replica job, the time-to-completion can be dramatically reduced.

Figure 1 summarizes the abstractions developed and used in this work to support the clustering of sub-jobs into larger big-jobs and the effective dispatching of the sub-jobs. The specific capability to cluster sub-jobs, is provided to the application via the *BigJob* abstraction. The SAGA Glide-In abstraction is used to support the commonly occurring place-holder job pattern. The BigJob abstraction defines a `big_job` and `sub_job` object; for each `big_job` object, a Glide-In job with the desired number of resources is started, and `sub_job` objects, which correspond to individual replicas, are mapped to a `big_job` using the jobid as reference. It is helpful to reiterate that although there is a `big_job` object, it is submitted as a Glide-In job. Also, the BigJob abstraction in turn utilises the Glide-In abstraction to map the individual big-job and sub-jobs to physical resources.

Implementation: The RE framework has been extended to support the BigJob and Glide-In abstractions. BigJob provides the ability to cluster sub-jobs; Glide-In allows the effective scheduling of the sub-jobs. As illustrated in Figure 2, the RE-Manager uses the `big_job`
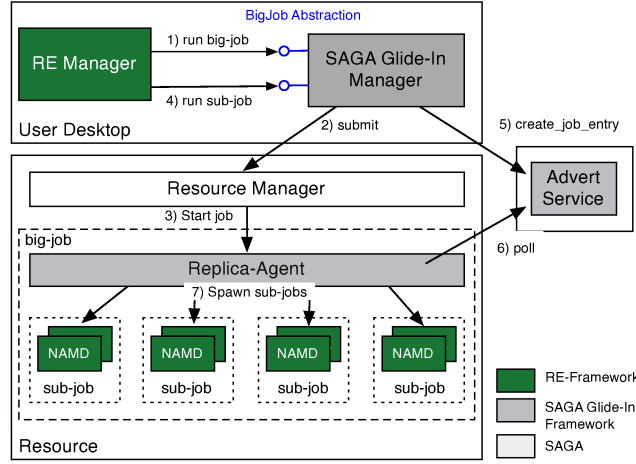
Figure 2: **RE-Manager and SAGA-Glide-In Framework: The Glide-In job (Replica-Agent) is used as place-holder job for all replica sub-jobs running on a single cluster. The RE-Manager can control both the Replica-Agent and the replica jobs.**

and `sub_job` objects as replacements for the `job` object defined by SAGA CPR. The `big_job` and `sub_job` objects behave similarly to regular SAGA CPR job objects. Thus, the RE-Manager, which is the application in this case, does not require any extensive modification; all that the RE-Manager has to do is to provide a mapping from a `sub_job` to a suitable `big_job` via a jobid.

The SAGA Glide-In implementation comprises of three components: 1) The *Glide-In Manager*, which provides the Glide-In abstraction and manages the orchestration and scheduling of Glide-In jobs (which in turn allows the management of both big-job objects and sub-jobs) and 2) The *Replica-Agent*, which represents the Glide-In job and thus, the application-level resource manager on the respective resource, and 3) the *Advert Service*, which is used for communication between the Glide-In Manager and Replica-Agent.

Before running regular jobs, an application, in this case the RE-Manager, must initialize a `big_job` object. The Glide-In Manager then queues a Glide-In job, which actually runs a Replica-Agent on the respective resource. For this `big_job` instance the specified number of resources is requested. Subsequently, `sub_job` objects can be submitted through the Glide-In Manager using the jobid of the `big_job` as reference. The Glide-In Manager ensures that the sub-jobs are launched onto the correct resource based upon the specified jobid using the right number of processes.

Communication between the Replica-Agent and Glide-In Manager is carried out using the SAGA Advert Service, a central key/value store. For each new job an advert entry is created by the Glide-In Manager. The Replica-Agent periodically polls for new jobs. If a new job is found and resources are available, the job is dispatched, otherwise it is queued. Further, the agent encapsulates local machine-specific settings. The Replica-Agent ensures e. g. that the right combination of compiler, MPI library and NAMD executable is used.

Adaptive Replica Scheduling: Distributed applications including RE simulations must be able to deal with time-varying resource availabilities. An application is referred to as *dynamic* when either its resource requirements, or the availability and utilisation of resources

changes during its runtime. *Adaptivity* is a mechanism to respond to dynamic changes; a dynamic application may deploy multiple adaptive strategies or choose between competing adaptive strategies. For an application to be adaptive, it is necessary for it to be able to effectively utilise an expanded or reduced set of resources; additionally for an adaptive application to be scalable it must also be able to determine which resources to utilise efficiently. For resource determination, our framework currently relies on a static, user-defined mapping of replicas and resources. In the remainder of this paper, we will focus on dynamic resource utilisation (and not on dynamic resource determination or optimisation).

For jobs that want to maximise their throughput the ability to adapt to dynamically changing resource loads is critical. It is equally important for long-running applications to be able to support an agile execution model allowing the effective utilisation of resources as they become available. Specifically, there are different ways a RE simulation can respond to a change in the number of resources required/available:

- *Scenario A:* By increasing the number of processes assigned to each replica the time-to-completion can be reduced. In addition, resources can be partitioned in a way that balances the different speeds of resources. For example, by adding processes to a delayed replica, bottlenecks due to synchronisation of replicas can be avoided.
- *Scenario B:* As resources become available, the number of replicas can be adjusted. Depending on the underlying physics model, the additional replicas can be used to either refine the temperature range (adaptive sampling) or to extend the temperature range (enhanced dynamics). This REMD approach is also referred to as *Cool Walking* (Brown and Head-Gordon 2003).

Our framework supports both adaptive strategies.

5. Distributed Replica-Exchange on the TeraGrid

To evaluate the performance of the RE-Manager several experiments have been conducted on TeraGrid (TG) and LONI resources. The resources used are: Ranger (TG), Abe (TG) and QueenBee (QB; both TG and LONI). The RE-Manager was configured to run a parallel NAMD simulation with up to 16 replicas sampling a temperature range between 300 and 450 K. Replica exchanges are carried out between pairs of replicas. Thus, there are up to 8 exchanges in each generation. Each test run comprises of up to 64 attempted exchanges; each replica can use up to 24 MPI processes and runs for 500 time steps between exchange attempts. The metric used is the time-to-completion for 64 attempted exchanges.

Initially, we investigated the effect of the Glide-In framework on the time-to-completion (T_c) using 16 replicas on QB. Using the Glide-In framework, there was a reduction in T_c from 52 minutes to 26 minutes per average, which corresponds to a decrease of 50 %. In the best case, improvements of up to 70 % were observed. This effect is attributed to the elimination of the queuing times for every sub-job. Once the Replica-Agents become active, replicas can be dispatched without requiring interactions with the local scheduler.

Further, we performed tests for the two adaptive scenarios. In the scenario A, the number of replicas was kept constant (Conventional REMD) and the *replica size*, i. e. the number of MPI processes assigned to each replica was varied as more resources became available. In scenario B, the *replica number*, i. e. the number of replicas participating in a generation, was varied (Cool Walking). We compare T_c for 64 attempted exchanges on different sets of distributed resources and Glide-In configurations.

In scenario A, up to 3 different resources are used, the number of Glide-Ins varies from 4 up to 12, whilst the number of replicas used is fixed at 16. Thus, the size of the
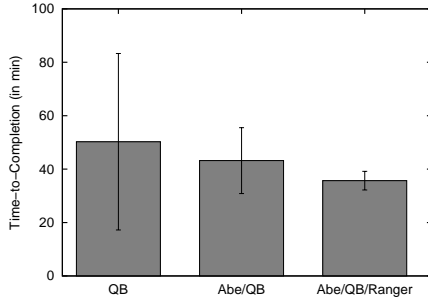
Figure 3: **Replica Size Adaptivity (Scenario A): As the number of resources available increases, the number of cores assigned to a replica (a NAMD job) is dynamically adjusted, while keeping the number of replicas constant at sixteen. There are 4 Glide-Ins with 32 cores each. As more resources become available, more cores are assigned to each replica, which leads to a reduction of T_c.**
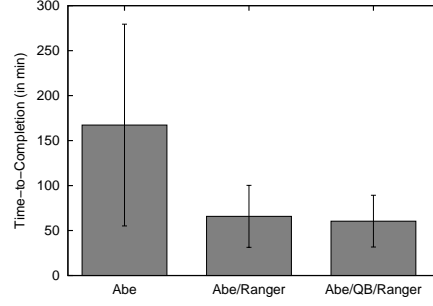
Figure 4: **Replica Number Adaptivity (Scenario B): In this scenario, the framework dynamically adjusts the number of replicas. A constant number of 4 Glide-Ins with 64 cores each are distributed across 1, 2 and 3 machines; the size of each replica is kept fixed at 16 cores. Once again, the greater the number of distributed resources that can be utilised, the smaller T_c.**

individual replicas is varied. When a resource is being used, it runs 4 Glide-Ins and each Glide-In job has a constant size of 32 cores. Although, the number of Glide-In jobs on a resource is fixed at 4 for reasons of simplicity, our results will hold for general values. If all resources (Abe, QB and Ranger) were being used, there would be 12 Glide-In jobs with 32 cores each, i. e. a total of 384 cores would be available. Glide-Ins are submitted to 1, 2 or 3 statically configured resources. After submission the Glide-Ins are subject to different queueing delays at the local schedulers. To reflect these different loads, the number of MPI processes assigned to each replica is dynamically increased as new resources become available. Depending on the number of available resources, between 8 (for 4 Glide-Ins) and 24 MPI processes (for 12 Glide-Ins) are used for each replica.

Figure 3 shows the results of the distributed run. In spite of the overhead for migrating replicas to newly available resources, a notable decrease in T_c of up to 15 minutes can be observed as the number of resources increases from one to three. Although, the efficiency, defined as runtime on one resource divided by the runtime on multiple resources scaled by the number of resources, is only about 0.5, this is a limitation of the used setup rather than a general scalability barrier. The setup comprises of rather short NAMD jobs; in particular during the initial phase jobs are often migrated to other resources, which mainly causes this overhead. During longer runs this overhead is negligible. What is also very important to note is the reduced fluctuation in the T_c when multiple resources are used. This is indicative of the fact that there is a reduction in the sensitivity to queue-loads – something that applications on real production environments have to battle with.

In scenario B, the capability of the RE-Manager to adaptively adjust the number of replicas by varying either the range of temperature or the specific temperatures simulated is evaluated. For this scenario, 4 Glide-Ins with 64 cores each are distributed across either 1, 2 or 3 different distributed resources. That means that the total of 256 cores is allocated on (i) Abe only, (ii) Abe and Ranger (iii) Abe, Ranger and QB. Each replica has a constant number of 16 MPI processes, and thus any Glide-In can run exactly 4 replicas when active.

At the beginning of the experiment, all 4 Glide-Ins are submitted to either 1, 2 or 3 resources (statically configured). Similar to scenario A, different queueing delays usually occur. Consequently, not all Glide-Ins start simultaneously. Using the adaptive temperature

sampling algorithm, the number of replicas is dynamically varied as the number of active Glide-Ins increases. Depending on the number of running Glide-Ins, the ensemble consists of 4 to 16 replicas (increase in steps of 4).

As shown in Figure 4, T_c decreases with the number of resources used. With the the simulation distributed onto greater number of resources, the probability that a single heavily-utilised resource delays the overall progress of the simulation is reduced. The results clearly demonstrate the benefits of the adaptive replica scheme – it is favourable to instantly use resources as they become active instead of waiting for the complete set of nodes to become available.

6. Results: Enhanced Sampling of Hepatitis-C Virus

In the previous sections we discussed the design and performance of the framework and showed how it enables the effective utilisation of multiple resources. In section 2, we provided motivation for why the RE approach is required to understand the energetics and conformational flexibility of the internal loop of the Hepatitis-C virus. The effectiveness of the RE approach in increasing the rate of convergence to equilibrium (Boltzmann distribution) or enhancing the sampling can be measured by the rate of attempted exchanges (Lei and Duan 2007), or equally by the inverse of the *time* between attempted exchanges.

REMD simulations (Scenario A) were performed for HCV IRES IIIb CA variant using the model described in section 2. The replicas covered a temperature range from 300 K to 450 K, thus fixing the number of replicas to 16. To demonstrate the effectiveness of our RE framework, we analyse a typical time-series of the number of resources used (available) and thus, the number of active Glide-Ins during a six hour run on a production infrastructure. Each Glide-In has 32 cores; thus, the number of cores for each replica is determined by the numbers of active Glide-In jobs. Figure 5 shows how the number of active Glide-Ins changed during the six hour interval. To start with, there were only enough processors to support one Glide-In job, but after approximately two hours, there were enough to activate two Glide-Ins; after three hours three Glide-Ins are running. Eight Glide-Ins were activated before the six hour time limit. As the number of Glide-Ins increases, the number of processors assigned to each replica increases, with the physically important consequence that there is a concomitant decrease in the average time between exchange attempts (upper line in Figure 5). The average time between exchange attempts as a function of the number of active Glide-Ins and the speedup – measured as the inverse of the average time normalised by the time taken with one Glide-In (thus the value of 1 for the one Glide-In case) – are shown in Figure 6.

7. Conclusion and Future Work

In summary, SAGA provides a well-defined and sufficiently powerful interface to develop the required abstraction to support adaptive distributed RE applications. SAGA allows the simple decoupling of the RE orchestration logic from the underlying distributed infrastructure. The SAGA Glide-In framework represents the first known instance of creating a runtime, system-level abstraction for distributed systems from basic programming interfaces. All this whilst remaining general purpose and extensible. The adaptive RE framework has been successfully deployed on TG and LONI resources. Using the BigJob abstraction, RE sub-jobs can efficiently be dispatched reducing the time-to-completion up to 70 %. The use of different adaptivity strategies to dynamically utilise additional resources led to a further
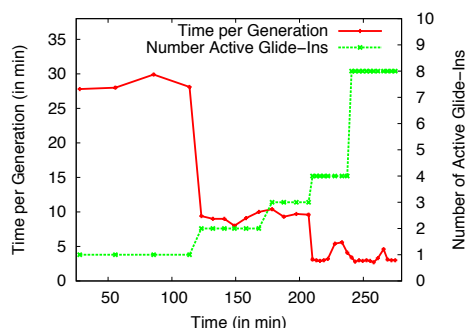
Figure 5: **The plots show the time-series of the average times between exchange attempts (upper line using the left-hand y axis) and the number of active Glide-Ins over a six-hour run on the TeraGrid.**
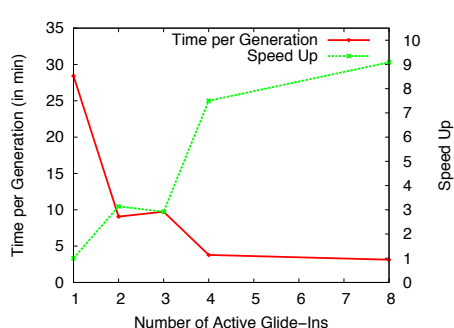
Figure 6: **The upper plot (using the left-hand y axis) illustrates how the average time between exchange attempts decreases as the number of Glide-Ins increases. At the same time the speedup grows.**

reduction of the time-to-completion. Importantly, we have shown how our REMD framework was used to enhance the sampling and rate of convergence for a biological system, the inner loop of the HCV IRES IIIb CA variant, which is an important drug-delivery target.

In the future, we will refine our RE framework making it more adaptive towards dynamic environments, e. g. by deploying an asynchronous RE scheme as described by Gallicchio et al. (2007). At the same time, we will improve our RE infrastructure to support further adaptive strategies for resource determination and utilisation. While it has been shown that resources can efficiently be allocated with the BigJob abstraction, a mechanism for dynamic resource discovery and for intelligent placements of jobs will be beneficial to further decrease the time-to-completion. Various approaches for the resources determination have been proposed, e. g. batch queue prediction (Nurmi et al. 2007, Chakraborty et al. 2008) and advance reservation-based schemes (Jeske et al. 2007).

References

Brown, S. and Head-Gordon, T. 2003, Cool Walking: A New Markov Chain Monte Carlo Sampling Method, *Journal of Computational Chemistry* **24**(1, 68-76).

Casanova, H., Obertelli, G., Berman, F. and Wolski, R. 2000, The AppLeS Parameter Sweep Template: User-level middleware for the Grid, *Scientific Programming* **8**(3), 111–126.

Chakraborty, P., Jha, S. and Katz, D. 2008, Novel Submission Modes for Tightly-Coupled Jobs Across Distributed Resources for Reduced Time to Solution, *UK e-Science All Hands Meeting*.

Collier, A., Gallego, J., Klinck, R., Cole, P., Harris, S., Harrison, G., Aboul-ela, F., Varani, G. and Walker, S. 2002, A Conserved RNA Structure Within the HCV IRES elF3-Binding Site, *Nature Structural Biology* **9**(5), 375–380.

Fol 2008, Folding at Home. http://folding.stanford.edu/.

Frey, J., Tannenbaum, T., Livny, M., Foster, I. and Tuecke, S. 2002, Condor-G: A Computation Management Agent for Multi-Institutional Grids, *Cluster Computing* **5**(3), 237–246.

Gallicchio, E., Levy, R. and Parashar, M. 2007, Asynchronous Replica Exchange for Molecular Simulations, *Journal of Computational Chemistry* **29**(5), 788–794.

Hansmann, U. 1997, Parallel Tempering Algorithm for Conformational Studies of Biological Molecules, *Chemical Physics Letters* **281**, 140–150.

Jeske, J., Luckow, A. and Schnor, B. 2007, Reservation-based Resource-Brokering for Grid Computing, *Proceedings of German e-Science Conference*, Baden-Baden, Germany.

Lei, H. X. and Duan, Y. 2007, Improved Sampling Methods for Molecular Simulation, *Current Opinion in Structural Biology* **17**(2), 187–191.

Luckow, A., Jha, S., Kim, J., Merzky, A. and Schnor, B. 2008, Distributed Replica-Exchange Simulations on Production Environments using SAGA and Migol, *Proceedings of 4th IEEE International Conference on e-Science*.

Manos, S., Mazzeo, M., Kenway, O., Coveney, P. V., Karonis, N. T. and Toonen, B. 2008, Distributed MPI Cross-Site Run Performance Using MPIg, *HPDC '08: Proceedings of the 17th international symposium on High performance distributed computing*, ACM, New York, NY, USA, pp. 229–230.

Nurmi, D., Brevik, J. and Wolski, R. 2007, QBETS: Queue Bounds Estimation From Time Series, *SIGMETRICS Perform. Eval. Rev.* **35**(1), 379–380.

Phillips, J. et al. 2005, Scalable Molecular Dynamics with NAMD, *Journal of Computational Chemistry* **26**, 1781–1802.

Shirts, M. and Pande, S. 2001, Mathematical Analysis of Coupled Parallel Simulations, *Phys. Rev. Lett.* **86**(22), 4983–4987.

Sugita, Y. and Okamoto, Y. 1999, Replica-Exchange Molecular Dynamics Method for Protein Folding, *Chemical Physics Letters* **314**, 141–151.

Woods, C. et al. 2005, Grid Computing and Biomolecular Simulation, *Philosophical Transactions of the Royal Society A: Mathematical, Physical and Engineering Sciences* **363**(1833), 2017–2035.