

# SAGA

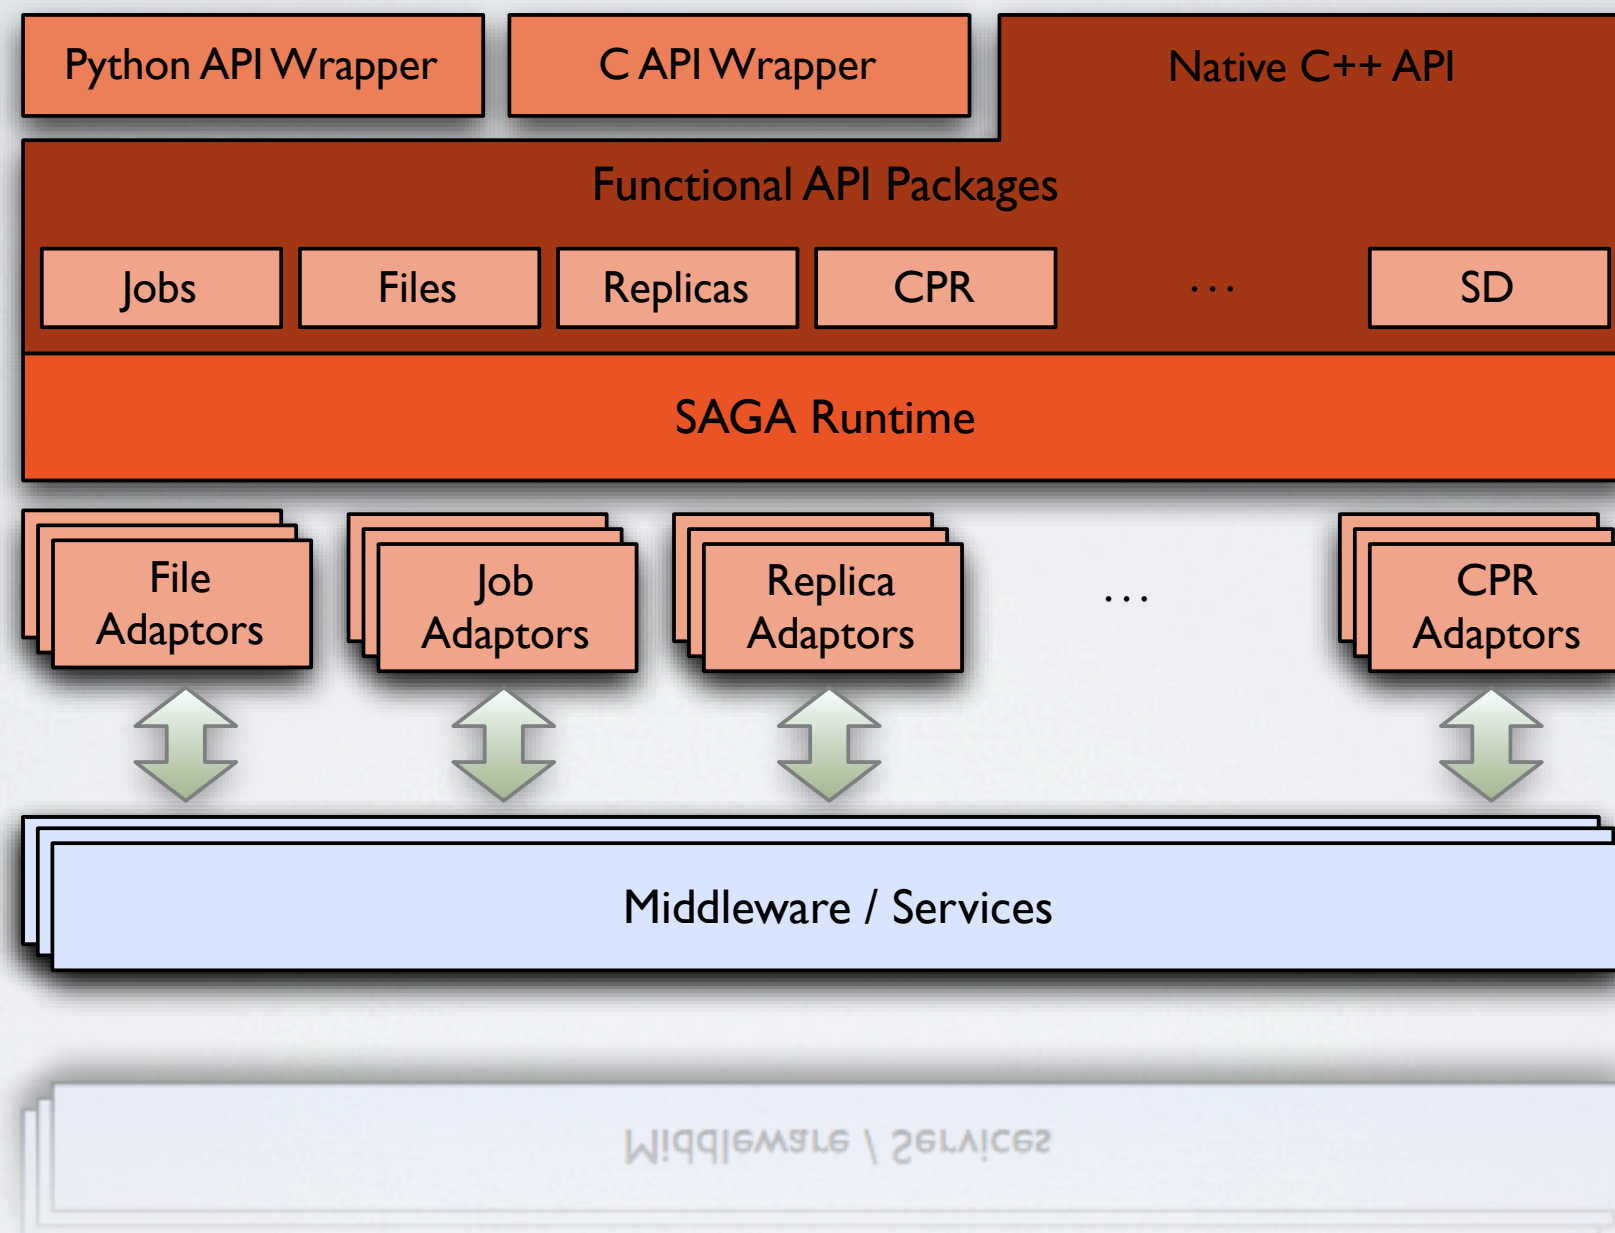
Introduction to the API and Hands-On Tutorial

<http://faust.cct.lsu.edu/trac/saga/wiki/ADSSS2009>

- Introduction to the API
- Requirements and Installation
- Tutorial Infrastructure
- Command Line Utilities
- Code Examples (C++)
  - Hello (Distributed) World
  - Chaining Jobs
  - Depending Jobs
- Python Language Bindings



- Introduction to the API
- Requirements and Installation
- Tutorial Infrastructure
- Command Line Utilities
- Code Examples (C++)
  - Hello (Distributed) World
  - Chaining Jobs
  - Depending Jobs
- Python Language Bindings





- Allows definition, submission, management and monitoring of interactive and batch jobs
- Consists of three main classes
  - `saga::job::description` - used to describe a saga job
  - `saga::job::service` - represents a (remote) computing resource
  - `saga::job::job` - represents the job itself
- Currently the following job package adaptors are available:
  - Fork (local), Globus GRAM2, SSH, OMII GridSAM, Condor, Amazon EC2, Platform LSF

```
try {  
  
    saga::url js_url ("gram://gatekeeper.lonestar.tacc.teragrid.org:2119/jobmanager-lsf");  
  
    saga::job::description jd;  
    jd.set_attribute (attributes::description_executable, "/home/oweidner/tests/heat_transfer");  
    jd.set_attribute (attributes::description_number_of_processes, "2");  
    jd.set_attribute (attributes::description_queue, "checkpoint");  
  
    saga::job::service js (js_url);  
    saga::job::job my_job = js.create_job (jd);  
  
    my_job.run();  
  
    std::cout << "Job ID      : " << my_job.get_job_id() << std::endl;  
    std::cout << "Job STATE : " << my_job.get_state() << std::endl;  
  
    my_job.suspend();  
    my_job.resume();  
  
    my_job.cancel();  
}  
  
catch (saga::exception const & e) {  
  
    std::cerr << "Ooops: " << e.what() << std::endl;  
}
```

- Can be used to traverse, modify, read and write local and remote filesystems
- Consists of two main classes:
  - `saga::filesystem::directory`
  - `saga::filesystem::file`
- Currently the following job package adaptors are available:
  - Local FS, Globus GridFTP, SSH, Hadoop Distributed Filesystem (HDFS), CloudStore KFS, OpenCloud Sector-Sphere



```
try {  
  
    saga::url file_url ("gsiftp://queenbee.loni-lsu.teragrid.org:2811//home/oweidner/.bashrc");  
    saga::filesystem::file f (file_url, saga::filesystem::Read);  
  
    while ( true )  
    {  
        saga::size_t const n = 1024*64;  
        saga::uint8_t data_buf[n+1];  
        for ( unsigned int i = 0; i <= n; ++i ) { data_buf[i] = '\\0'; }  
  
        if ( f.read (saga::buffer (data_buf, n), n) )  
        {  
            // output buffer content  
            std::cout << data_buf << std::flush;  
        }  
        else  
        {  
            break;  
        }  
    }  
}  
catch (saga::exception const & e) {  
  
    std::cerr << "Ooops: " << e.what() << std::endl;  
}
```



- An *advert* service can be used for persistent hierarchical storage of application level information and saga objects
- Semantics and usage-mode is defined by the application:
  - e.g. result storage, synchronization of application components, ...
- API is very similar to the file package:
  - `saga::advert::directory` - represents the hierarchical structure
  - `saga::advert::entry` - represents an advert object
- Currently the following advert package adaptors are available:
  - PostgreSQL / SQLite3

```
try {  
  
    saga::url advert_url ("advert://macpro01.cct.lsu.edu//users/oweidner/project01/result");  
    saga::advert::entry e (file_url, saga::advert::ReadWrite | saga::advert::Create)  
  
    e.set_attribute("Iteration", "120");  
    e.set_attribute("Dataset", "sim_42");  
  
    // You can store just a string  
    e.store_string("123.33f");  
  
    // or a saga object, e.g. a file  
    saga::url file_url ("gsiftp://queenbee.loni-lsu.teragrid.org:2811//home/oweidner/.bashrc");  
    saga::filesystem::file f (file_url, saga::filesystem::Read);  
    e.store_object(f);  
}  
catch (saga::exception const & e) {  
  
    std::cerr << "Ooops: " << e.what() << std::endl;  
}
```



- `saga::replica` - Replica management
  - Adaptors: Globus RLS, PostgreSQL / SQLite3
- `saga::stream` - Stream client and server
  - Adaptors: BSD Sockets
- `saga::sd` - Service discovery
  - Adaptors: default SD
- `saga::cpr` - Checkpoint and recover
  - Adaptors: default CPR / MiGOL



- Introduction to the API
- Requirements and Installation
- Tutorial Infrastructure
- Command Line Utilities
- Code Examples (C++)
  - Hello (Distributed) World
  - Chaining Jobs
  - Depending Jobs
- Python Language Bindings

- Open Source - released under the Boost Software License 1.0
- Implemented as a set of libraries
  - SAGA Core - A light-weight engine / runtime that dispatches calls from the API to the appropriate middle-ware adaptors
  - SAGA functional packages - Groups of API calls for: jobs, files, service discovery, advert services, RPC, replicas, CPR, ... (extensible)
  - SAGA language wrappers - Thin Python and C layers on top of the native C++ API
  - SAGA middle-ware adaptors - Take care of the API call execution on the middle-ware
- Can be configured / packaged to suit your individual needs!



- In order to build & install SAGA you need the following:
  - A UNIX operating system (Linux, MacOS, etc.)
  - A C++ Compiler (preferably gcc  $\geq$  3.4)
  - The Boost C++ Libraries (  $\geq$  1.33.1) from <http://boost.org>
  - Python (if you want to build the Python language bindings)
- Adaptors may have additional requirements
  - PostgreSQL / SQLite client libraries
  - Globus / Condor / LSF installations
  - etc ...



- Download the latest source release from:

<http://saga.cct.lsu.edu/cpp/download>

- Checkout the latest source from Subversion:

`svn co https://svn.cct.lsu.edu/repos/saga/trunk`

- Configure/make - based build system. It's as simple as:

```
./configure --prefix=/usr/local && make install
```

- Top-level configure/make recursively calls configures/makes for each adaptor, language bindings, etc...
- Fault tolerant: if one of the sub-level packages can't be configured (missing prerequisites, etc...) it is simply skipped
- Sub-level packages can be built individually and even outside the source tree
- XCode and VisualStudio project files available for developers



# SAGA

A Simple API for Grid Applications [<http://saga.cct.lsu.edu>]

## DEPLOYMENT 2009/2010



Ole Weidner, Shantenu Jha

ADSSS'09 Abingdon Sept. 03, 2009




- Introduction to the API
- Requirements and Installation
- **Tutorial Infrastructure**
- Command Line Utilities
- Code Examples (C++)
  - Hello (Distributed) World
  - Chaining Jobs
  - Depending Jobs
- Python Language Bindings

# SAGA

A Simple API for Grid Applications [<http://saga.cct.lsu.edu>]

## INFRASTRUCTURE





- Login (ssh) to one of the following machines:
  - tc09.nesc.ed.ac.uk
  - tc11.nesc.ed.ac.uk
  - tc15.nesc.ed.ac.uk
  - tc16.nesc.ed.ac.uk
  - tc17.nesc.ed.ac.uk
- Usernames:
  - course01, course02, course03, ..., course24
- Password: to be announced



- Set up the SAGA environment

`source /usr/local/saga/share/saga/saga-env.sh`

```
#!/bin/bash

export SAGA_LOCATION=/usr/local/saga/
export LD_LIBRARY_PATH=${SAGA_LOCATION}/lib/:${LD_LIBRARY_PATH}
export PATH=${SAGA_LOCATION}/bin:${PATH}

export PYTHONPATH=${SAGA_LOCATION}/lib/python2.6/site-packages/:${PYTHONPATH}
```

- Introduction to the API
- Requirements and Installation
- Tutorial Infrastructure
- **Command Line Utilities**
- Code Examples (C++)
  - Hello (Distributed) World
  - Chaining Jobs
  - Depending Jobs
- Python Language Bindings



- Provides basic functionality of the file package
- Examples:
  - List the contents of a directory  
`saga-file list_dir file://localhost/tmp/`
  - Get the size of a file  
`saga-file get_size file://localhost/etc/passwd`
  - Copy a file  
`saga-file copy /etc/passwd ssh://tc11/tmp/courseXX_etcpasswd_copy`
  - Print the contents of a file  
`saga-file cat /etc/passwd ssh://tc11/tmp/courseXX_etcpasswd_copy`

- Provides basic functionality of the job package
- Examples
  - Submit a non-interactive job  
`saga-job submit ssh://localhost /bin/touch /tmp/blah`
  - Run an interactive job  
`saga-job run ssh://tc09 /bin/cat /proc/cpuinfo`



- Provides basic functionality of the advert package
- Examples:
  - List the content of an advert directory  
`saga-advert list_directory advert://macpro01.cct.lsu.edu//ADSSS09 ?`
  - Create an advert entry  
`saga-advert add_entry advert://macpro01.cct.lsu.edu//ADSSS09/aloha`
  - Attach an attribute  
`saga-advert set_attribute advert://macpro01.cct.lsu.edu//ADSSS09/aloha  
Foo Bar`
  - List all attributes  
`saga-advert list_attributes advert://macpro01.cct.lsu.edu//ADSSS09/aloha`

- Try and run command line tools
  - Copy a file, move it, delete it, read its contents (local / remote)
  - Run a job (/bin/sleep 20), monitor its status (local / remote)
  - Use the advert service to create directories, entries, store data, set attributes

### NOTES:

- Please stay within <advert://macpro01.cct.lsu.edu//ADSSS09/>
- Create a sub-directory for your username, e.g. [course24](#)



- Introduction to the API
- Requirements and Installation
- Tutorial Infrastructure
- Command Line Utilities
- **Code Examples (C++)**
  - Hello (Distributed) World
  - Chaining Jobs
  - Depending Jobs
- Python Language Bindings

- We will go over three different examples
- For each example:
  - Create a subdirectory, e.g. `~/ex01` `~/ex02` `~/ex03`
  - Download the Makefile from the tutorial wiki
  - Download the source files from the tutorial wiki
  - Compile
  - Run



- Introduction to the API
- Requirements and Installation
- Tutorial Infrastructure
- Command Line Utilities
- Code Examples (C++)
  - Hello (Distributed) World
  - Chaining Jobs
  - Depending Jobs
- Python Language Bindings

- Spawn 3 (remote) jobs (/bin/echo) with the 3 words "Hello", "distributed", and "world!" as their arguments.
- You can change these three lines:

```
#define HOST1 "fork://localhost"  
#define HOST2 "fork://localhost"  
#define HOST3 "fork://localhost"
```

- What do you observe when you run it multiple times?



- Introduction to the API
- Requirements and Installation
- Tutorial Infrastructure
- Command Line Utilities
- Code Examples (C++)
  - Hello (Distributed) World
  - Chaining Jobs
  - Depending Jobs
- Python Language Bindings

- Introduces dependencies (ordered execution) between the jobs
- Each job receives the output of the previous job (integer number), increments it by one and passes it to the next job
- Again, change these three lines:

```
#define HOST1 "fork://localhost"  
#define HOST2 "fork://localhost"  
#define HOST3 "fork://localhost"
```

- Can you come up with other, real-life uses-cases (maybe even from your own field of work/research) for this?



- Introduction to the API
- Requirements and Installation
- Tutorial Infrastructure
- Command Line Utilities
- Code Examples (C++)
  - Hello (Distributed) World
  - Chaining Jobs
  - Depending Jobs
- Python Language Bindings





- Introduction to the API
- Requirements and Installation
- Tutorial Infrastructure
- Command Line Utilities
- Code Examples (C++)
  - Hello (Distributed) World
  - Chaining Jobs
  - Depending Jobs
- Python Language Bindings

# SAGA

A Simple API for Grid Applications [<http://saga.cct.lsu.edu>]

---

# THANKS

Questions / Comments ?