



CENTER FOR COMPUTATION  
& TECHNOLOGY

# SAGA

A Simple API for Grid Applications

## Introduction to the SAGA API



omii-uk  
[www.omii.ac.uk](http://www.omii.ac.uk)



## Outline

- ▣ API structure and scope
- ▣ API walkthrough
- ▣ implementation details
- ▣ API extensions

## SAGA Design Principles

- ▣ SAGA: Simple API for Grid Applications
  - ▣ OGF approach to a uniform API layer (facade)
- ▣ governing principle: 80:20 rule
  - ▣ simplicity versus control!
- ▣ top-down approach
  - ▣ use case driven!
  - ▣ defines application level abstractions
- ▣ extensible
  - ▣ stable look & feel
  - ▣ API packages
- ▣ API Specification is language independent (IDL)
  - ▣ Renderings exist in C++, Python, Java
  - ▣ Examples here are in C++

## SAGA Intro: Example 1

```
// SAGA: File Management example

saga::filesystem::directory dir ("any://remote.host.net//data/");

if ( dir.exists ("a") && ! dir.is_dir ("a") )
{
    dir.copy ("a", "b", Overwrite);
}

list <saga::url> names = dir.find ("*-{123}.txt");

saga::filesystem::directory tmp  = dir.open_dir ("tmp/", Create);
saga::filesystem::file      file = dir.open      ("tmp/data.txt");
```

## SAGA Intro: Example 1

- ▣ API is clearly POSIX (libc + shell) inspired
- ▣ where is my security??
- ▣ what is 'any:/' ???
- ▣ usage should be intuitive (hopefully)

## SAGA Intro: Example 2

```
// SAGA: Job Submission example

saga::job::description jd;
// details left out

saga::job::service js ("any://remote.host.net/");
saga::job::job      j = js.create_job (jd);

j.run ();

cout << "Job State: " << j.get_state () << endl;

j.wait ();

cout << "RetVal " << j.get_attribute ("ExitCode") << endl;
```

## SAGA Intro: Example 2'

```
// SAGA: Job Submission example

saga::job::service js ("any://remote.host.net");
saga::job::job      j = js.run_job ("touch /tmp/touch.me");

cout << "Job State: " << j.get_state () << endl;

j.wait ();

cout << "RetVal " << j.get_attribute ("ExitCode") << endl;
```

## SAGA Intro: Example 2

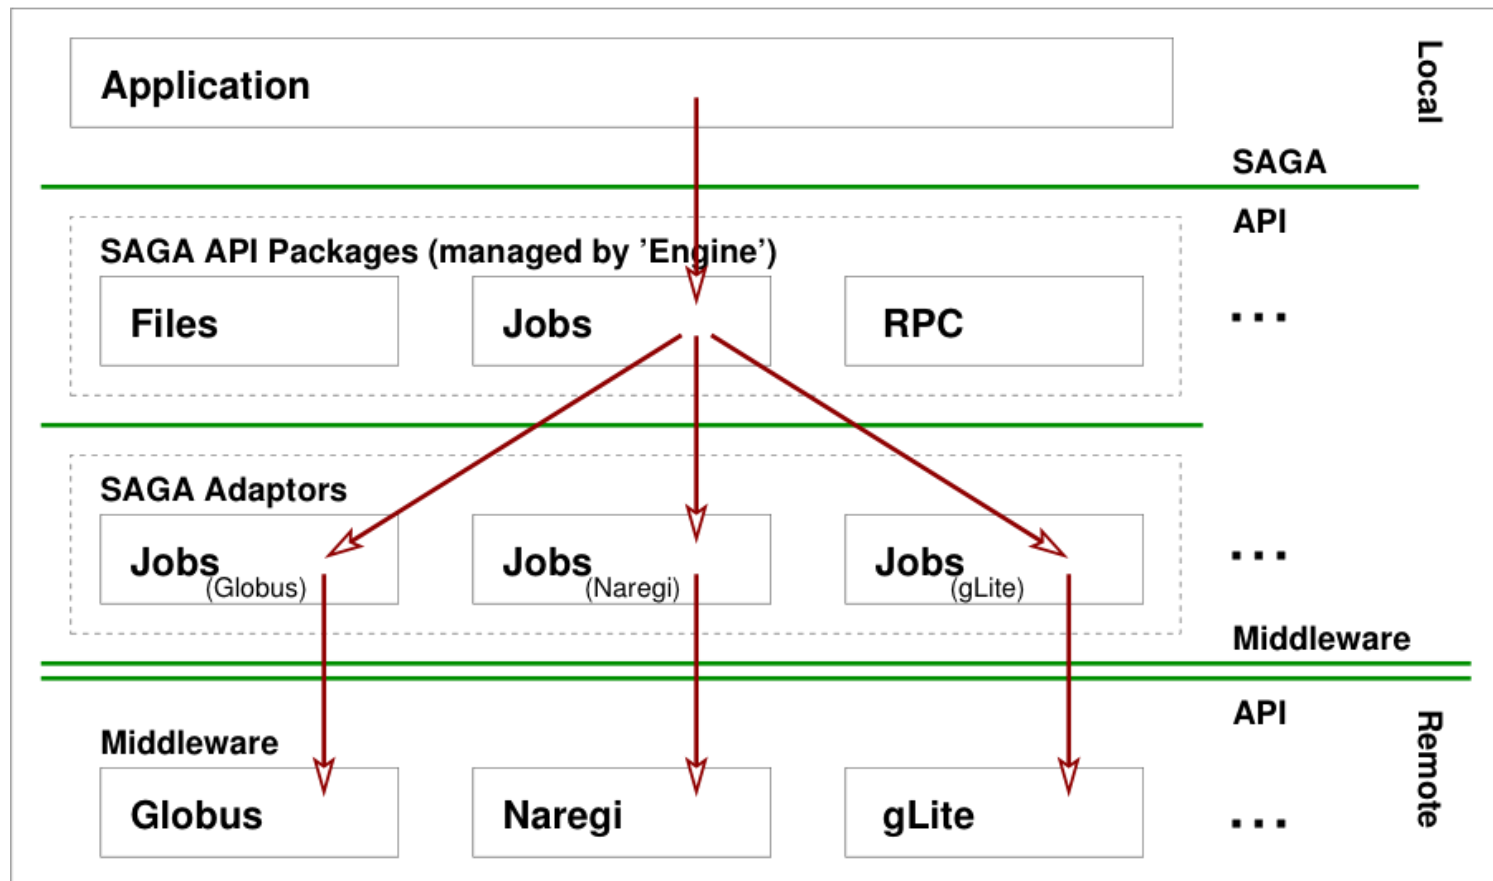
- ▣ stateful objects!
- ▣ yet another job description language? :-(
- ▣ many hidden/default parameters
  - ▣ keeps call signatures small
- ▣ 'any:/' again!
- ▣ TIMTOWTDI (there is more than one way to do it)



## SAGA Intro: 10.000 feet

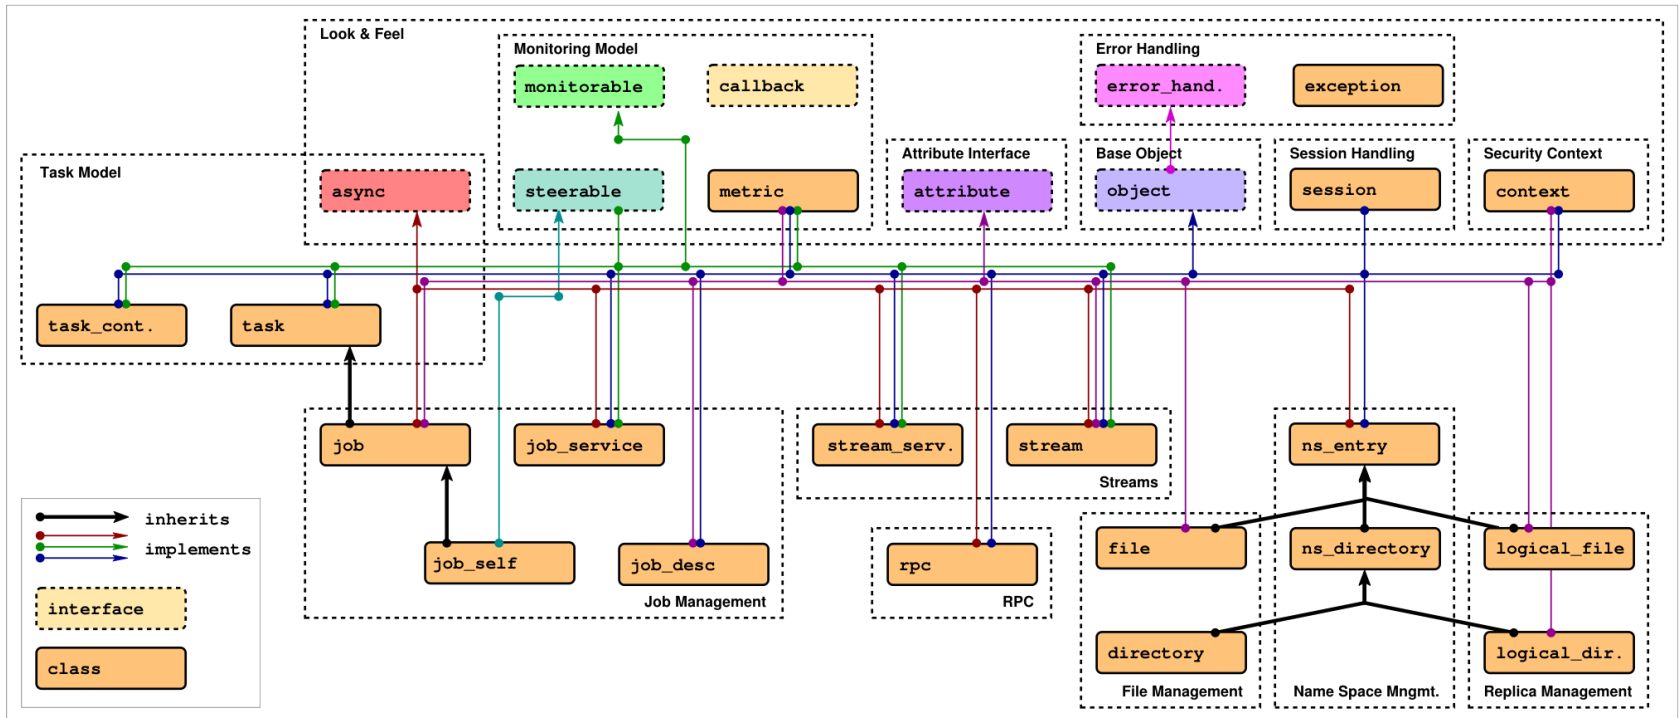
- ▣ object oriented:
  - ▣ uses inheritance and interfaces
  - ▣ very moderate use of templates though!
- ▣ functional and non-functional elements strictly separated
  - ▣ functional API:
    - ▣ typically mappable to remote operations
    - ▣ ordered in API 'Packages': extensible
  - ▣ non-functional API:
    - ▣ typically not mappable to explicit remote operations
    - ▣ "look & feel": orthogonal to functional API
    - ▣ security, asynchronous ops, notifications, ...
- ▣ few inter-package dependencies - allows for partial implementations

## Implementation



# SAGA Class Hierarchy

## SAGA: Class hierarchy



# SAGA: Class hierarchy

Functional API Packages


## SAGA Job Package: Overview

- ▣ running jobs is **use case #1**
- ▣ all middlewares support it, one way or the other
- ▣ well stablished patterns exist
  - ▣ job description
  - ▣ job state
  - ▣ submission endpoint
  - ▣ ...

## SAGA Job Package: Example 1

```
saga::job::description jd;  
saga::job::service      js ("gram://remote.host.net");  
saga::job               j = js.create_job (jd);  
  
j.run ();  
  
cout << "Job State: " << j.get_state () << endl;  
  
j.wait ();  
  
cout << "Retval " << j.get_attribute ("ExitCode") << endl;
```

## SAGA Job Package: job states





## SAGA Job Package: job operations

```
j.run      ();  
j.wait     ();  
j.cancel   ();  
  
j.suspend  ();  
j.resume   ();  
  
j.signal    (SIGUSR1);  
j.checkpoint ();  
j.migrate   (jd);
```

## SAGA Job Package: job description

```
saga::job::description jd;  
  
jd.set_attribute ("Executable",      "/bin/tail");  
jd.set_attribute ("WorkingDirectory", "data/");  
jd.set_attribute ("Cleanup",         "False");  
  
// pseudo code *blush*  
jd.set_vector_attribute ("Arguments",  ["-f", "my_log"]);  
jd.set_vector_attribute ("Environment", ["TMPDIR=/tmp/"]);  
jd.set_vector_attribute ("FileTransfer", ["my_log >> all_logs"]);
```

## SAGA Job Package: job description

Executable	WorkingDirectory	TotalPhysicalMemory
Arguments	<i>Interactive</i>	CPUArchitecture
Environment	Cleanup	OperatingSystemType
CandidateHosts	Input	<i>Queue</i>
SPMDVariation	Output	JobProject
TotalCPUCount	Error	<i>JobContact</i>
NumberOfProcesses	<i>JobStartTime</i>	FileTransfer
ProcessesPerHost	WallTimeLimit	
ThreadsPerProcess	TotalCPUTime	

## SAGA Job Package: job description

- ▣ leaning heavily on JSDL, but flat
- ▣ borrowing from DRMAA
- ▣ mixes hardware, software and scheduling attributes!
- ▣ cannot be extended
- ▣ no support for 'native' job descriptions (RSL, JDL, ...)
- ▣ only 'Executable' is required
- ▣ backend MAY ignore unsupported keys  
`cd /tmp/data && rm -rf *`

## SAGA Job Package: job service

```
saga::job::service js ("gram://remote.host.net/");

vector <string> ids = js.list (); // list known jobs

while ( ids.size () )
{
    string    id = ids.pop_back (); // fetch one job id
    saga::job j  = js.get_job (id); // reconnect to job

    cout << id << " : " << j.get_state () << endl;
}
```

## SAGA Job Package: job service

- ▣ represents a specific job submission endpoint
- ▣ job states are maintained on that endpoint (usually)
- ▣ full reconnect may not be possible (I/O streaming)
- ▣ lifetime of state up to backend
- ▣ reconnected jobs may have different job description (lossy translation)

## SAGA Namespace Package

- ▣ interfaces for managing entities in name spaces
- ▣ files, replicas, information, resources, steering parameter, checkpoints, . . .
- ▣ manages hierarchy (mkdir, cd, ls, . . . )
- ▣ entries are assumed to be opaque (copy, move, delete, ...)

## SAGA Namespace Package: example

```
saga::name_space::directory d ("ssh://remote.host.net//data/");

if ( d.is_entry ("a") && ! d.is_dir ("a") )
{
    d.copy ("a", "../b");
    d.link ("../b", "a", Overwrite);
}

list <saga::url> names = d.find ("*-{123}.text.");

saga::name_space::directory tmp = d.open_dir ("tmp/data/1",
                                              saga::name_space::CreateParents);

saga::name_space::entry data = tmp.open ("data.txt");

data.copy ("data.bak", Overwrite); // uses cwd
```



## SAGA Namespace Package

- ▣ name space entries are opaque: the name space package can never look inside
- ▣ directories *are* entries (inheritance)
- ▣ inspection: `get_cwd()`, `get_url()`, `get_name()`, `exists()`,  
`is_entry()`, `is_dir()`, `is_link()`, `read_link()`
- ▣ manipulation: `create()`, `copy()`, `link()`, `move()`, `remove()`
- ▣ permissions: `permissions_allow()`, `permissions_deny()`
- ▣ wildcards are supported (POSIX influence...)

## SAGA Filesystem Package

- ▣ implements name space interface
- ▣ adds access to **content** of namespace::entries (files)
- ▣ POSIX oriented: `read()`, `write()`, `seek()`
- ▣ optimizations: for distributed file access:
  - ▣ scattered I/O
  - ▣ pattern based I/O
  - ▣ extended I/O (from GridFTP)

## SAGA Filesystem Package: Example

```
saga::filesystem::file f ("any://remote.host.net/data/data.bin");

char mem[1024];
saga::mutable_buffer buf (mem);

if ( f.get_size () >= 1024 )
{
    buf.set_data (mem + 0, 512);
    f.seek (512, saga::filesystem::Start);
    f.read (buf);
}

if ( f.get_size () >= 512 )
{
    buf.set_data (mem + 512, 512);
    f.seek (0, saga::filesystem::Start);
    f.read (buf);
}
```

## SAGA Filesystem Package

- provides access to the **content** of filesystem entries (sequence of bytes)
- saga buffers are used to wrap raw memory buffers
- saga buffers can be allocated/managed by the SAGA implementation
- several incarnations of read/write: posix style, scattered, pattern based

## SAGA Filesystem Package: Flags

```
enum flags {  
    None           = 0,  
    Overwrite      = 1,  
    Recursive      = 2,  
    Dereference    = 4,  
    Create         = 8,  
    Exclusive      = 16,  
    Lock           = 32,  
    CreateParents  = 64,  
    Truncate       = 128, // not on name_space  
    Append         = 256, // not on name_space  
    Read           = 512,  
    Write          = 1024,  
    ReadWrite      = 1536, // Read | Write  
    Binary         = 2048 // only on filesystem  
}
```

## SAGA Advert Package

- ▣ persistent storage of application level information
- ▣ semantics of information defined by application
- ▣ allows storage of serialized SAGA objects (object persistency)
- ▣ ***very useful for bootstrapping and coordinating distributed application components***

## SAGA Advert Package: Example

```
// example for browsing my task adverts
saga::advert::directory todo ("any//remote.host.net/my_tasks/");

// pseudo vector code
list <saga::url> urls = todo.find ("*", ["priority=urgent"]);

while ( urls.size () )
{
    saga::advert ad (urls.pop_front ());
    std::cout << ad.get_attribute ("title")          << std::endl;
    std::cout << ad.get_attribute (« deadline")      << std::endl;
    std::cout << ad.get_attribute ("description") << std::endl << std::endl;
}
```

## SAGA Advert Package: Example

```
// master side code: advertise (publish) a saga::file instance
saga::file f (url);

saga::advert ad ("any//remote.host.net/files/my_file_ad", Create);

ad.store_object (f);
```

```
// client side code: retrieve file instance
saga::advert ad ("any//remote.host.net/files/my_file_ad");

saga::file f = ad.retrieve_object ();
```



# SAGA: Class hierarchy

Look & Feel Packages

## SAGA Session: Example – default session

```
saga::ns_dir dir ("any://remote.host.net//data/");

if ( dir.is_entry ("a") && ! dir.is_dir ("a") )
{
    dir.copy ("a", "../b");
    dir.link ("../b", "a", Overwrite);
}

list <saga::url> names = dir.find ("*-{123}.text.");

saga::name_space::directory sub    = dir.open_dir ("tmp/");
saga::name_space::entry      entry = dir.open      ("data.txt");

entry.copy ("data.bak", Overwrite);
```

## SAGA Session: Properties

- ▣ by default hidden (default session is used)
- ▣ session is identified by lifetime of security credentials, and by objects in this session (jobs etc.)
- ▣ session is used on object creation (optional)
- ▣ `saga::context` can attach security tokens to a session
  - ▣ the default session has default contexts

## SAGA Session: Example – explicit session

```
saga::context c1 (saga::context::X509);  
saga::context c2 (saga::context::X509);  
  
c2.set_attribute ("UserProxy", "/tmp/x509up_u123.special");  
  
saga::session s;  
  
s.add_context (c1);  
s.add_context (c2);  
  
saga::name_space::dir dir (s, "any://remote.host.net/data/");
```

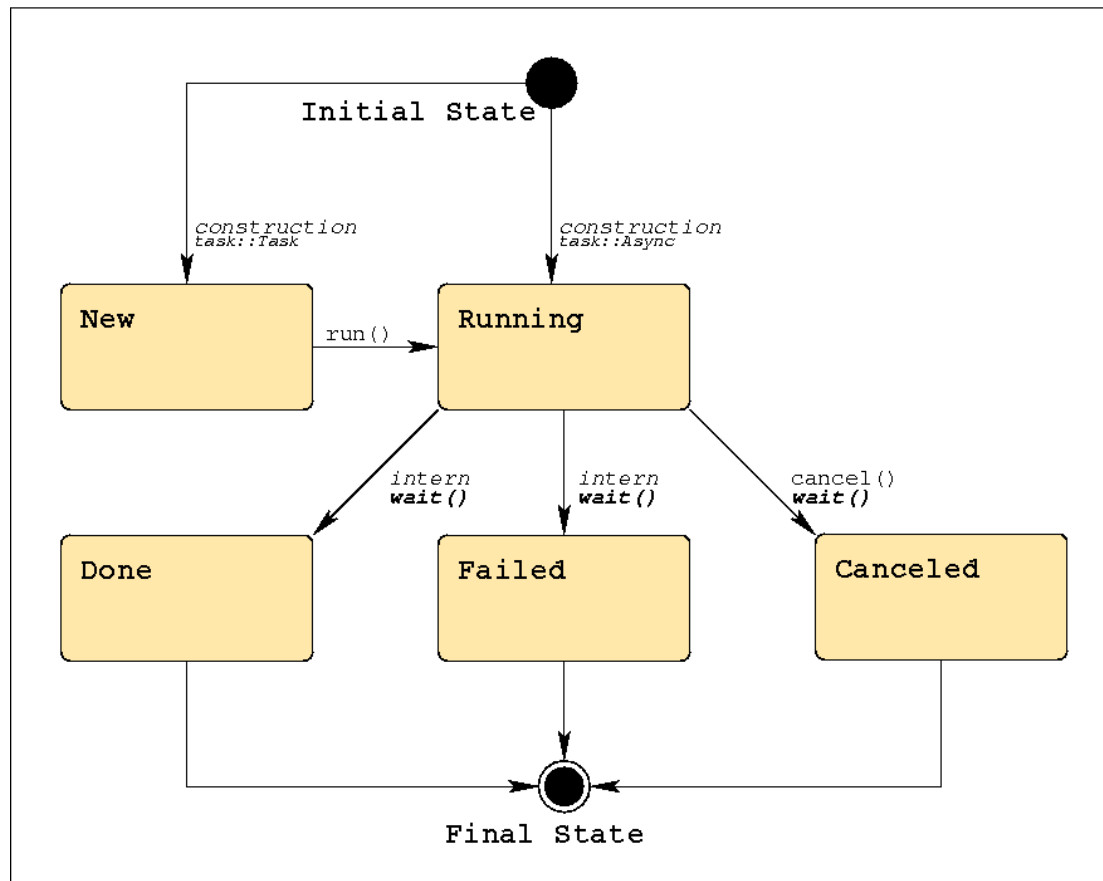
## SAGA Session: Lifetime

```
saga::dir  dir (s, "gridftp://remote.host.net/data/");  
  
saga::file file = dir.open ("data.bin");  
  
s.remove_context (c1);  
s.remove_context (c2);  
  
file.copy ("data.bin.bak");  // this works - session is sticky!
```

## SAGA Tasks

- ▣ asynchronous operations are a *MUST* in distributed systems, and thus Grids
- ▣ `saga::task` represents an asynchronous operation (e.g. `file.copy ()`)
- ▣ `saga::task_container` manages multiple tasks
- ▣ tasks are stateful (similar to jobs)

## SAGA Tasks: States



## SAGA Tasks

- ▣ different versions for each method call: Sync, Async, Task
- ▣ signature basically the same
- ▣ differ in state of the task returned by that method
  - ▣ Sync: task is Done
  - ▣ Async: task is Running -> wait();
  - ▣ Task: task is New -> run(); wait();

- ▣ delayed exception delivery

```
if ( saga::task::Failed == task.get_state () )  
{  
    task.rethrow ();  
}
```



## SAGA Task: Example

```
// normal method call, synchronous
/* void */ file.copy ("data.bak");

// async versions, never throw (use 'rethrow()' on failure)
saga::task t1 = file.copy      <saga::task::Sync>  ("data.bak.1");
saga::task t2 = file.copy      <saga::task::Async> ("data.bak.2");
saga::task t3 = file.copy      <saga::task::Task>  ("data.bak.3");

// t1: Done
// t2: Running
// t3: New

t3.run ();    // t3 now Running, too

t2.wait ();
t3.wait ();

// t1, t2, t3 are final (Done or Failed)
```

## SAGA Task: Example

```
// normal method call, synchronous
size_t s = file.get_size ();

// async versions, never throw (use 'rethrow()' on failure)
saga::task t1 = file.get_size <saga::task::Sync> ();
saga::task t2 = file.get_size <saga::task::Async> ();
saga::task t3 = file.get_size <saga::task::Task> ();

// get_result: implies wait() and rethrow(), and thus can throw!
size_t s1 = t1.get_result <size_t> ();
size_t s2 = t2.get_result <size_t> ();
size_t s3 = t3.get_result <size_t> ();
```

## SAGA Task Container: Example

```
// create task container
saga::task_container tc;

// add tasks
tc.add (t1);
tc.add (t2);
tc.add (t3);

// collective operations on all tasks in container
tc.run ();

saga::task done_task = tc.wait (saga::task::Any);

tc.wait (saga::task::All);
```

## SAGA Task Container: Tasks and Jobs

```
// NOTE:  
// class saga::job : public saga::task  
// task container can thus manage tasks *and* jobs:  
  
saga::task task = file.copy <saga::task::Async> ("b");  
saga::job job = js.run_job ("remote.host.net", "/bin/date");  
  
saga::task_container tc;  
  
tc.add (task);  
tc.add (job);  
  
tc.wait (saga::task::All);
```