

SAGA – a Simple API for Grid Applications

Introduction to the API

Agenda

- SAGA API structure and scope
- walkthrough
- SAGA extensions

- coding applications for Grids is notoriously difficult:
 - complex infrastructure,
 - diversity of middlewares
 - environment evolves very quickly
- diversity of Grid Middleware implies diversity of APIs
- difficult to keep up with MW development, and to stay **simple**
- few APIs try to generalize Grid programming concepts

- some API standards exist in OGF, and are successful
- OGF APIs do **not** cover the complete OGF scope
- the various API standards are disjunct
- WSDL as service interface specification cannot replace an application level API (wrong level of abstraction)
- **SAGA tries to address these issues**

OGF: top-down vs. bottom-up

- bottom-up often agrees on (semantic) LCD + backend specific extensions
- top-down usually focuses on semantics of application requirements
- bottom-up APIs tend to be more powerful
- top-down APIs tend to be simpler and more concise
- **we very much prefer top-down!**

SAGA

Simple API for Grid Applications

SAGA Design Principles

- **SAGA: Simple API for Grid Applications**
- OGF approach to a uniform API layer (facade)
- **governing principle: 80:20 rule**
simplicity versus control!
- **top-down approach:** use case driven!
→ defines **application level** abstractions
- **extensible:** stable look & feel + API packages
- **influenced by:** DRMAA, GridRPC, OREP, JSDL, POSIX, GAT, CoG, LSF, Globus, ...
- API Specification is **Language Independent (IDL)**
Renderings exist in C++, Python, Java
Examples here are in C++

SAGA Intro: Example 1

SAGA: File Management

```
saga::filesystem::directory dir ("any://remote.host.net//data/");  
  
if ( dir.exists ("a") && ! dir.is_dir ("a") )  
{  
    dir.copy ("a", "b", Overwrite);  
}  
  
list <saga::url> names = dir.find ("*-{123}.txt");  
  
saga::filesystem::directory tmp  = dir.open_dir ("tmp/", Create);  
saga::filesystem::file      file = dir.open      ("tmp/data.txt");
```


SAGA Intro: Example 1

- API is clearly POSIX (libc + shell) inspired
- where is my security??
- what is 'any: / /' ???

SAGA Intro: Example 2

SAGA: Job Submission

```
saga::job::description jd; // details left out
saga::job::service      js ("any://remote.host.net/");
saga::job::job           j = js.create_job (jd);

j.run ();

cout << "Job State: " << j.get_state () << endl;

j.wait ();

cout << "Retval " << j.get_attribute ("ExitCode") << endl;
```

SAGA Intro: Example 2'

SAGA: Job Submission

```
saga::job::service      js ("any://remote.host.net");  
saga::job::job          j = js.run_job ("touch /tmp/touch.me");  
  
cout << "Job State: " << j.get_state () << endl;  
  
j.wait ();  
  
cout << "Retval " << j.get_attribute ("ExitCode") << endl;
```

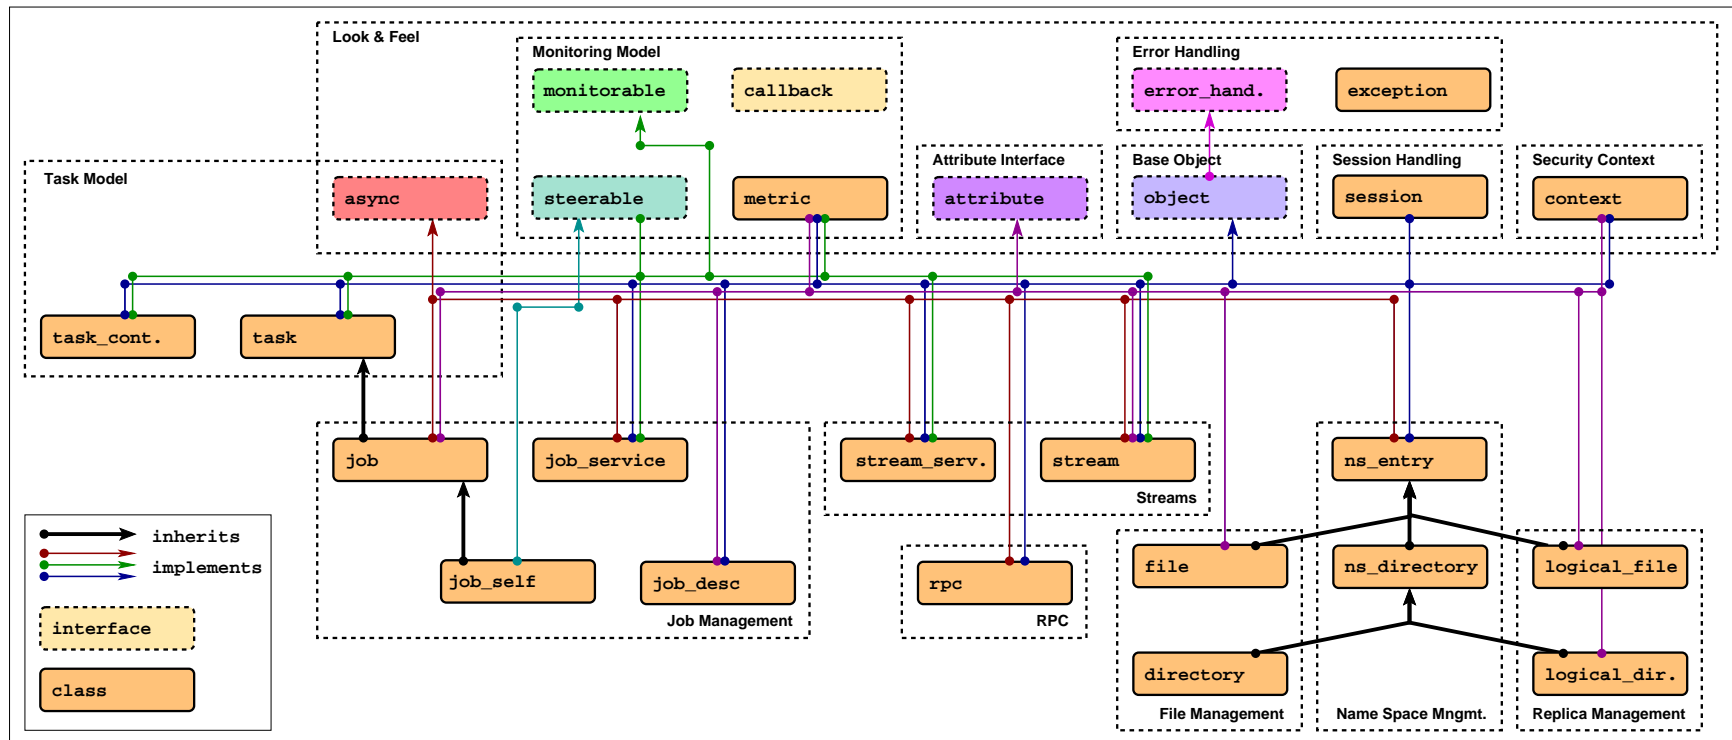
SAGA Intro: Example 2

- stateful objects!
- yet another job description language? : – (
- many hidden/default parameters (to keep call signatures small)
- 'any: / /' again!
- TIMTOWTDI (there is more than one way to do it)

SAGA Intro: 10.000 feet

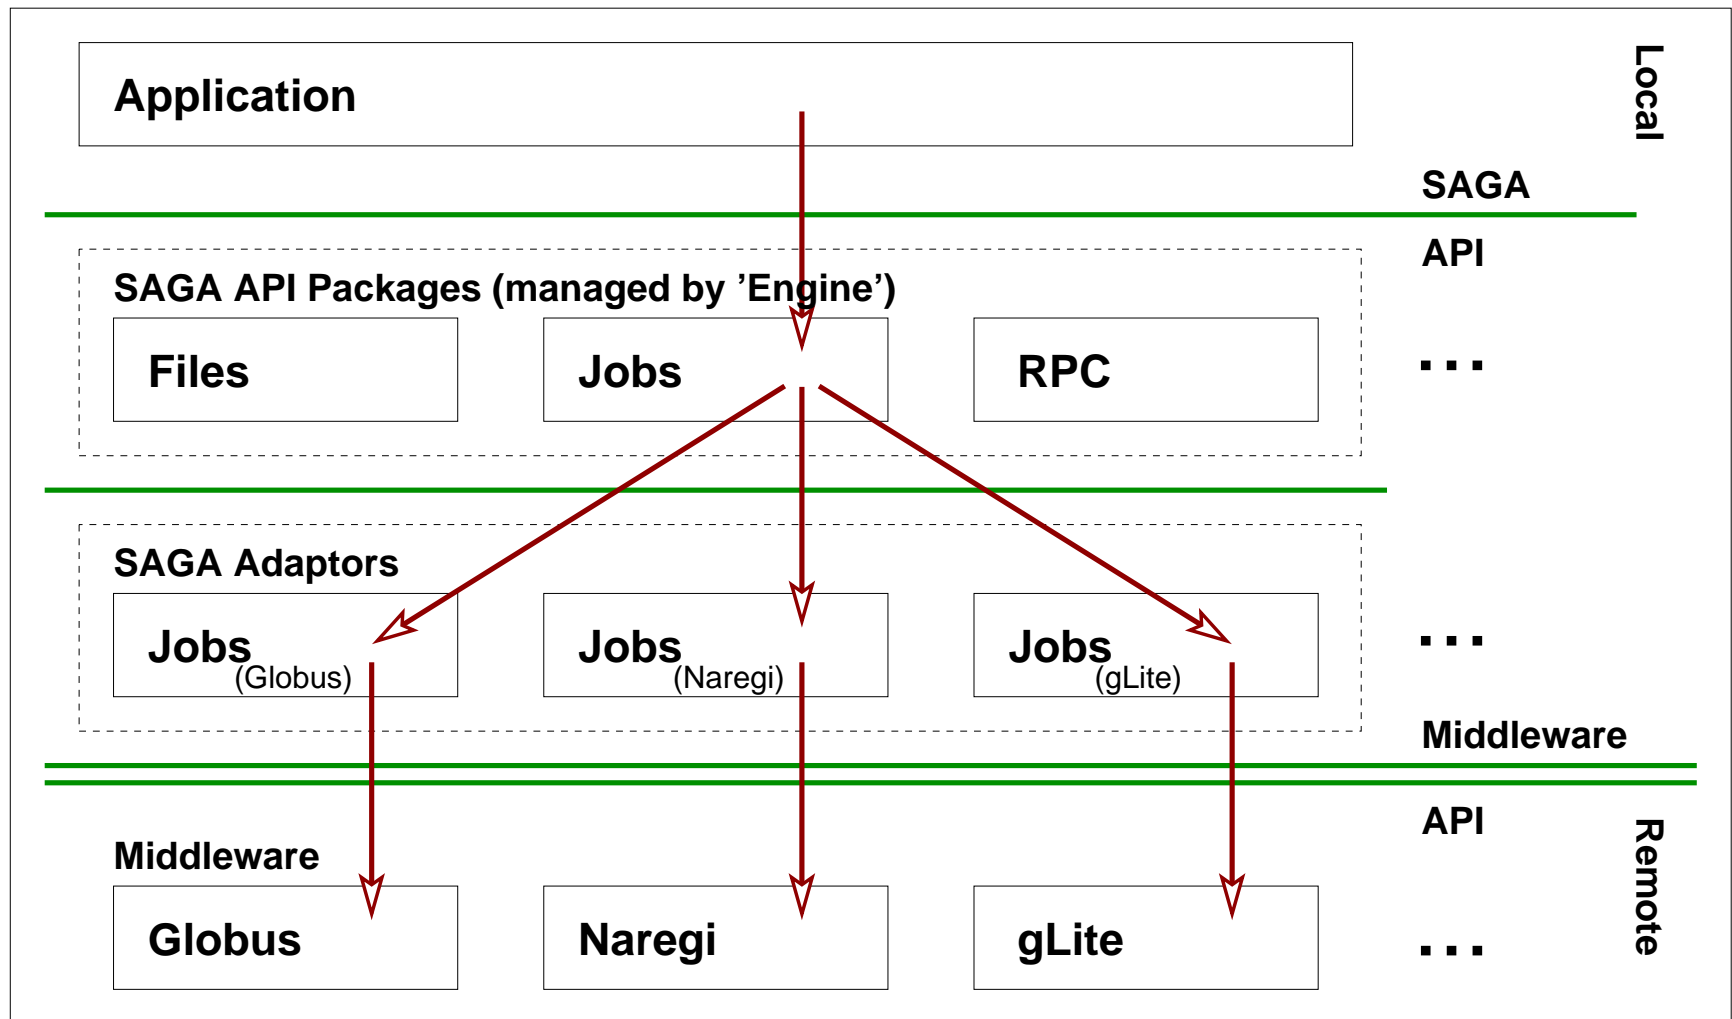
- **object oriented:** inheritance, interfaces
very moderate use of templates though!
- functional and non-functional elements strictly separated
 - *non-functional API:* look & feel- orthogonal to functional API
often not mappable to remote operations
 - *functional API:* API 'Packages' - extensible
typically mappable to remote operations
- few inter-package dependencies - allows for partial implementations

SAGA: Class hierarchy

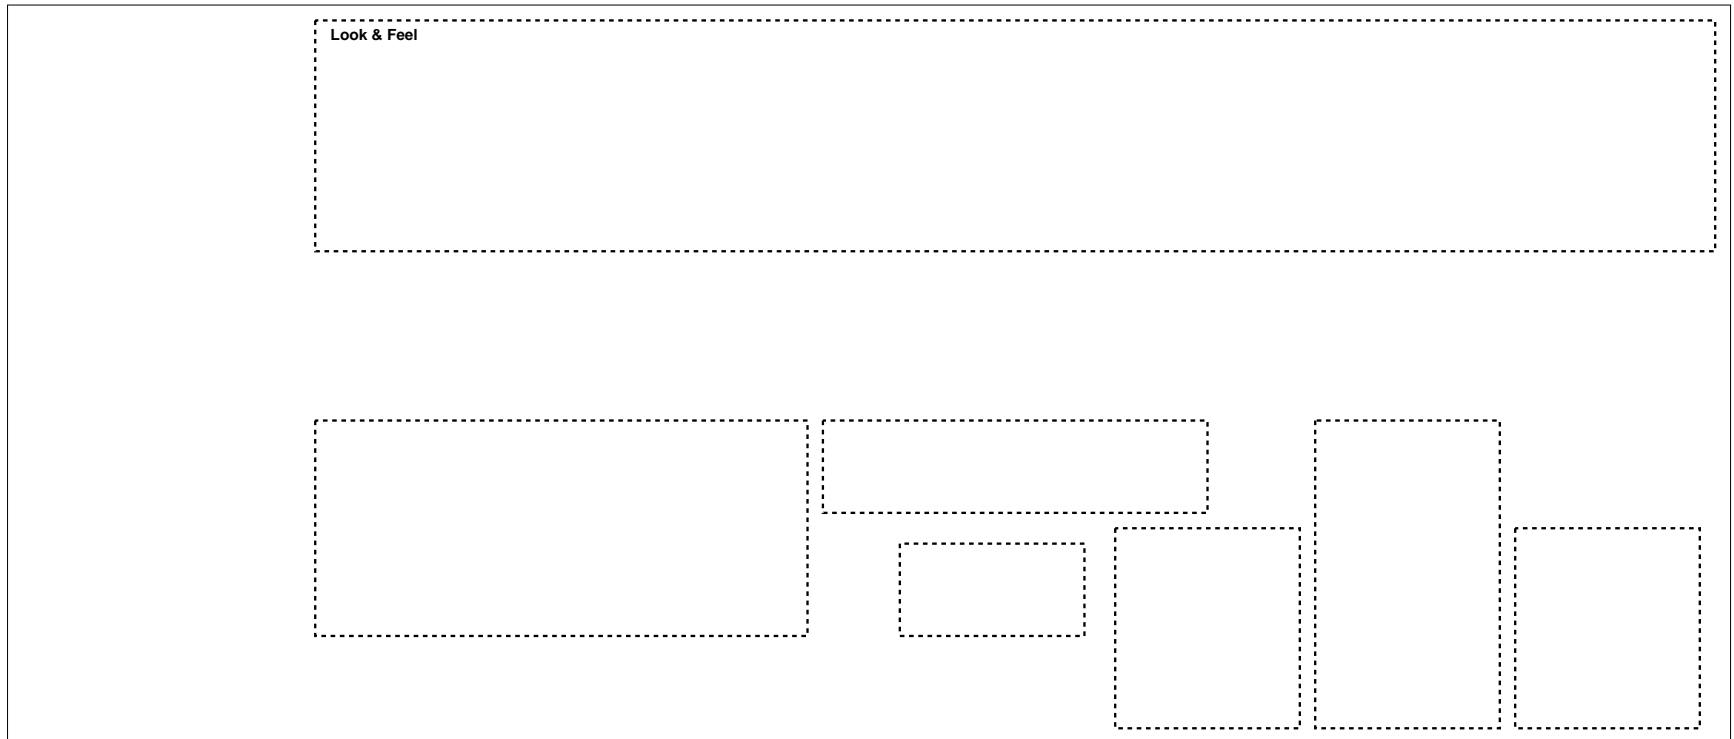


SAGA API Structure: look & feel (top) + API packages (bottom)

Implementation



SAGA: Class hierarchy



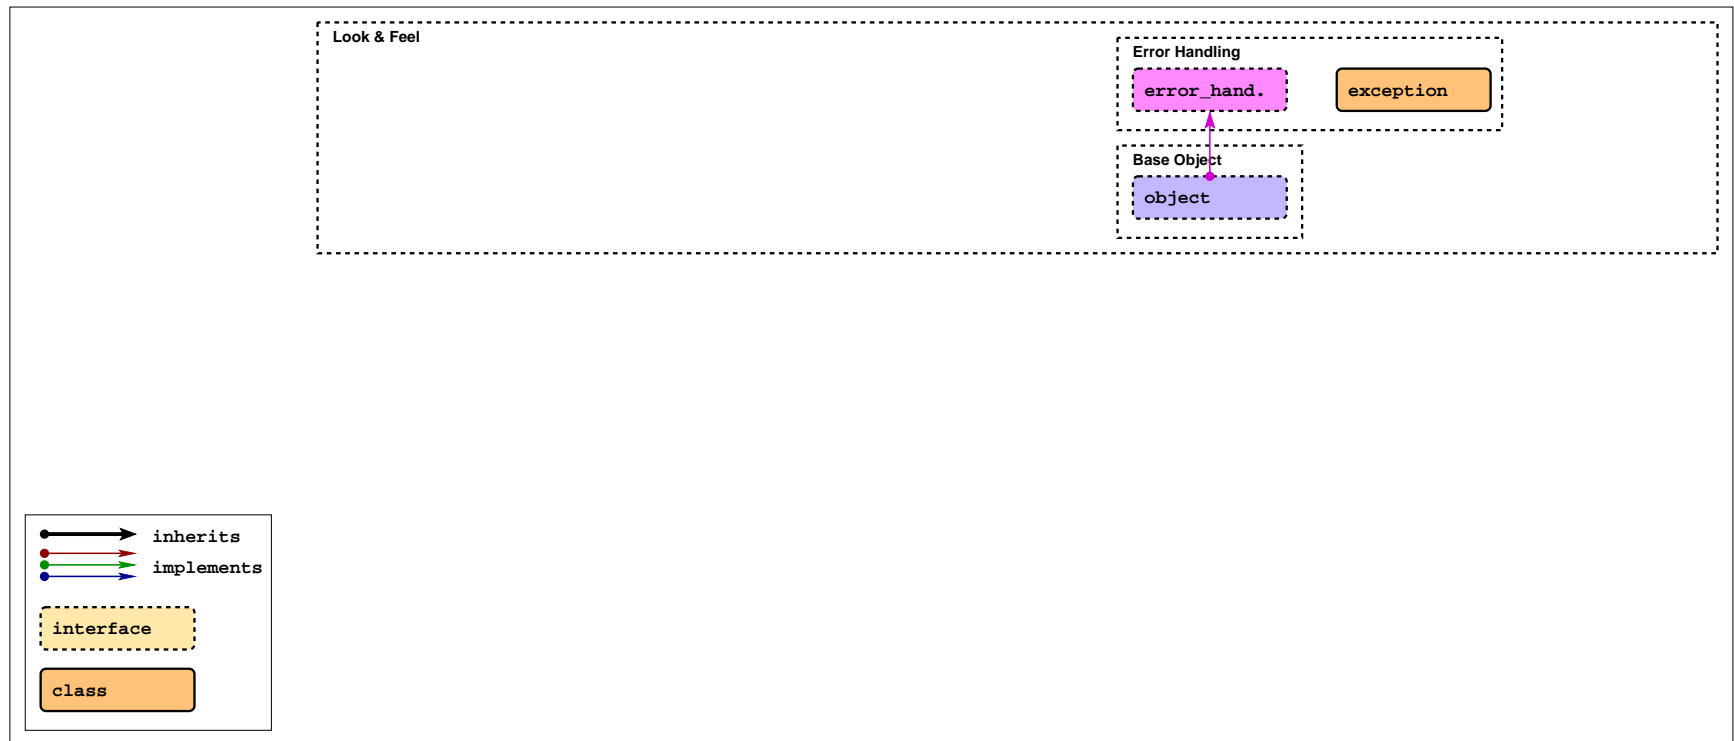
SAGA: Class hierarchy



SAGA Look & Feel:

`saga::object` allows for object uuids, `clone()` etc.

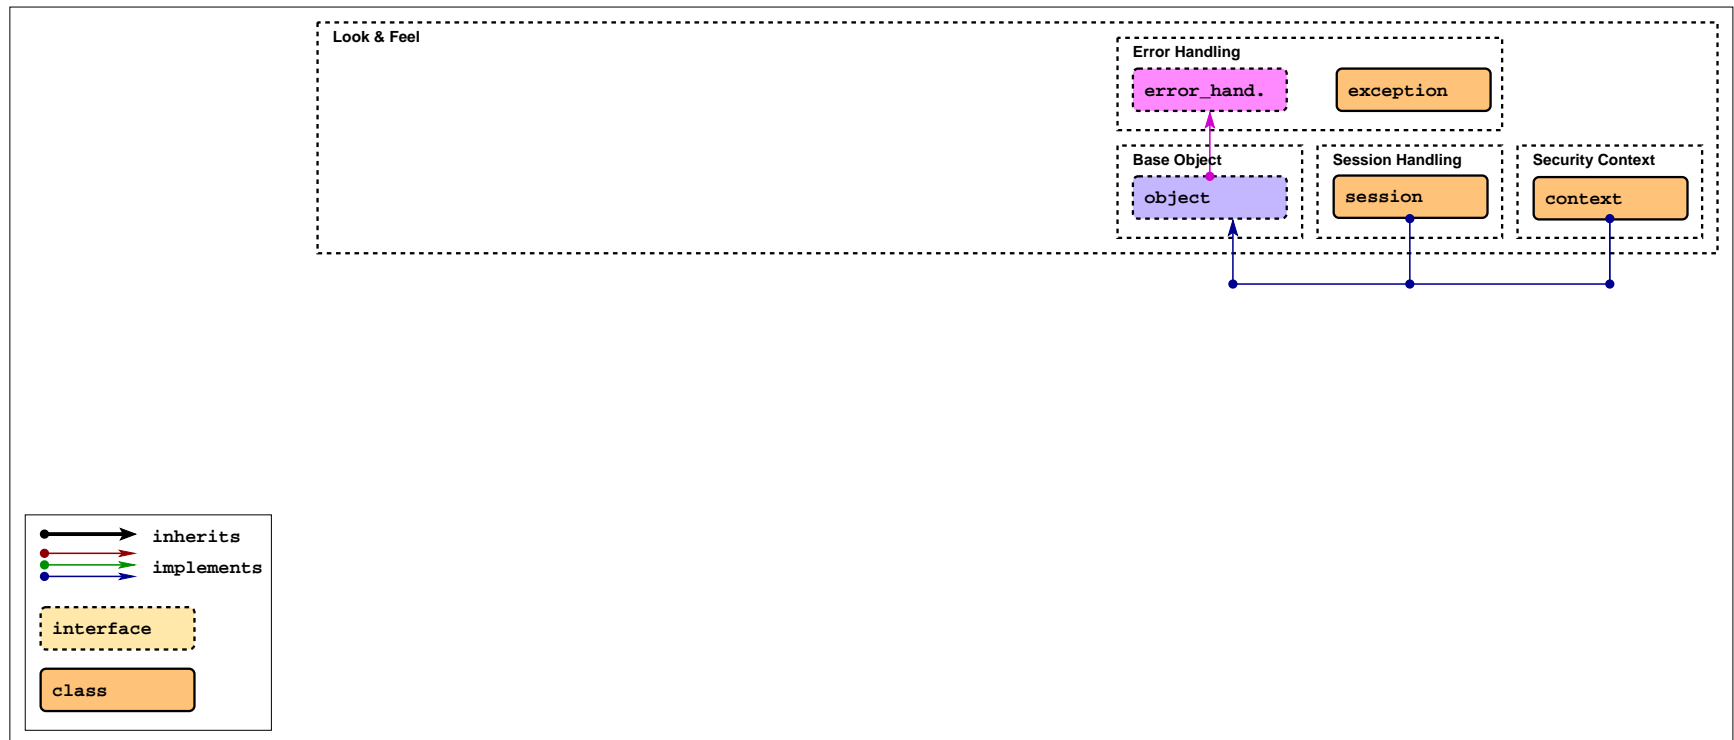
SAGA: Class hierarchy



SAGA Look & Feel:

errors are based on exceptions or error codes.

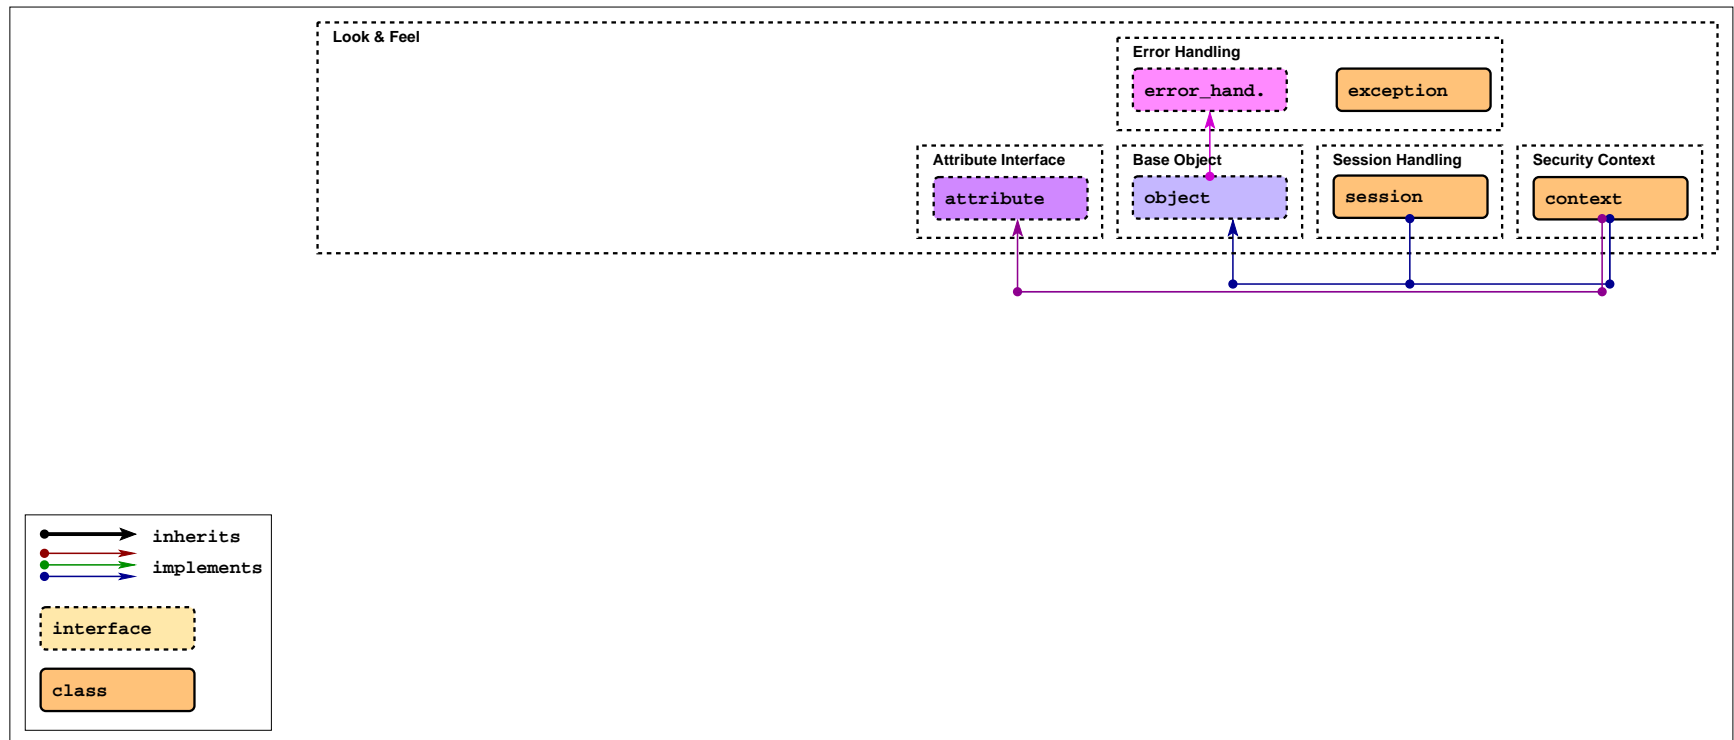
SAGA: Class hierarchy



SAGA Look & Feel:

session and credential management is hidden.

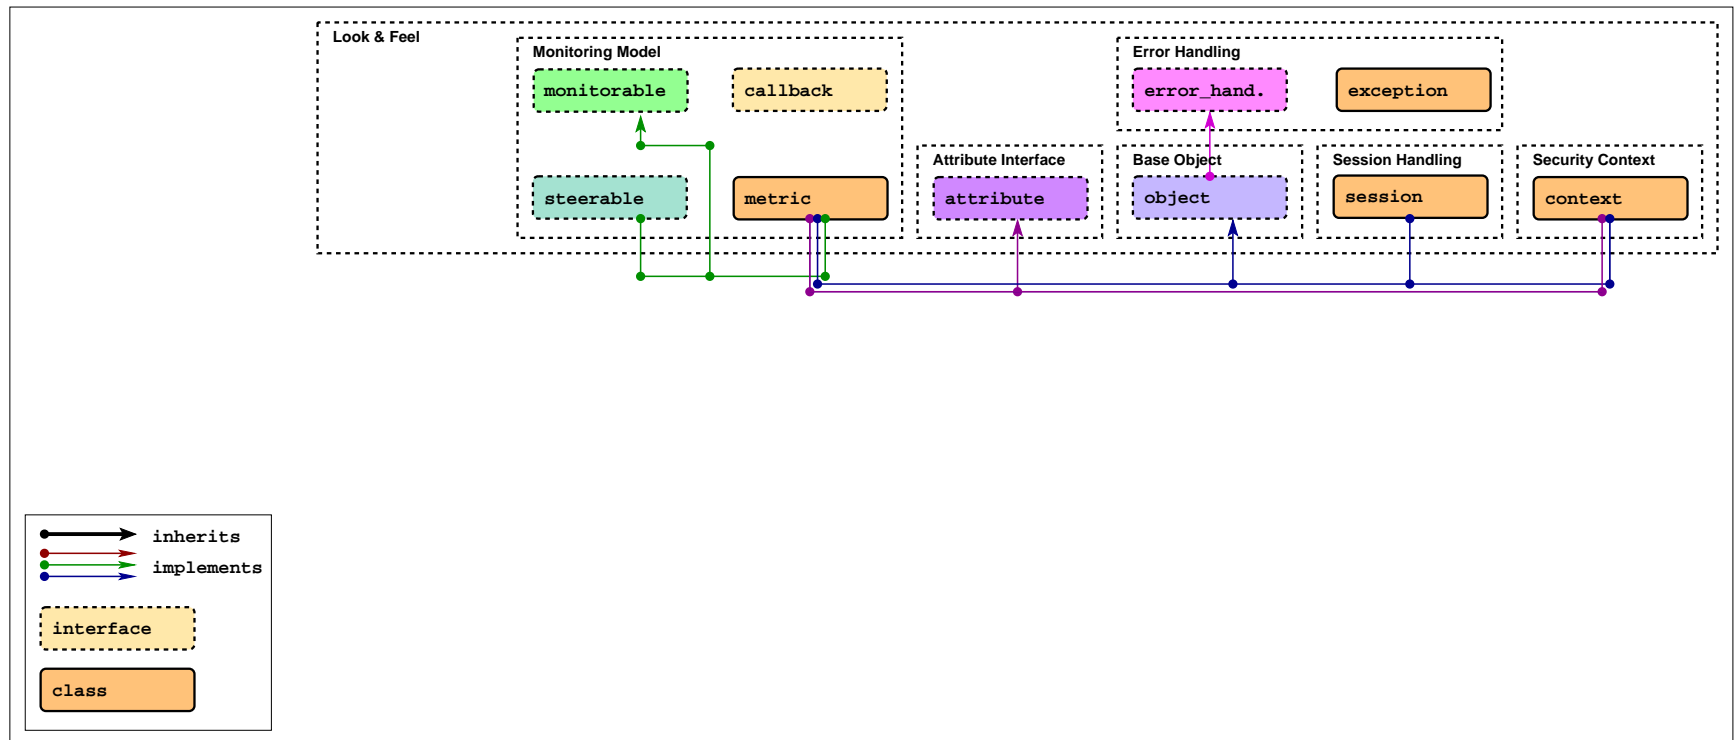
SAGA: Class hierarchy



SAGA Look & Feel:

Attribute interface for meta data.

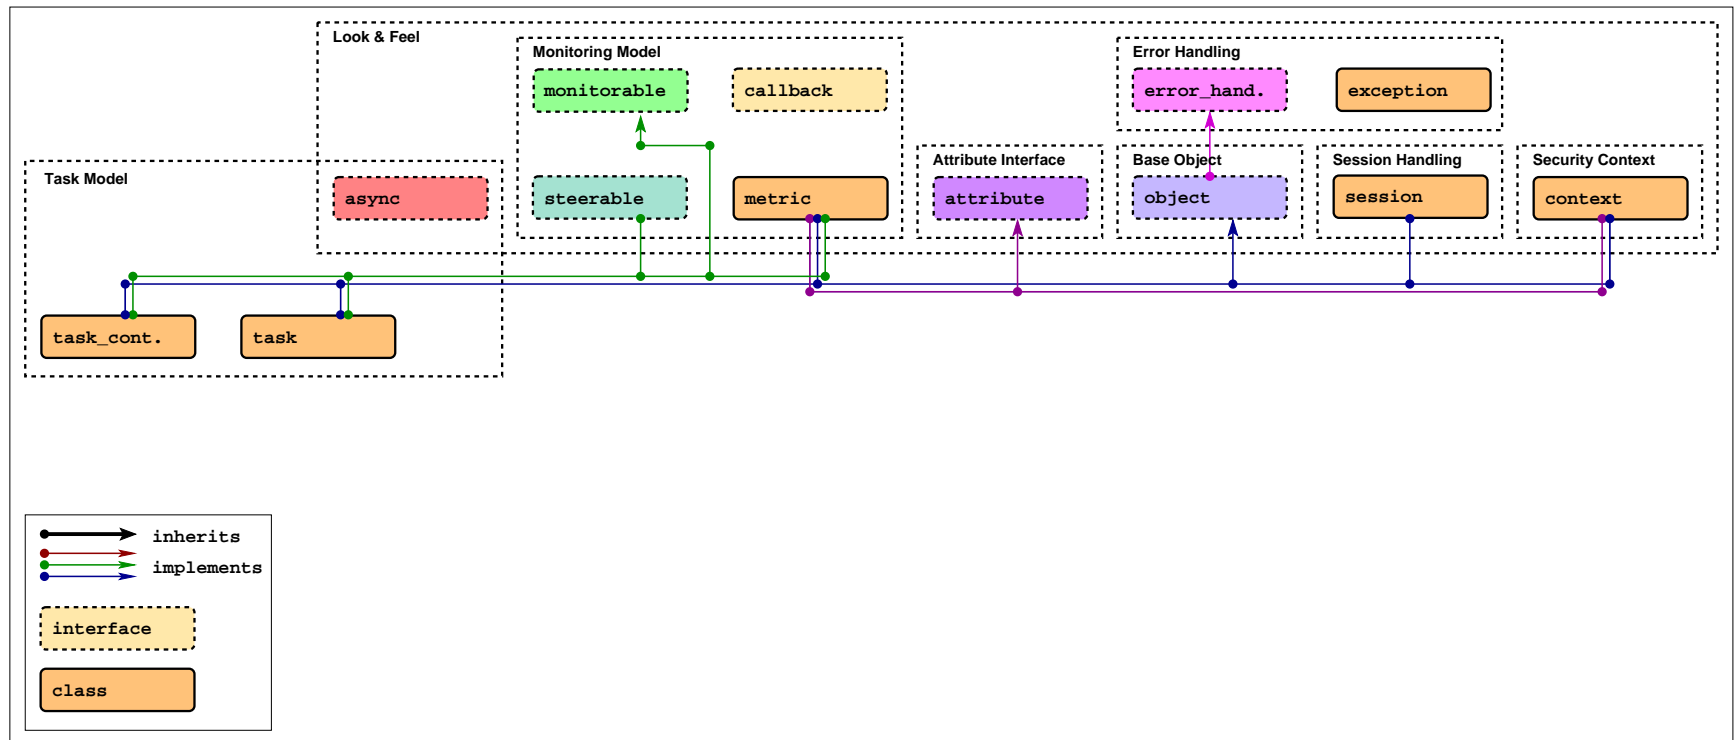
SAGA: Class hierarchy



SAGA Look & Feel:

Monitoring includes asynchronous notifications.

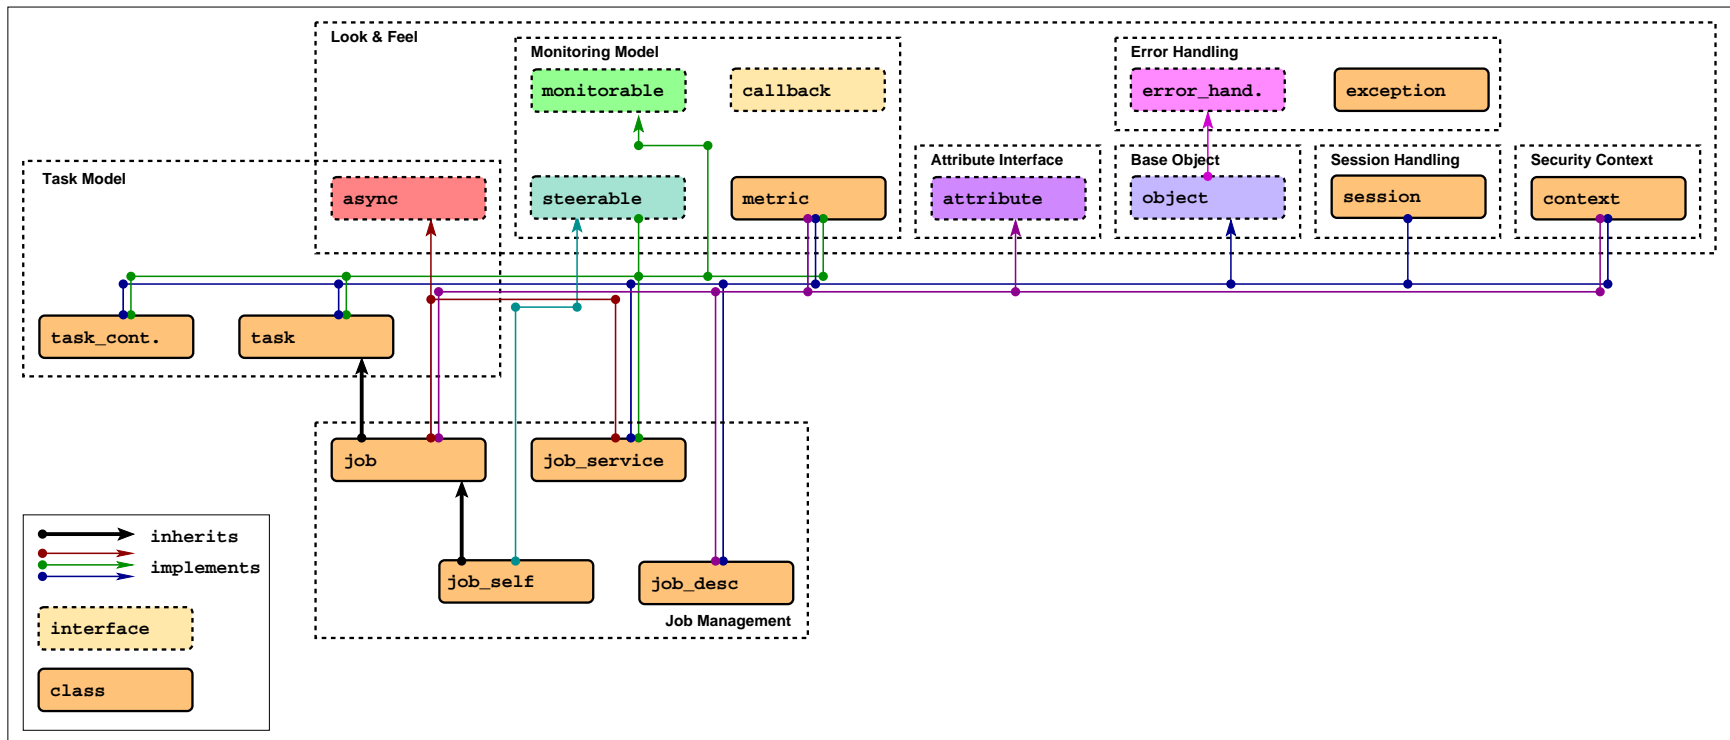
SAGA: Class hierarchy



SAGA Look & Feel:

the task model adds asynchronous operations.

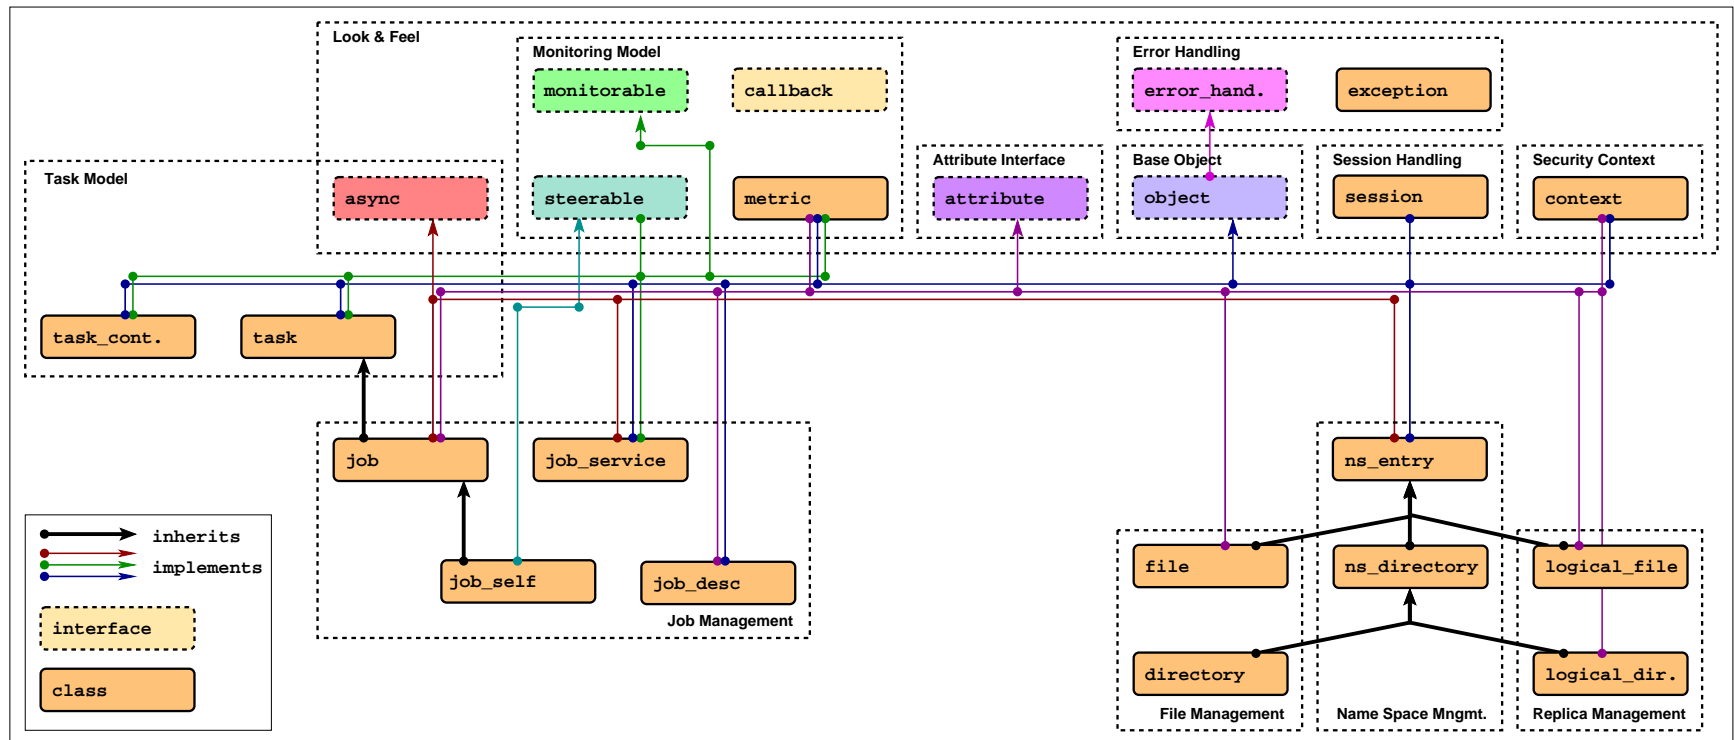
SAGA: Class hierarchy



SAGA API Package 'job':

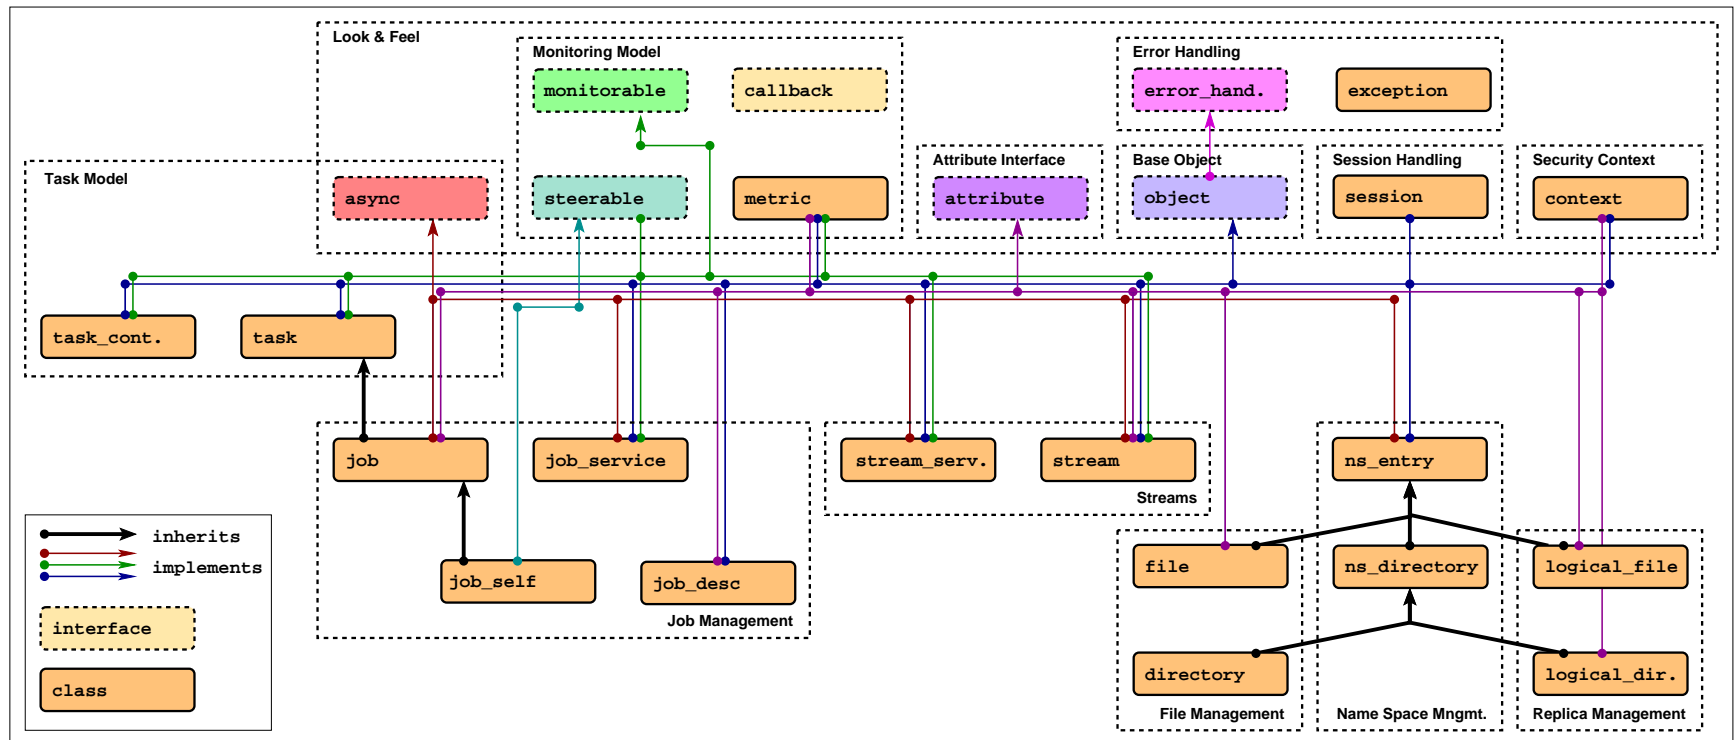
create and manage remote processes.

SAGA: Class hierarchy



SAGA API Package 'name_spaces':
manage files, replicas, etc.

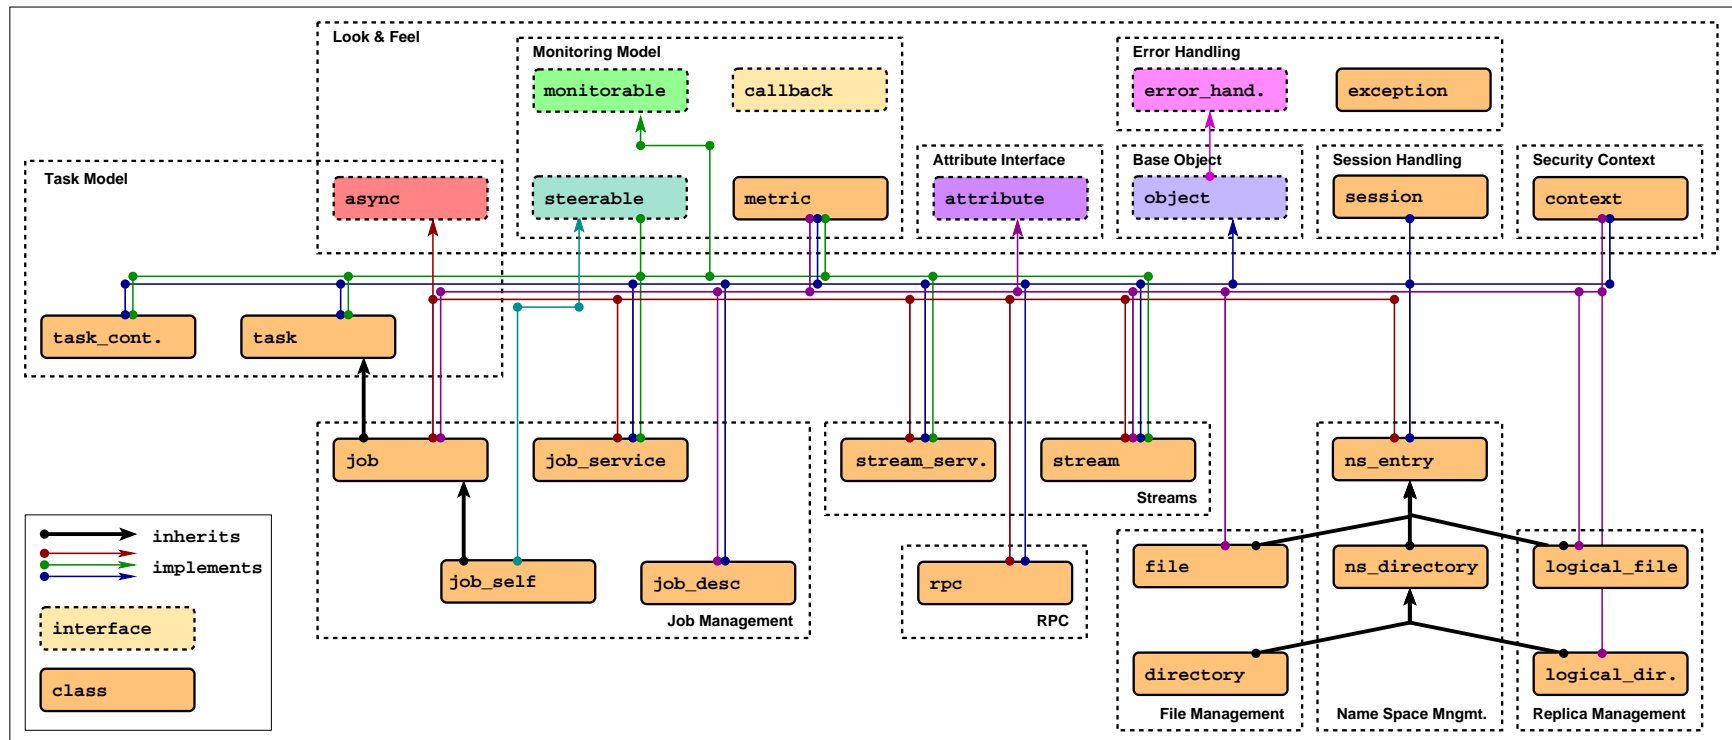
SAGA: Class hierarchy



SAGA API Package 'stream':

SAGA rendering of BSD streams.

SAGA: Class hierarchy

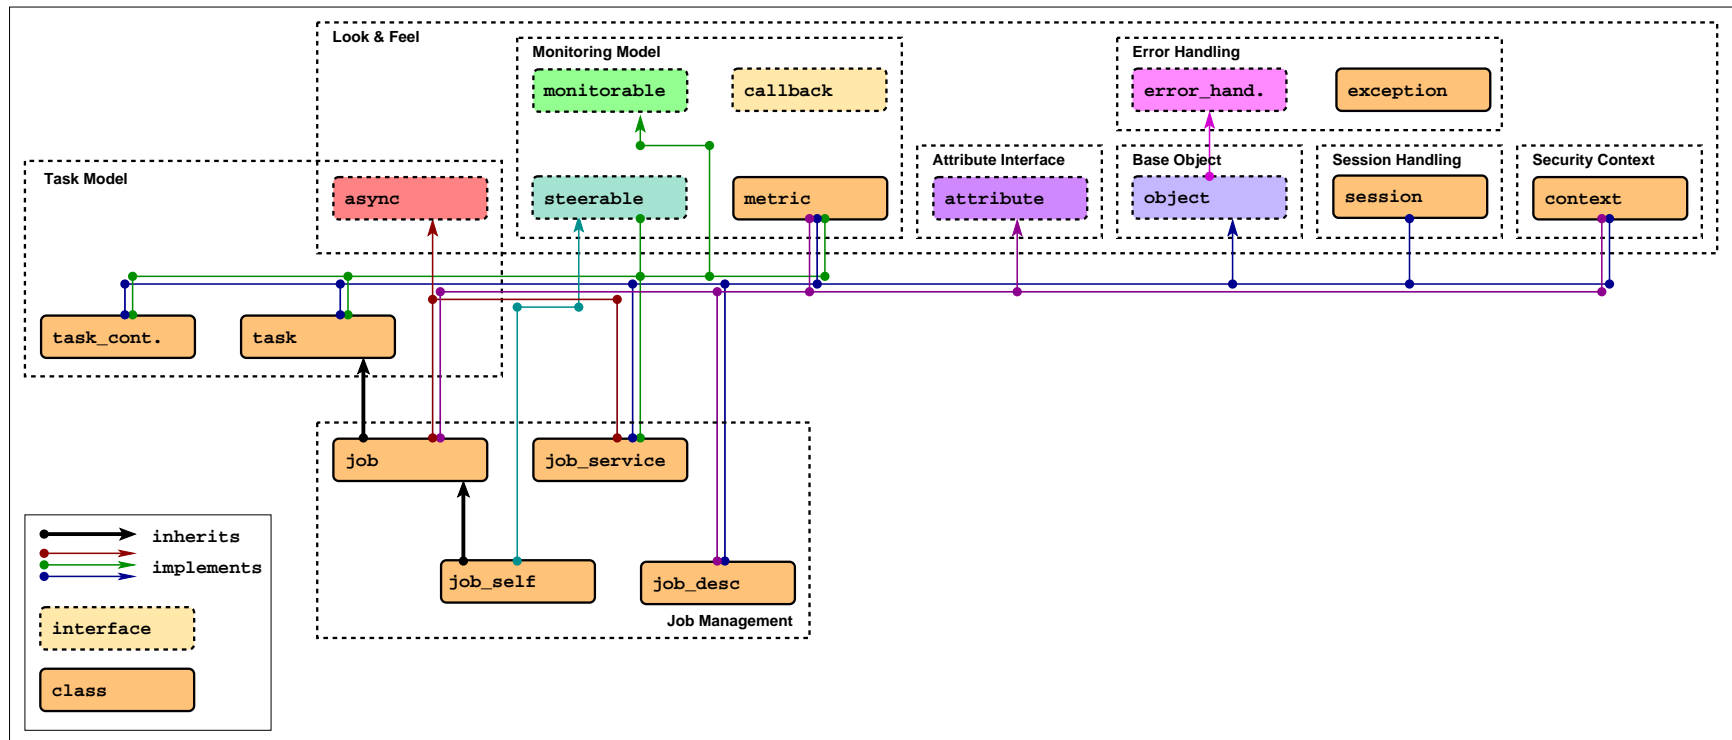


SAGA API Package 'rpc':

remote procedure calls.

Functional API Packages


SAGA: Jobs



SAGA: Jobs - Overview

- `job_service` uses `job_description` to create job instances
- `job_description` attributes are based on JSDL
- state model is based on / synced with BES
- `job_self` represents the SAGA application
- job submission and management, but no resource discovery, job dependencies, or workflows

SAGA: Job States



SAGA Examples: Jobs

job submission

```
saga::job::service      js ("gram://headnode.gram.net");
saga::job::job          j = js.run_job ("/bin/sleep 10",
                                         "clusternode-2.gram.net");

cout << "Job State: " << j.get_state () << endl;

j.wait ();

cout << "Retval " << j.get_attribute ("ExitCode") << endl;
```

SAGA Examples: Jobs

job submission

```
saga::job::description jd;  
saga::job::service      js ("gram://remote.host.net");  
saga::job                j = js.create_job (jd);  
  
j.run ();  
  
cout << "Job State: " << j.get_state () << endl;  
  
j.wait ();  
  
cout << "Retval " << j.get_attribute ("ExitCode") << endl;
```


SAGA Examples: Jobs

jobs (cont.)

```
j.run      ();  
j.wait     ();  
j.cancel   ();  
  
j.suspend  ();  
j.resume   ();  
  
j.signal    (SIGUSR1);  
j.checkpoint ();  
j.migrate   (jd);
```

SAGA Examples: Job Descr.

_____ job description - JSDL based _____

```
saga::job::description jd;

jd.set_attribute ("Executable",      "/bin/tail");
jd.set_attribute ("WorkingDirectory", "data/");
jd.set_attribute ("Cleanup",         "False");

// pseudo code *blush*
jd.set_vector_attribute ("Arguments",  ["-f", "my_log"]);
jd.set_vector_attribute ("Environment", ["TMPDIR=/tmp/"]);
jd.set_vector_attribute ("FileTransfer", ["my_log >> all_logs"]);
```

SAGA Job Description

- leaning heavily on **JSDL**, but flat
- borrowing from DRMAA
- mixes hardware, software and scheduling attributes!
- cannot be extended
- no support for 'native' job descriptions (RSL, JDL, ...)
- only 'Executable' is required
- backend MAY ignore unsupported keys!

```
cd /tmp/data && rm -rf *
```

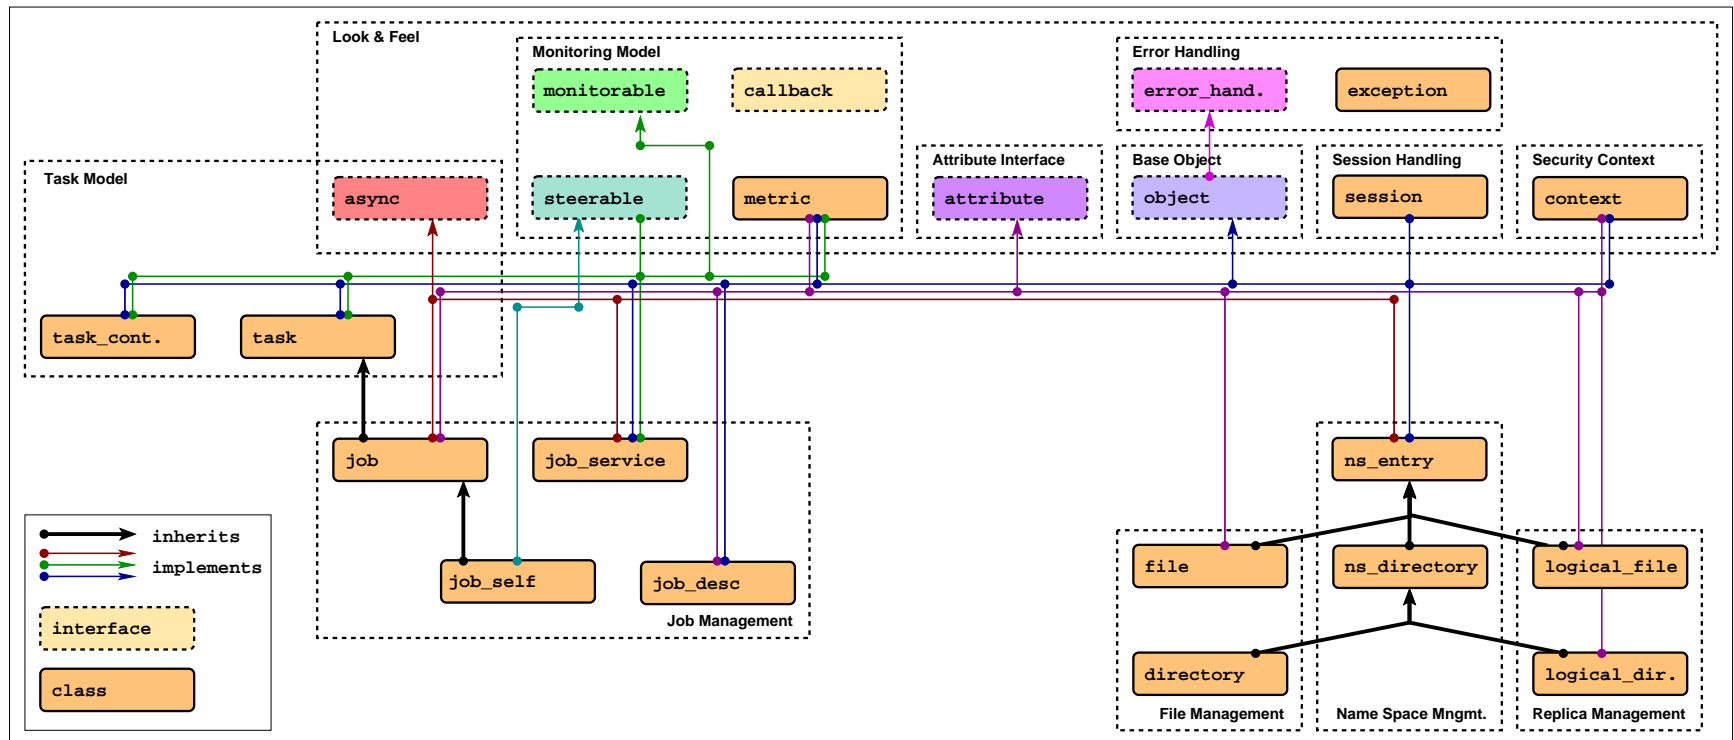
SAGA Example: job service

_____ job service _____

```
saga::job::service js ("gram://remote.host.net/");  
  
vector<string> ids = js.list (); // list known jobs  
  
while ( ids.size () )  
{  
    string id = ids.pop_back ();  
  
    saga::job j = js.get_job (id); // reconnect to job  
  
    cout << id << " : " << j.get_state () << endl;  
}
```

- represents a specific job submission endpoint
- job states are maintained on that endpoint (usually)
- full reconnect may not be possible (I/O streaming)
- lifetime of state up to backend
- reconnected jobs may have different job description (lossy translation)

SAGA: Name Spaces etc.



SAGA: Name Spaces

- interfaces for managing entities in name spaces
- files, replicas, information, resources, steering parameter, checkpoints, ...
- manages hierarchy (mkdir, cd, ls, ...)
- manages NS entries as opaque (copy, move, delete, ...)

- implements name space interface, and adds access to content of NS entries (files)
- Posix oriented: read, write seek
- Grid optimizations: scattered I/O, pattern based I/O, extended I/O

- implements name space interface, and adds access to properties of NS entries (logical files / replicas)
- O/REP oriented: list, add, remove replicas; manage meta data
- Grid optimizations are hidden (replica placement strategies, consistency and version management, ...)

- implements name space interface, and adds access to arbitrary key/value pairs on each entry.
- entries also allow to store a **serialized** SAGA Object!
- utterly useful for application bootstrapping, communication between different application modules, application persistency, etc etc.
- **not part of the SAGA Specification**, but is getting standardized as an extension

SAGA Examples: NameSpaces



_____ name space management _____

```
saga::name_space::directory d ("ssh://remote.host.net//data/");

if ( d.is_entry ("a") && ! d.is_dir ("a") )
{
    d.copy ("a", "../b");
    d.link ("../b", "a", Overwrite);
}

list <saga::url> names = d.find ("*-{123}.text.");

saga::name_space::directory tmp  = d.open_dir ("tmp/data/1",
                                                saga::name_space::CreateParents);
saga::name_space::entry      data = tmp.open    ("data.txt");

data.copy ("data.bak", Overwrite);      // uses cwd
```

- name space entries are opaque: the name space package can never look inside
- directories are entries (inheritance)
- **inspection:**
`get_cwd, get_url, get_name, exists, is_entry, is_dir, is_link, read_link`
- **manipulation:**
`create (c'tor, open), copy, link, move, remove`
- **permissions:**
`permissions_allow, permissions_deny`
- wildcards are supported (remember POSIX influence...)

SAGA Examples: Files

```
_____ file access _____  
  
saga::filesystem::file f ("any://remote.host.net/data/data.bin");  
  
char mem[1024];  
saga::mutable_buffer buf (mem);  
  
if ( f.get_size () >= 1024 )  
{  
    buf.set_data (mem + 0, 512);  
    f.seek (512, saga::filesystem::Start);  
    f.read (buf);  
}  
  
if ( f.get_size () >= 512 )  
{  
    buf.set_data (mem + 512, 512);  
    f.seek (0, saga::filesystem::Start);  
    f.read (buf);  
}
```

- provides access to the **content** of filesystem entries (sequence of bytes)
- saga buffers are used to wrap raw memory buffers
- saga buffers can be allocated by the engine
- several incarnations of read/write: posix style, scattered, pattern based

SAGA Name Spaces: Flags

```
enum flags {  
    None          = 0,  
    Overwrite     = 1,  
    Recursive     = 2,  
    Dereference   = 4,  
    Create        = 8,  
    Exclusive     = 16,  
    Lock          = 32,  
    CreateParents = 64,  
    Truncate      = 128,    // not on name_space  
    Append        = 256     // not on name_space  
    Read          = 512,  
    Write         = 1024,  
    ReadWrite     = 1536    // Read | Write  
    Binary        = 204     // only on filesystem  
}
```

SAGA Examples: Replicas

_____ replica management _____

```
saga::replica::directory dir ("raptor://remote.host.net/data/");

if ( dir.is_entry ("a") || dir.is_link ("a") )
{
    dir.copy ("a", "../b");
    dir.link ("../b", "a");
}

saga::replica::file file = dir.open ("tmp/data.txt");
list <string>  locations = file.list_locations ();

file.replicate ("gridftp://other.host.net/tmp/a.dat");
```


- provides access to the **content** of replica system entries (list of physical locations, plus attributes)
- saga attribute interface is used for entry meta data. Meta data are maintained by application, and/or backend.
- `replicate()` creates a new copy, and adds new location to list

SAGA Examples: Replicas

```
_____ replica meta data _____  
  
saga::replica::directory dir ("raptor://remote.host.net/data/");  
  
list <saga::url> files = dir.find ("*", "type=jpg");  
  
while ( file.size () )  
{  
    saga::logical_file lf (file.pop_front ());  
  
    lf.replicate ("file://localhost/data/images/",  
                 saga::replica::Overwrite);  
}
```

- persistent storage of application level information
- semantics of information defined by application
- allows storage of serialized SAGA objects (object persistency)

SAGA Examples: Adverts

Adverts

```
saga::advert::directory todo ("any//remote.host.net/my_tasks/");

// pseudo vector code
list <saga::url> urls = todo.find ("*", ["priority=urgent"]);

while ( urls.size () )
{
    saga::advert ad (urls.pop_front ());

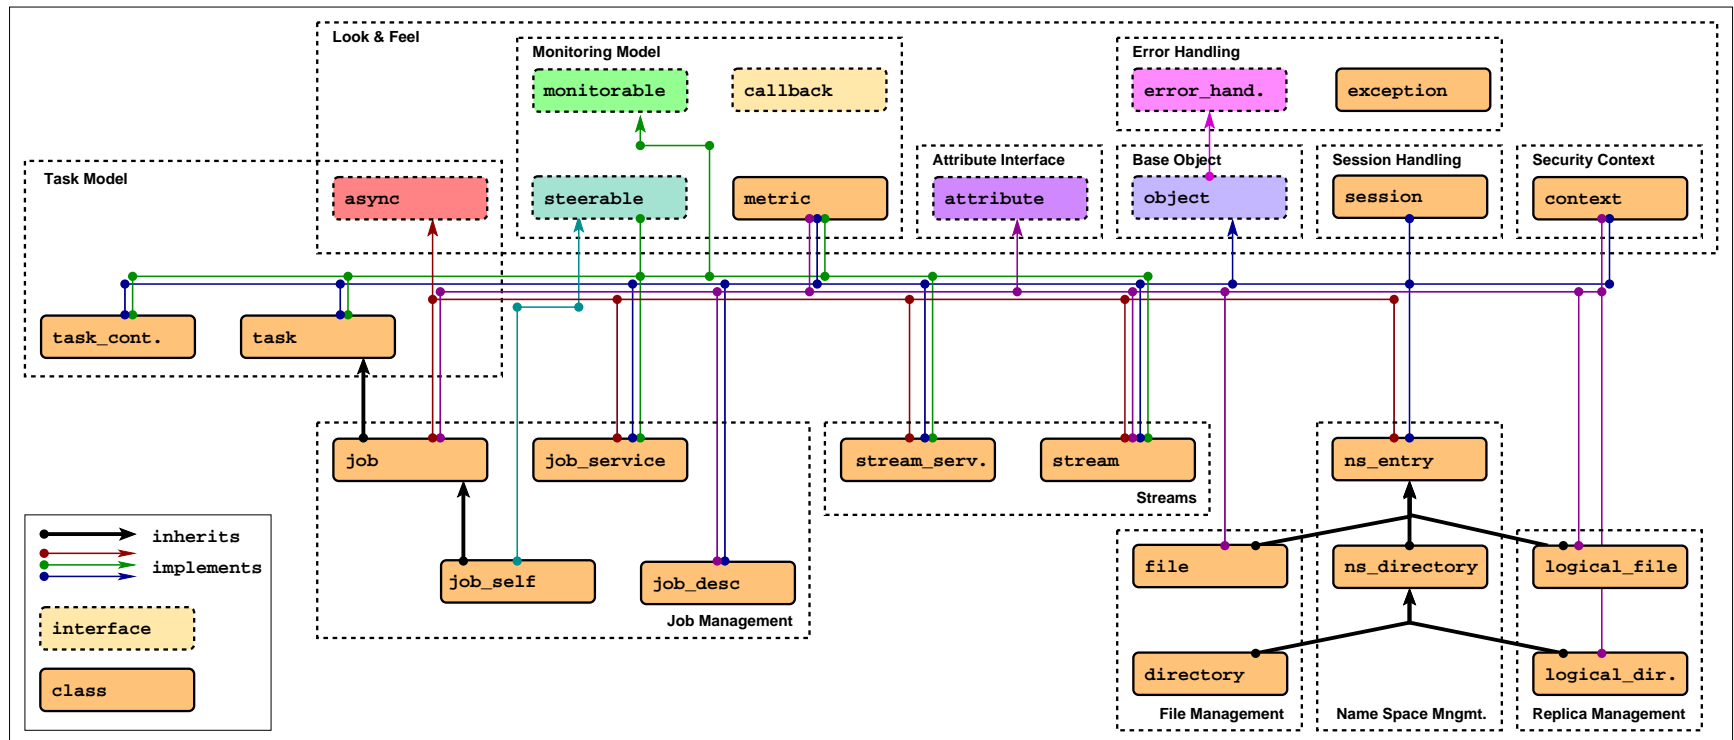
    std::cout << ad.get_attribute ("description") << std::endl;
}
```

SAGA Examples: Adverts

Persistent SAGA Objects

```
saga::file    f    (url);  
saga::advert  ad ("any//remote.host.net/files/my_file_ad", Create);  
  
ad.store_object (f);  
  
-----  
  
saga::advert  ad ("any//remote.host.net/files/my_file_ad");  
saga::file    f = ad.retrieve_object ();
```

SAGA: Streams



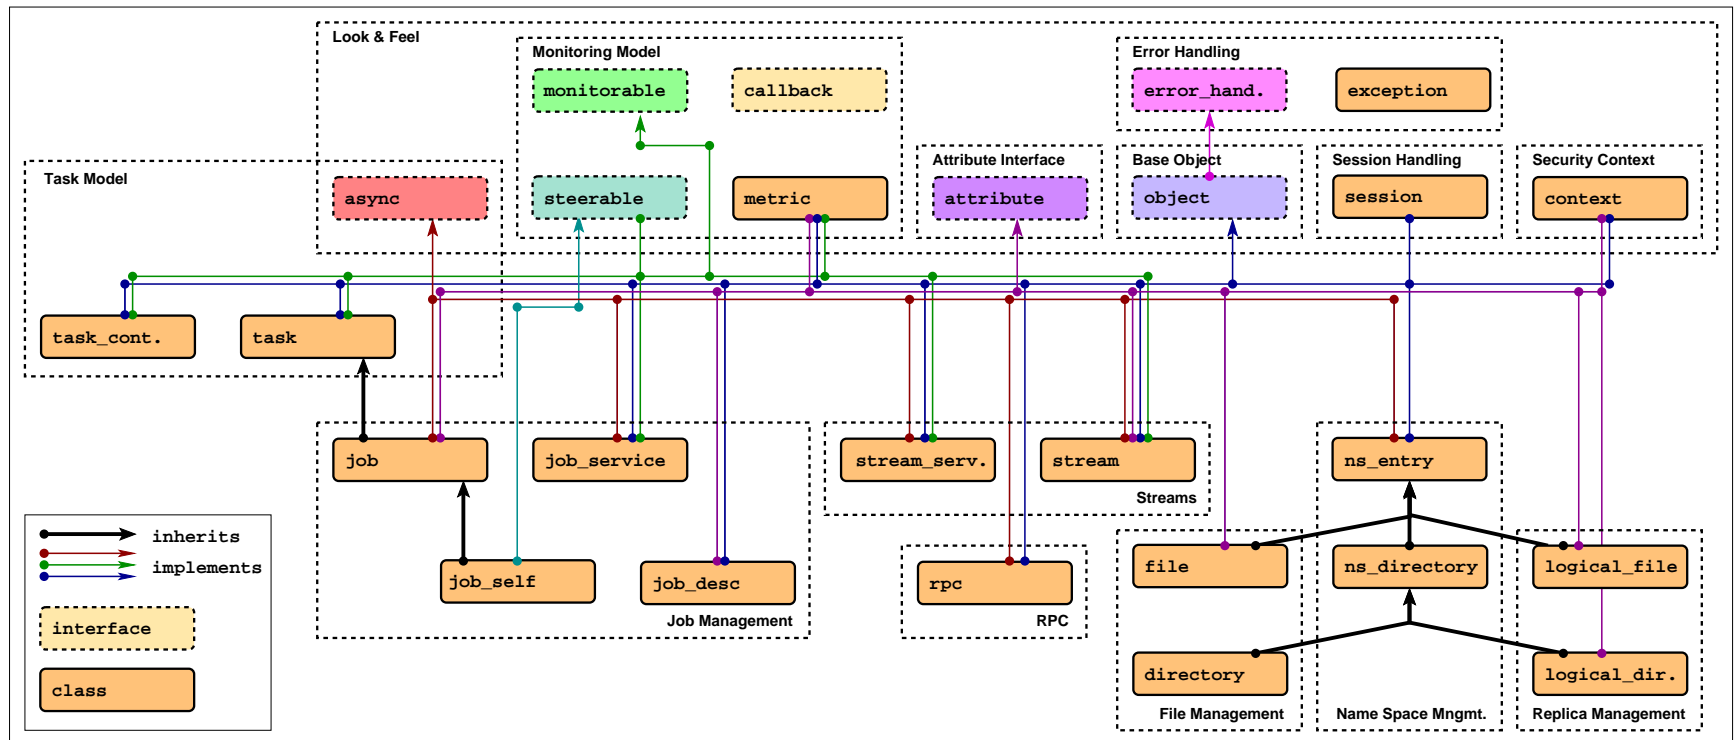
SAGA Examples: Streams

```
_____ stream server _____  
saga::stream_service ss ("tcp://localhost:1234");  
  
saga::stream_client sc = ss.serve ();  
  
sc.write ("Hello client", 13);
```

```
_____ stream client _____  
char buf [13];  
saga::stream_client sc ("tcp://remote.host.net:1234");  
  
sc.connect ();  
sc.read      (buf, 13);  
  
cout << buf << endl;
```

- simple and BSD socket oriented
- not supposed to replace MPI etc, but allows for simple application level communication
- will be superceded by message package

SAGA: RPC



SAGA Examples: RPC

_____ remote procedure call _____

```
saga::rpc rpc ("ninfgr://remote.host.net:1234/random");

list <saga::rpc::parameter> params;
params.push_back (new saga::rpc::parameter (Out, 10));

rpc.call (params);

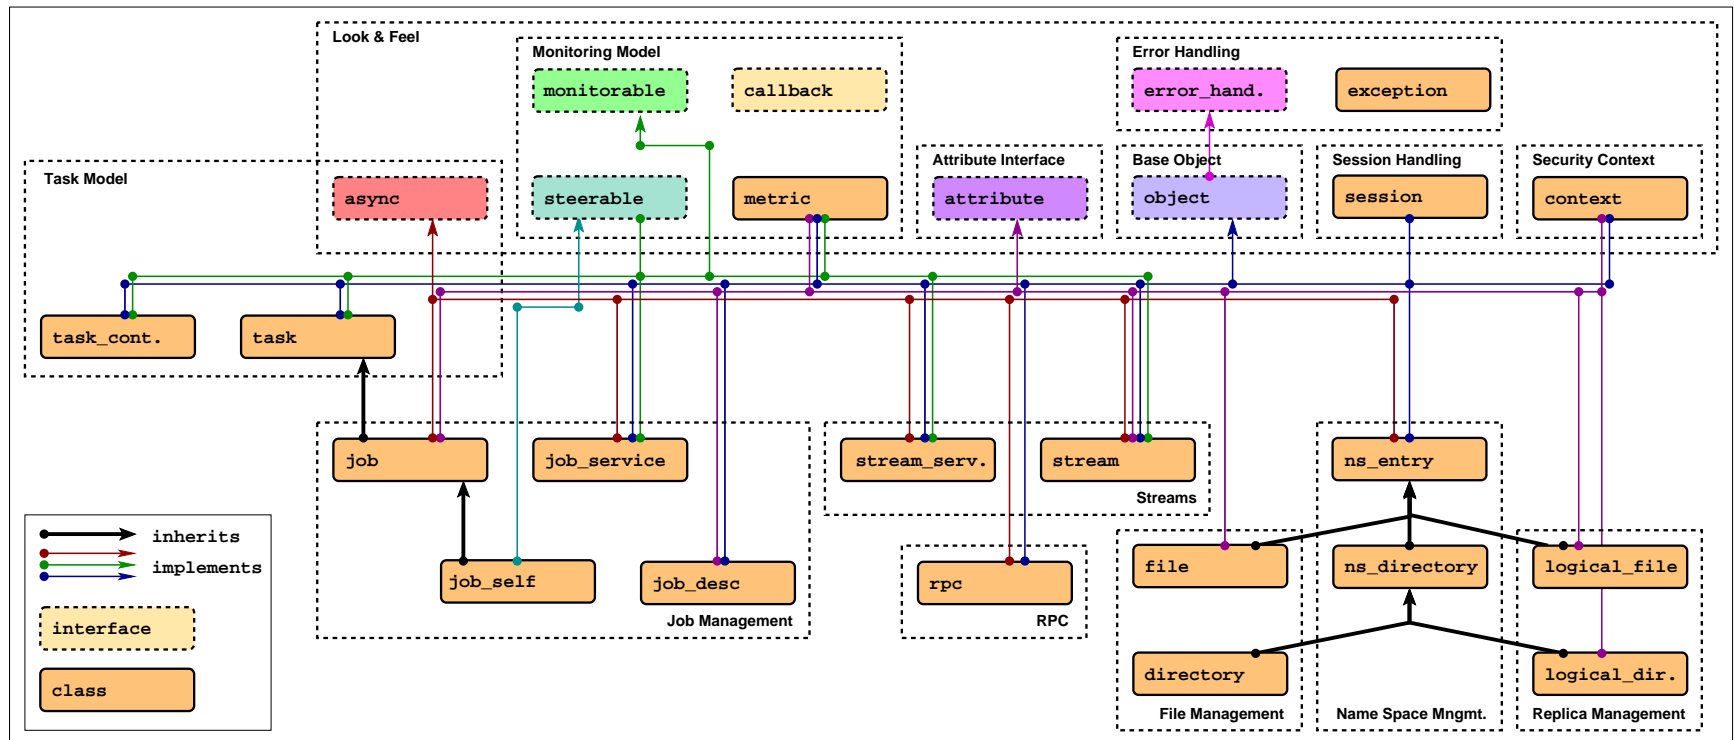
cout << "found random number: " << ::atoi (param.buffer) << endl;

delete (params.pop_front ());
```

- maps GridRPC standard into the SAGA look & feel
- parameters are stack of structures (similar to scattered I/O)
- future revision will work on optimized data handling

Non-Functional API Packages

SAGA: Session and Context



SAGA Examples: Session

_____ default sessions _____

```
saga::ns_dir dir ("any://remote.host.net//data/");

if ( dir.is_entry ("a") && ! dir.is_dir ("a") )
{
    dir.copy ("a", "../b");
    dir.link ("../b", "a", Overwrite);
}

list <saga::url> names = dir.find ("*-{123}.text.");

saga::name_space::directory tmp = dir.open_dir ("tmp/");
saga::name_space::entry entry = tmp.open ("data.txt");

entry.copy ("data.bak", Overwrite);
```

SAGA Examples: Session

context management

```
saga::context c1 (saga::context::X509);  
saga::context c2 (saga::context::X509);  
  
c2.set_attribute ("UserProxy", "/tmp/x509up_u123.special");  
  
saga::session s;  
  
s.add_context (c1);  
s.add_context (c2);  
  
saga::name_space::dir dir (s, "any://remote.host.net/data/");
```

SAGA: Session Management

- by default hidden (default session is used)
- session is identified by lifetime of security credentials and by objects in this session (jobs etc.)
- session is used on object creation (optional)
- `saga::context` is used to attach security tokens to a session
- the default session has default contexts

SAGA Examples: Session

_____ session inheritance _____

```
saga::dir  dir (s, "gridftp://remote.host.net/data/");  
  
saga::file file = dir.open ("data.bin");  
  
s.remove_context (c1);  
s.remove_context (c2);  
  
file.copy ("data.bin.bak");    // works - state is sticky!
```

SAGA Examples: Contexts

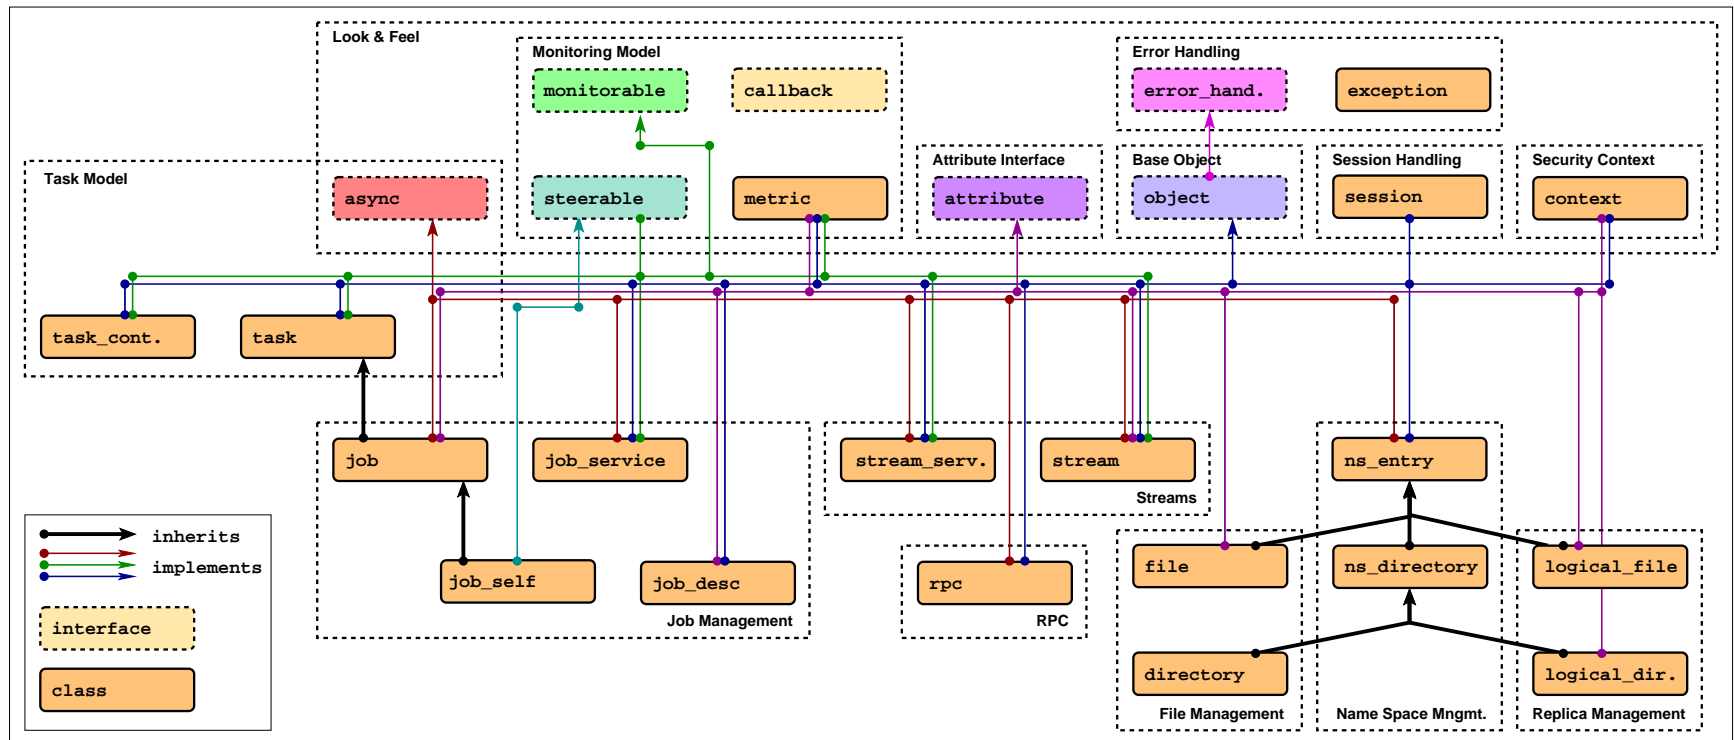
```
                                authorization
// server side code
saga::stream_service ss ("tcp://localhost:1234");

saga::stream_client sc = ss.serve ();

saga::context c = sc.get_context ();

if ( c.get_type == Globus &&
      c.attribute_equals ("RemoteID", "O=MyCA, O=MyOrg, CN=Joe" ) )
{
    sc.write ("welcome!", 9);
}
else
{
    sc.write ("bugger off!", 12);
    sc.close ();
}
```

SAGA: Monitoring



- monitoring of Grid entities (jobs, files, . . .)
- monitoring of interactions (task state, notification, . . .)
- `monitorables` have metrics
- `metrics` can be pulled, or subscribed to (callbacks)
- some metrics can be written (basic steering)

SAGA Examples: Monitoring

pull monitoring

```
saga::job::job j = js.create_job (jd);  
  
j.run ();  
  
saga::metric m = j.get_metric ("MemoryUsage");  
  
while ( 1 )  
{  
    cout << "Memory Usage: " << m.get_value () << endl;  
    sleep (1);  
}
```

SAGA Examples: Monitoring

```
_____ callbacks _____  
class my_cb : public saga::callback  
{  
    public:  
        bool cb (saga::monitorable obj,  
                 saga::metric      m,  
                 saga::context      c)  
        {  
            cout << "Memory Usage: " << m.get_value () << endl;  
            return (true);  
        }  
};  
  
my_cb cb;  
saga::job::job j = js.create_job (jd);  
j.run ();  
  
saga::metric m = j.get_metric ("MemoryUsage");  
m.add_callback ("MemoryUsage", cb);
```

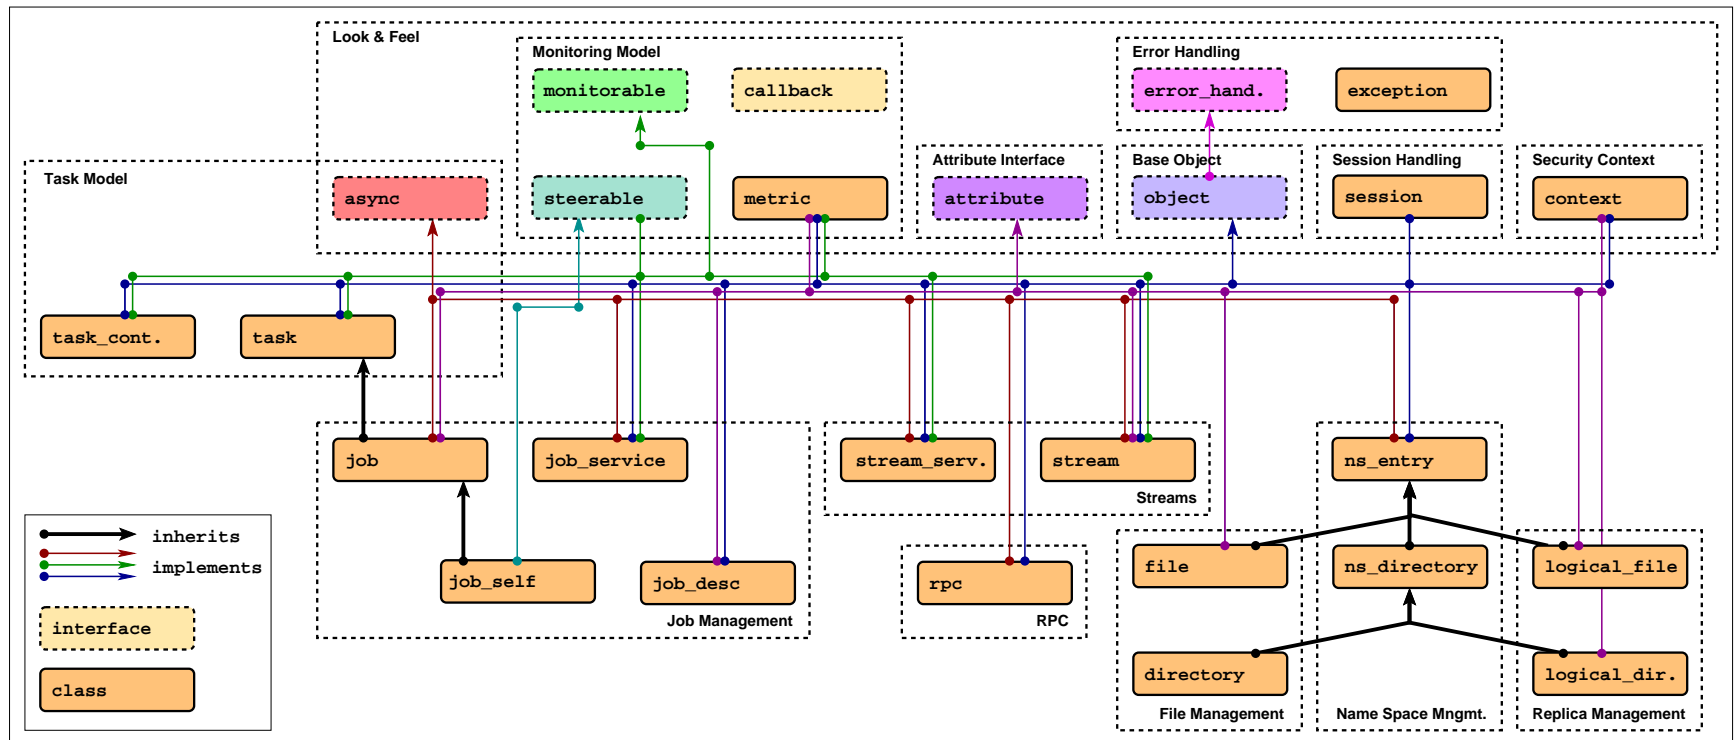
SAGA Examples: Monitoring

```
_____ callbacks _____  
class my_cb : public saga::callback  
{  
    public:  
        bool cb (saga::monitorable obj,  
                 saga::metric      m,  
                 saga::context      c)  
        {  
            cout << "Memory Usage: " << m.get_value () << endl;  
            return (true);  
        }  
};  
  
my_cb cb;  
saga::job::job j = js.create_job (jd);  
j.run ();  
  
j.add_callback ("MemoryUsage", cb);
```

SAGA Examples: Monitoring

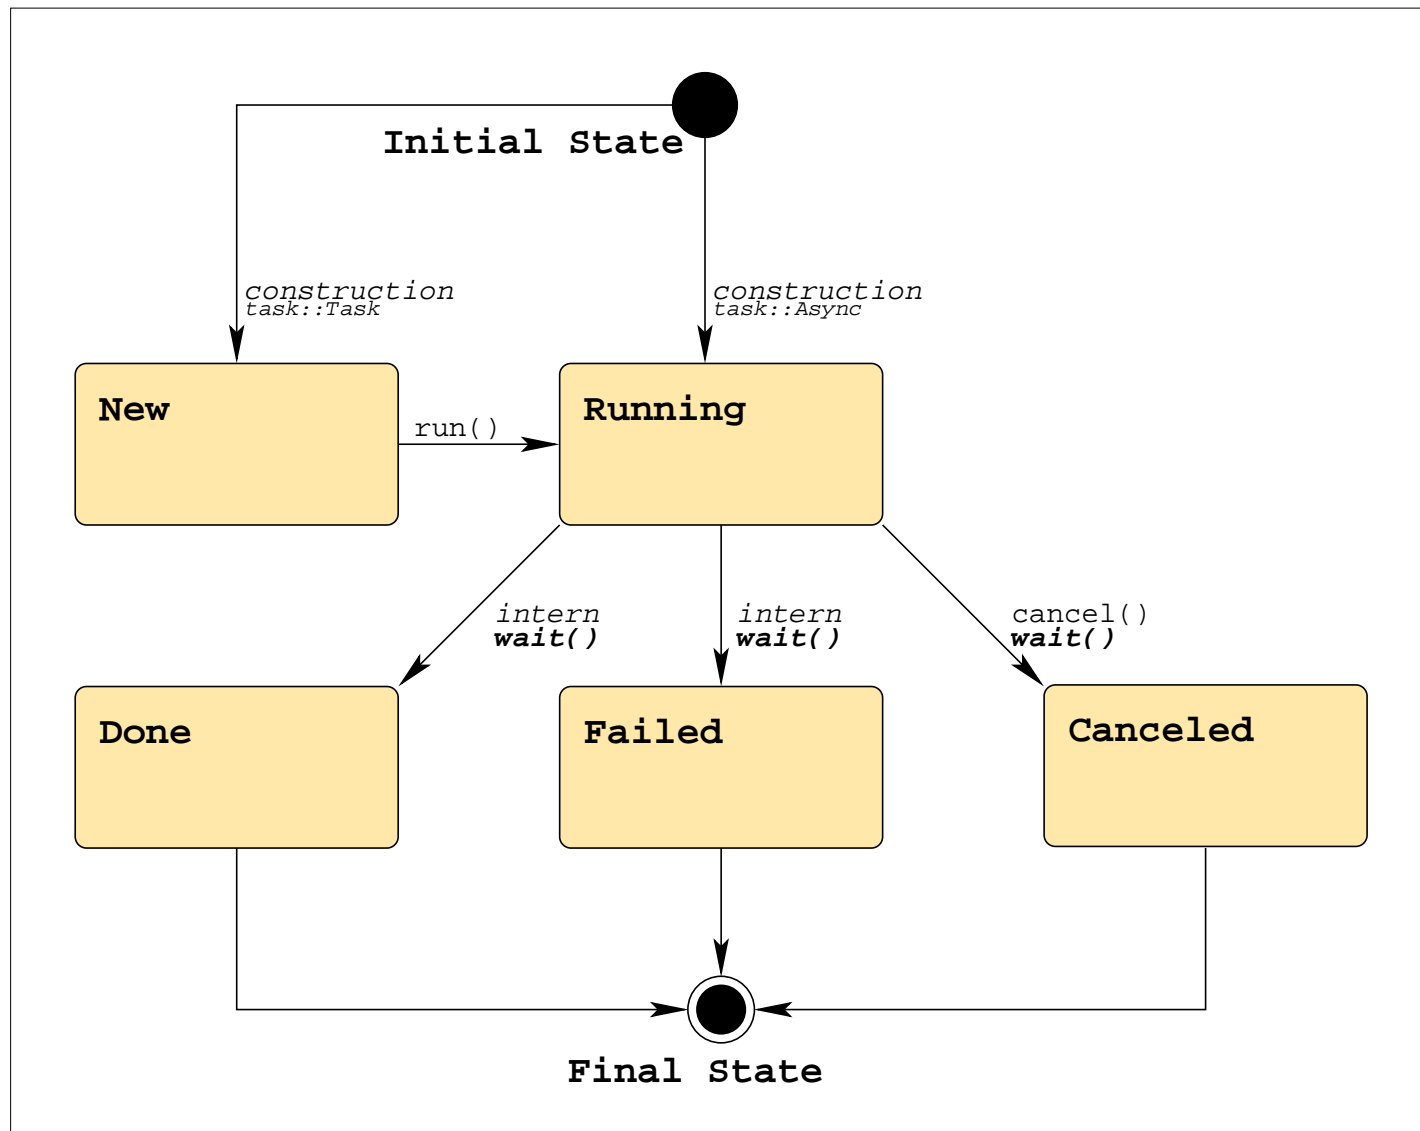
```
_____ callbacks (cont.) _____  
class my_cb : public saga::callback  
{  
    public:  
        bool cb (saga::monitorable obj,  
                 saga::metric      m,  
                 saga::context      c)  
        {  
            cout << m.get_name () << " : " << m.get_value () << endl;  
            return (true);  
        }  
};  
  
list <string> metrics = j.list_metrics ();  
  
while ( metrics.size () )  
{  
    j.add_callback (metrics.pop_front (), cb);  
}
```


SAGA: Tasks




- asynchronous operations are a MUST in distributed systems, and Grids
- `saga::task` represents an synchronous operation (e.g. `file.copy ()`)
- `saga::task_container` manages multiple tasks
- tasks are stateful (similar to jobs)

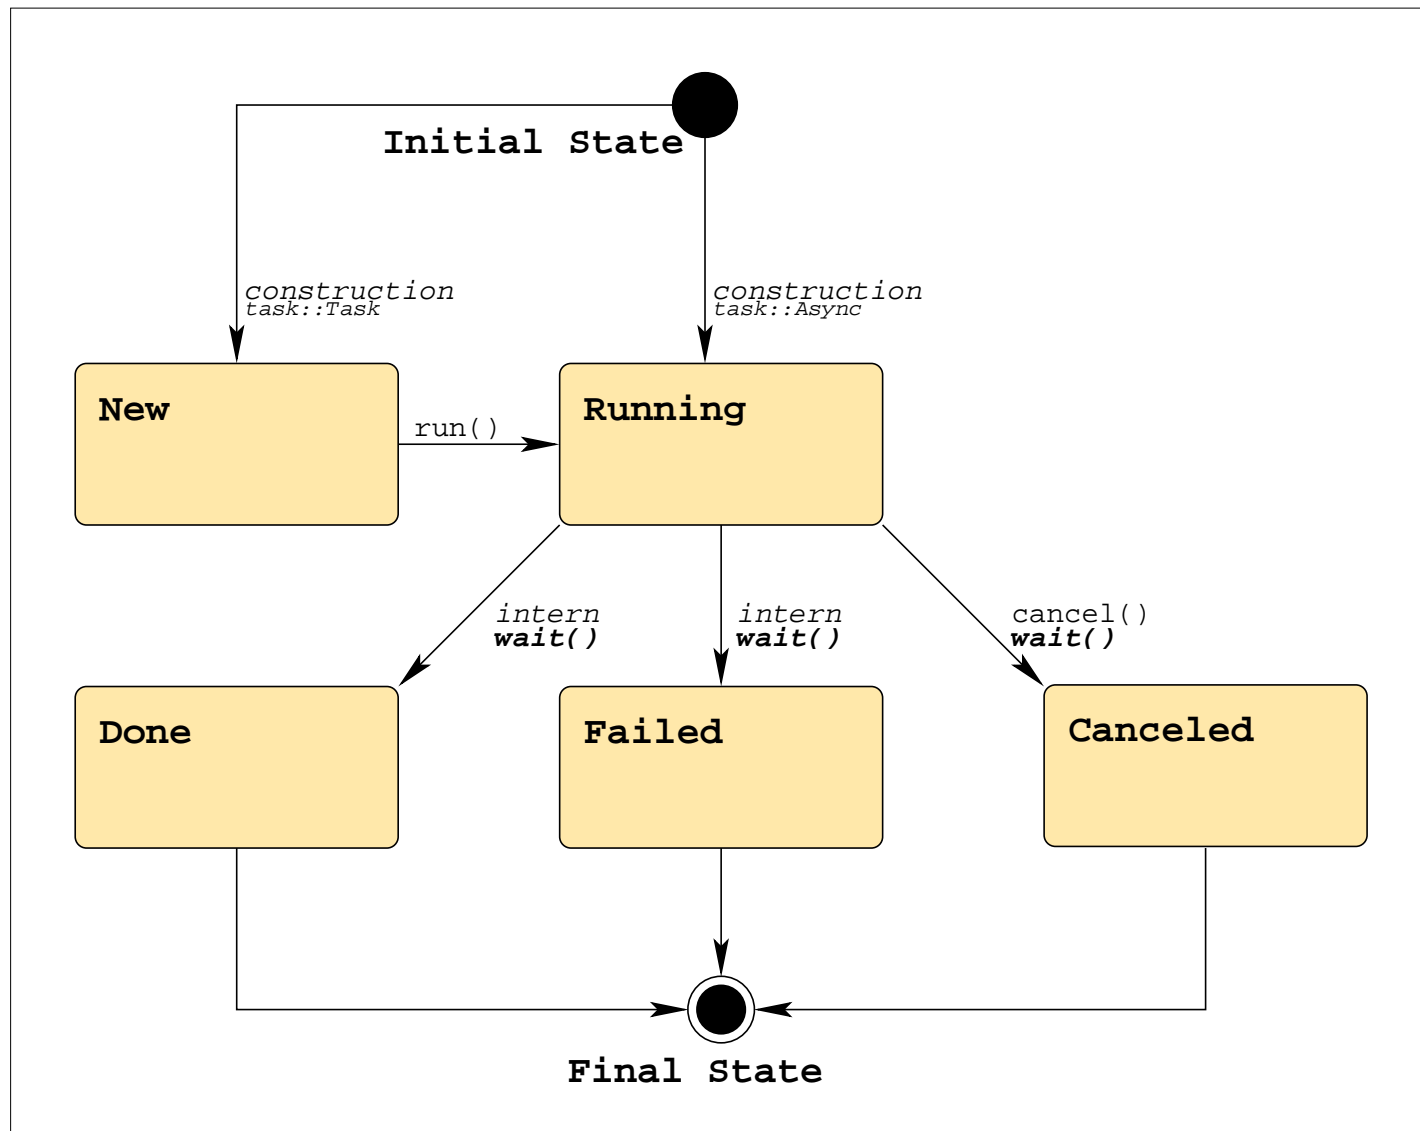
SAGA: Task States



SAGA: Job States



SAGA: Task States



- different versions for each method call: sync, async, task
- signature basically the same
- differ in state of task returned by that method

SAGA Examples: Tasks

```
_____ tasks (i) _____

saga::file file ("gsiftp://remote.host.net/data/data.bin");

// normal, synchronous
file.copy ("data.bak"); // void

// async versions, never throw (use 'rethrow' on failure)
saga::task t1 = file.copy <saga::task::Sync>  ("data.bak.1");
saga::task t2 = file.copy <saga::task::Async>  ("data.bak.2");
saga::task t3 = file.copy <saga::task::Task>   ("data.bak.3");

// t1: Done
// t2: Running
// t3: New
```

SAGA Examples: Tasks

```
_____ tasks (ii) _____

saga::file file ("gsiftp://remote.host.net/data/data.bin");

// normal, synchronous
size_t s = file.get_size ();

// async versions
saga::task t1 = file.get_size <saga::task::Sync> ();
saga::task t2 = file.get_size <saga::task::Async> ();
saga::task t3 = file.get_size <saga::task::Task> ();

// get_result implies wait, and can throw!
size_t s1 = t1.get_result <size_t> ();
size_t s2 = t2.get_result <size_t> ();
size_t s3 = t3.get_result <size_t> ();
```


SAGA Examples: Tasks

```
tasks (iii)

t3.run ();

cout << t3.get_state () << endl;  // Running

t2.wait ();
t3.wait ();

// t1, t2, t3: Done (or Failed...)
```

SAGA Examples: Tasks

tasks container

```
saga::task_container tc;  
  
tc.add (t1);  
tc.add (t2);  
tc.add (t3);  
  
tc.run  ();  
  
saga::task done_task = tc.wait (Any);  
  
tc.wait (All);
```

SAGA Examples: Tasks

_____ tasks and jobs _____

```
saga::task task = file.copy <saga::task::Async> ("b");  
saga::job  job  = js.run_job ("remote.host.net", "/bin/date");  
  
task.add_callback ("State", my_cb);  
job.add_callback  ("State", my_cb);  
  
saga::task_container tc;  
  
tc.add (task);  
tc.add (job);  
  
tc.wait ();
```

— — —

SAGA planned extensions

- service discovery
- message based communication
- information service (Advert Service)
- resource discovery and management
- checkpoint & recovery (GridCPR)

Questions?