

SAGA: Simple Examples, Programming Manual SAGA-Shell, Example Applications

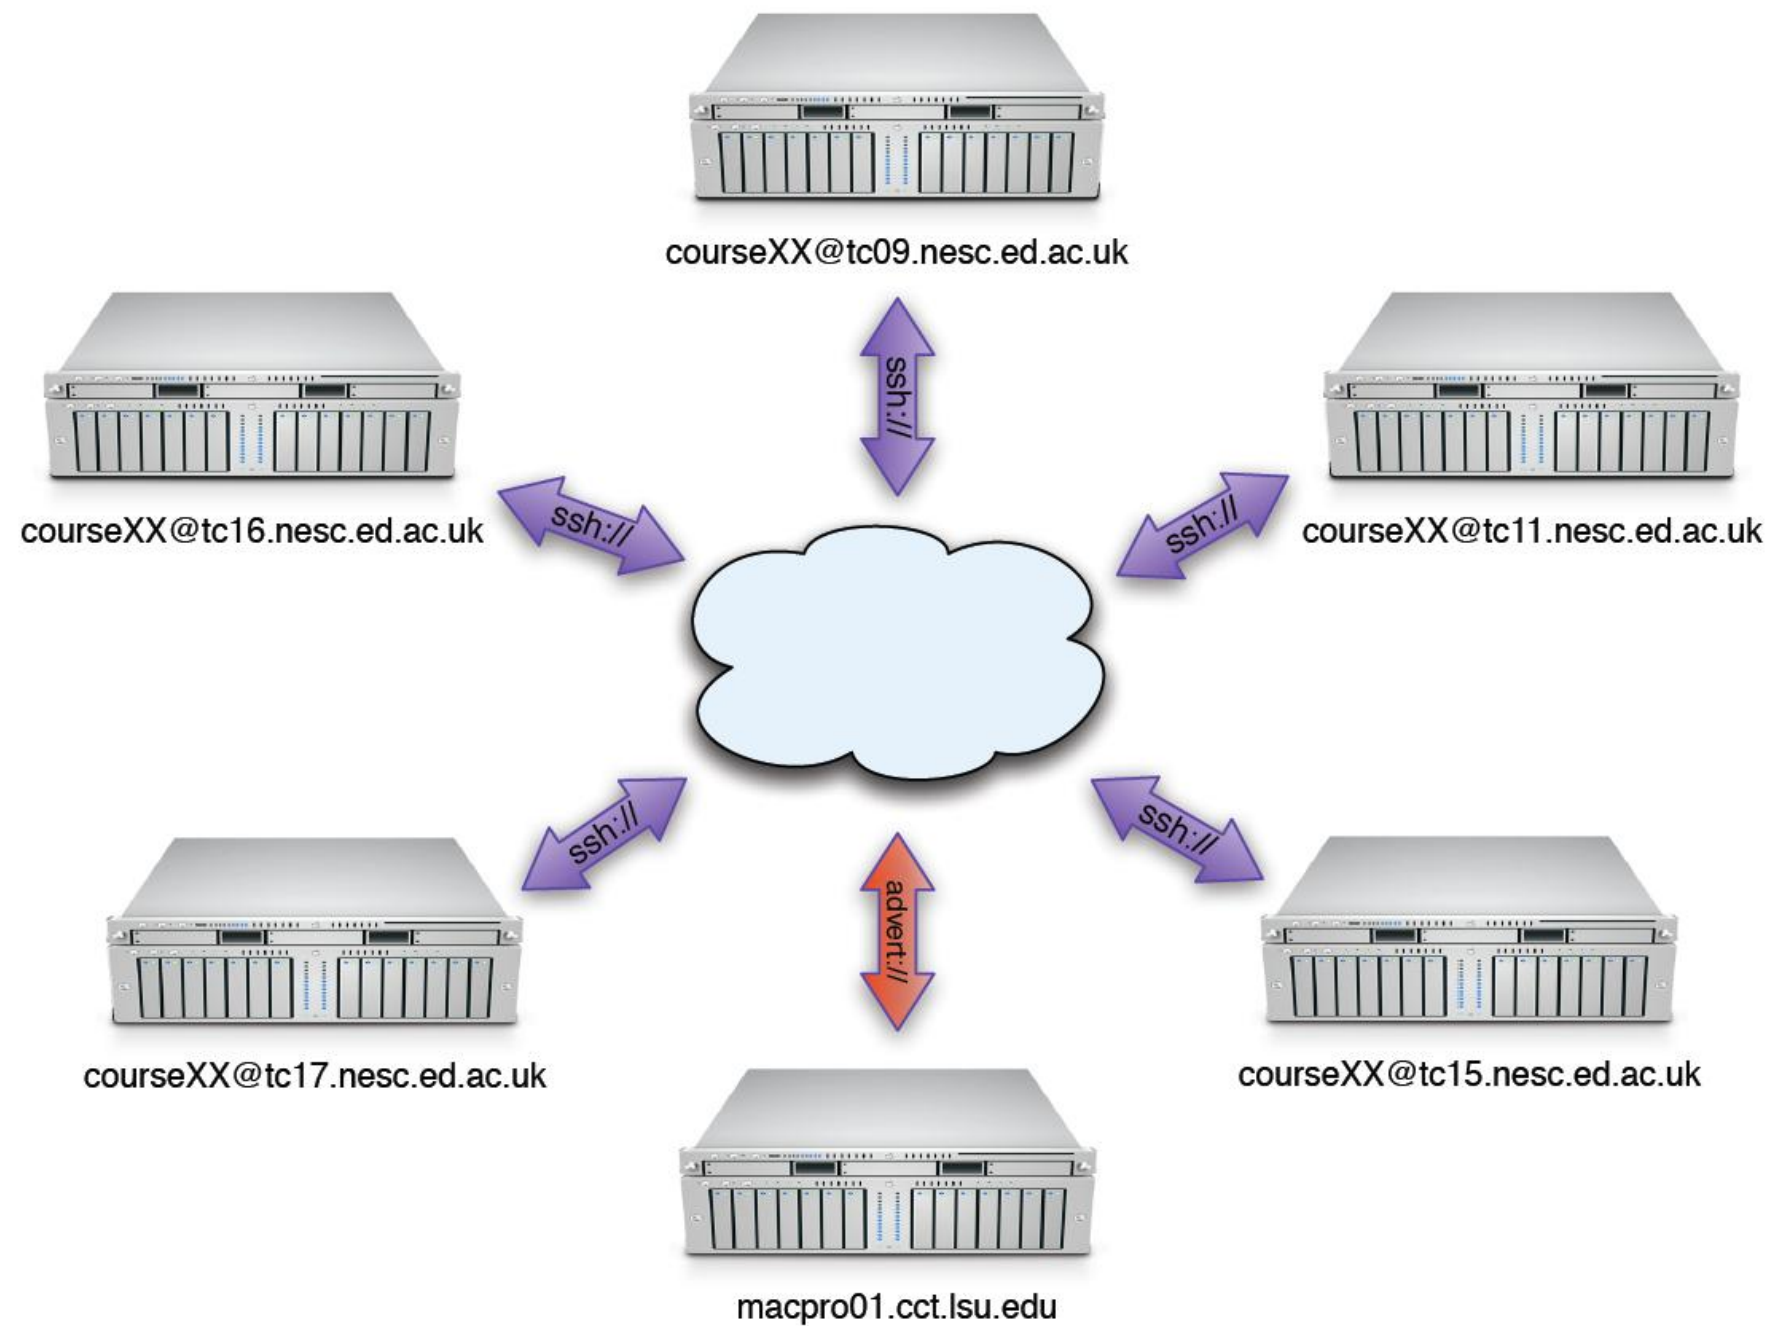
NeSC 2009, September 3rd/4th

Hartmut Kaiser, Shantenu Jha, Ole Weidner, Andre Merzky, Andre Luckow

Agenda

- ▶ Overview, resources, accounts, etc.
- ▶ SAGA command line tools
- ▶ Examples

Infrastructure



Login

- ▶ Login (ssh) to one of the following machines:

tc09.nesc.ed.ac.uk

tc11.nesc.ed.ac.uk

tc15.nesc.ed.ac.uk

tc16.nesc.ed.ac.uk

tc17.nesc.ed.ac.uk

- ▶ Credentials:

course01, course02,..., course12

Enter the SAGA world

```
source /usr/local/saga/share/saga/saga-env.sh
```

Make sure everything works:

```
saga-job run ssh://user@tcXX.nesc.ed.ac.uk /bin/hostname  
saga-advert list_directory advert://FIXME/FIXME
```

Enter the SAGA world

- ▶ If something goes wrong:

```
setenv SAGA_VERBOSE = 0..6
```

- ▶ Will print logging information about adaptors, settings, API calls, etc.

Documentation

- ▶ General information

- ▶ <http://faust.cct.lsu.edu/trac/saga/wiki/NeSC2009>

- ▶ API documentation

- ▶ <http://saga.cct.lsu.edu/cpp/apidoc/>

- ▶ Programming manual

- ▶ <http://tinyurl.com/saga-manual>

Command line tools

- ▶ SAGA comes with simple command line tools that allow to access basic package functionality.
- ▶ The source code is very simple and a great starting point to explore the SAGA package APIs:

- ▶ `saga-file` `$SAGA_ROOT/saga/tools/cltools/file/`
- ▶ `saga-job` `$SAGA_ROOT/saga/tools/cltools/job/`
- ▶ `saga-advert` `$SAGA_ROOT/saga/tools/cltools/advert/`
- ▶ `saga-replica` `$SAGA_ROOT/saga/tools/cltools/replica/`
- ▶ `saga-shell` `$SAGA_ROOT/saga/tools/shell/`

Command line tools

- ▶ ‘Shell bindings’
 - ▶ Package specific (file, job, advert, replica)
- ▶ SAGA shell
 - ▶ All in one solution
 - ▶ Filesystem navigation (filesystem, advert, replica)
 - ▶ Job launching
 - ▶ Scripting

Command line tool: saga-file

- ▶ Supported protocols

- ▶ Depends on SAGA adaptors
- ▶ We will use ssh and local adaptors
- ▶ Also available: Globus GridFTP, Curl (subset), KFS, Amazon EC2, Opencloud (Sector/Sphere), Hadoop (HDFS)

- ▶ Supported commands:

| Command | Arguments |
|----------|---------------------|
| copy | <url from> <url to> |
| move | <url from> <url to> |
| remove | <url> |
| cat | <url> |
| list_dir | <url> |

Command line tool: saga-job

▶ Supported protocols

- ▶ Depends on SAGA adaptors
- ▶ We will use ssh and local adaptors
- ▶ Also available: Globus Gram, Condor, OMII-GridSAM, LSF, Amazon EC2, Opencloud (Sector/Sphere)

▶ Supported commands:

| Command | Arguments |
|---------|--------------------------------|
| run | <rm url> <command> <arguments> |
| submit | <rm url> <command> <arguments> |
| state | <rm url> <jobid> |
| suspend | <rm url> <jobid> |
| resume | <rm url> <jobid> |
| cancel | <rm url> <jobid> |

Command line tool: saga-advert

▶ What is it?

- ▶ Central data store with
 - ▶ Hierarchical keys
 - ▶ Attributes
- ▶ Filesystem like structure

▶ Supported protocols

- ▶ Depends on SAGA adaptors
- ▶ We will use local adaptor
 - ▶ Local backend: SQLite3
 - ▶ Remote backend: PostgreSQL
- ▶ Also available: Hadoop H-Base, Hypertable

Command line tool: saga-advert

► Supported commands:

| Command | Arguments |
|-----------------------------------|----------------------------|
| list_directory | <advert-url> <pattern> |
| add_directory remove_directory | <advert-url> |
| add_entry remove_entry | <advert-url> |
| store_string | <advert-url> <string> |
| retrieve_string | <advert-url> |
| list_attributes | <advert-url> |
| set_attribute | <advert-url> <key> <value> |
| remove_attribute | <advert-url> <key> |

Command line tool: saga-replica

- ▶ What is it?
 - ▶ Central data store allowing to map logical file names to a set of physical files (i.e. different instances of same file on different machines)
 - ▶ Hierarchical keys
 - ▶ Attributes
 - ▶ Filesystem like structure
- ▶ Supported protocols
 - ▶ Depends on SAGA adaptors
 - ▶ We will use local adaptor
 - ▶ Local backend: SQLite3
 - ▶ Remote backend: PostgreSQL
 - ▶ Also available: Globus RLS (subset)

Command line tool: saga-replica

► Supported commands:

| Command | Arguments |
|-----------------------------------|------------------------|
| list_directory | <lfname> <pattern> |
| add_directory remove_directory | <lfname> |
| add_lfn remove_lfn list_pfn | <lfname> |
| add_pfn | <lfname> <pfn> |
| remove_pfn | <lfname> <pfn> |
| list_attributes | <lfname> |
| set_attribute | <lfname> <key> <value> |
| remove_attribute | <lfname> <key> |

Command line tool: saga-shell

- ▶ All in one of all command line tools as mentioned earlier
- ▶ Keeps context in between commands
- ▶ Navigate (remote) filesystems (advert, replica too!)
- ▶ Launch (remote) jobs, uses io redirection to access in/out
- ▶ All commands are implemented using SAGA

Command line tool: saga-shell

► Some of the supported commands

| Type | Commands |
|------------------------|--|
| File system navigation | pwd, ls, mv, cp, cd, mkdir, rmdir, touch, cat |
| Job package | run, suspend, resume, kill, status, ps |
| replica | rep_find, rep_list, rep_add, rep_remove, rep_update, rep_replicate |
| environment | setenv, getenv, env |
| permissions | add_proxy, remove_proxy |

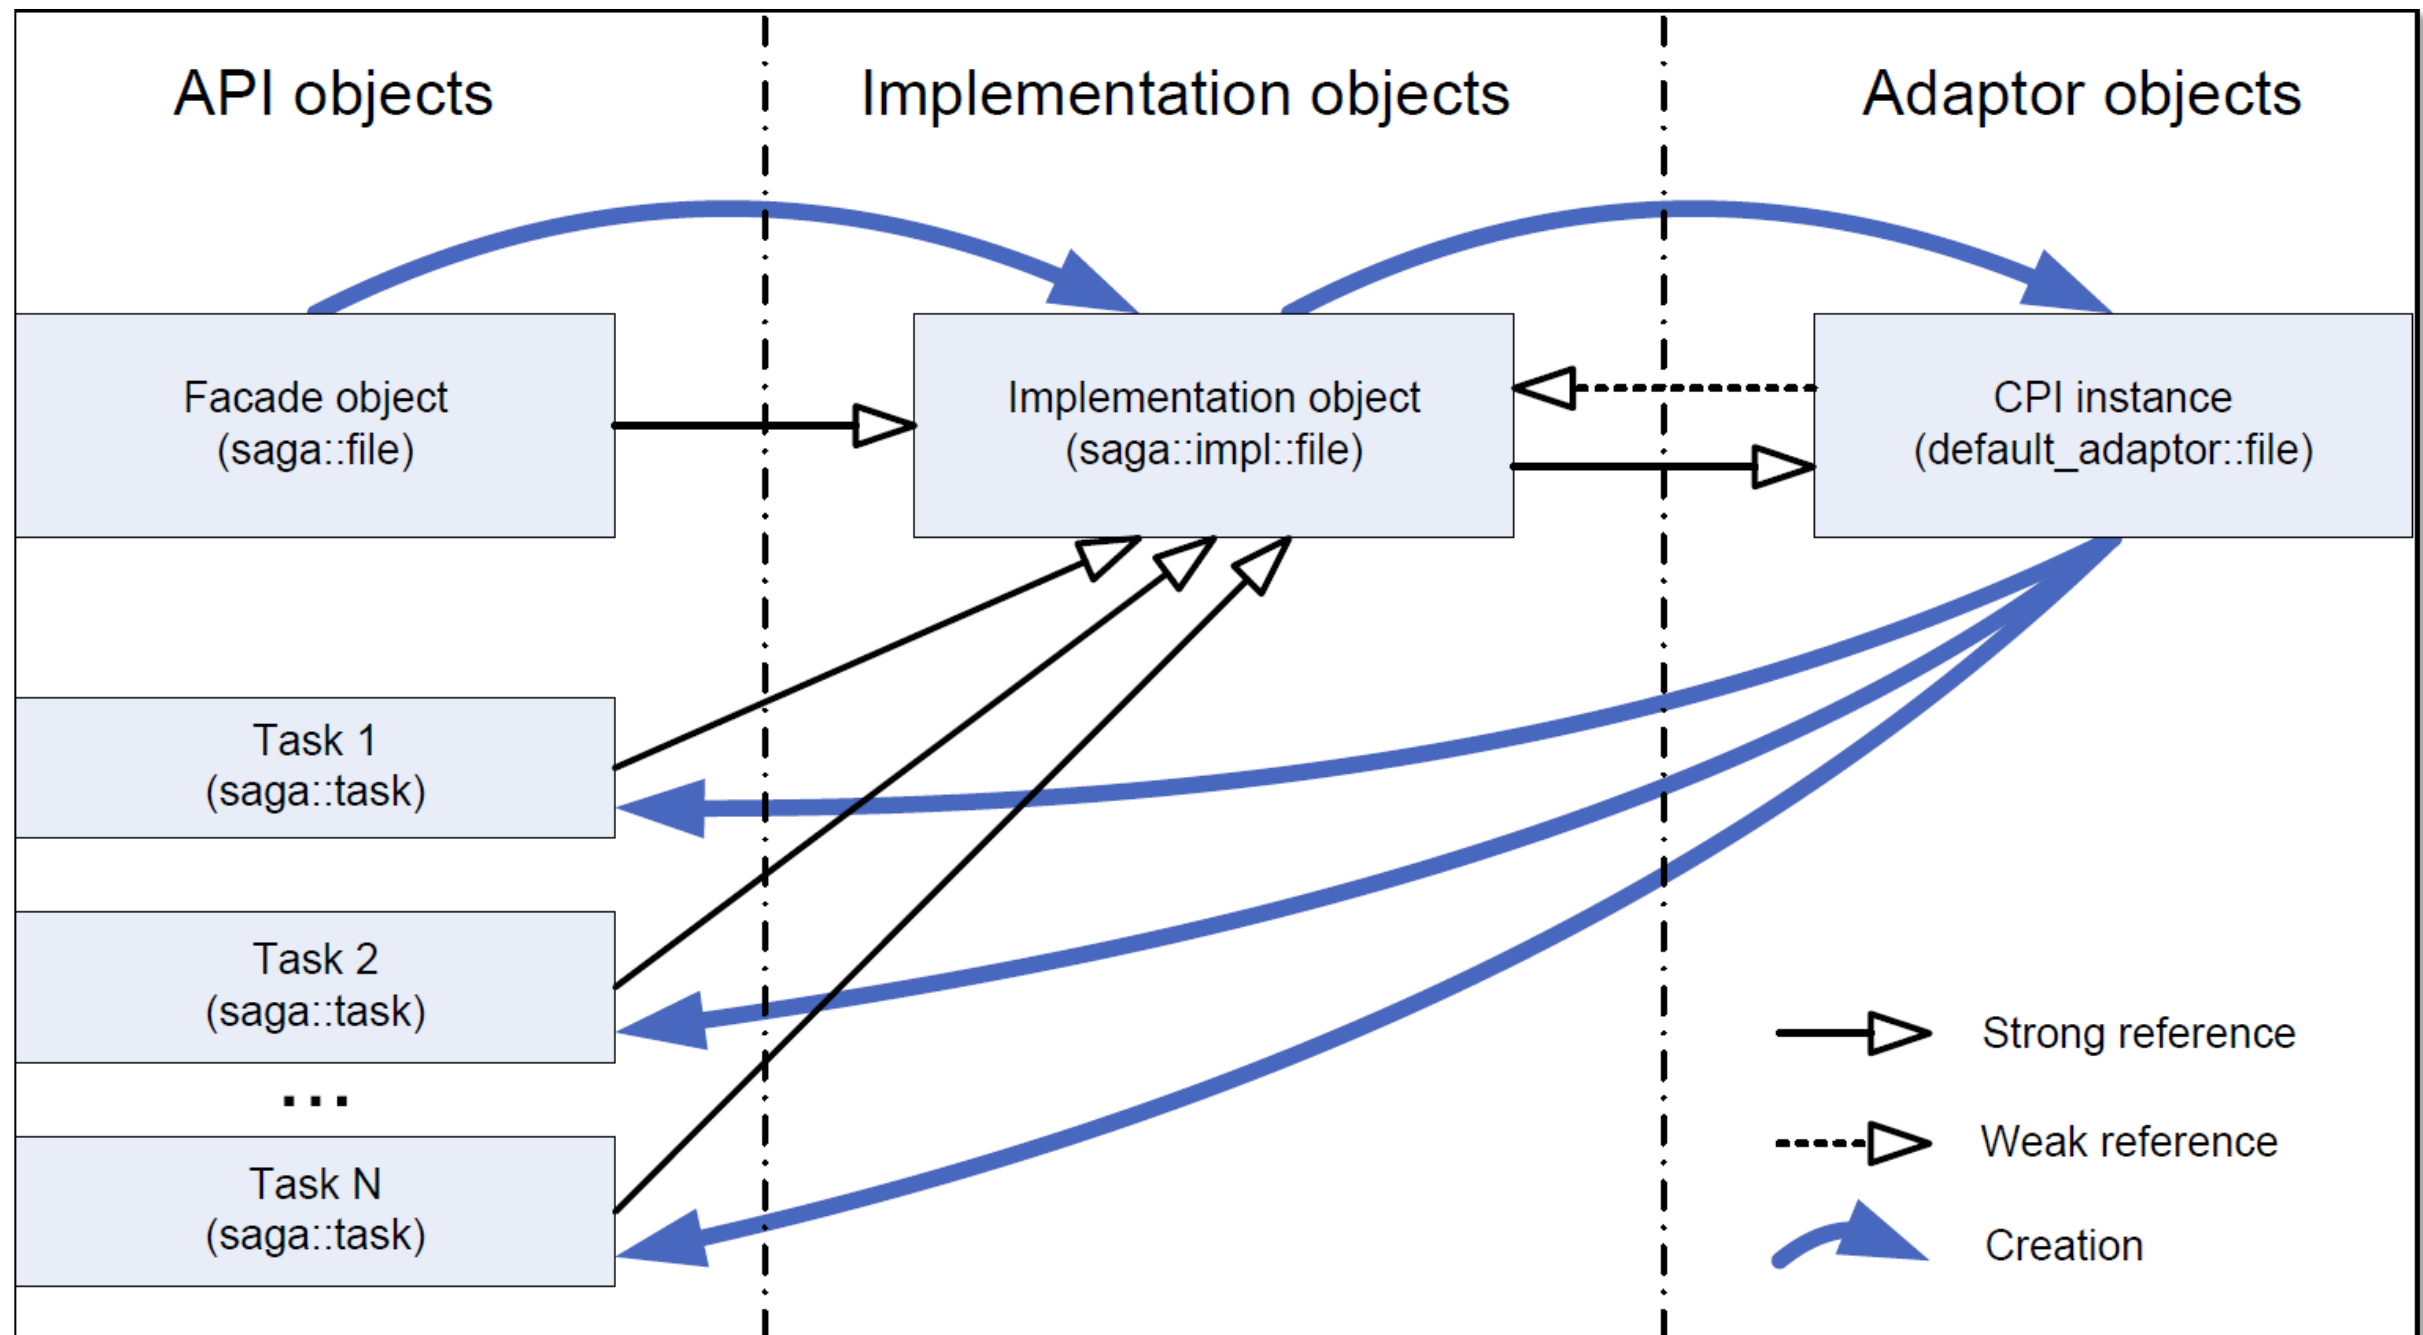
Hands on

- ▶ Try and run command line tools
- ▶ Copy a file, move it, delete it
- ▶ Run a job (/bin/sleep 10), monitor its status
- ▶ Play with advert service, create, directories, entries, store date, attributes

General Guidelines

- ▶ Pimpl paradigm, shared_ptr
- ▶ Sync/async API's
- ▶ Task container
- ▶ Error handling

Pimpl paradigm, shared_ptr



Pimpl paradigm, shared_ptr

- ▶ All SAGA API objects are very lightweight (except `saga::url`)
 - ▶ Cheap to copy, passed as arguments, or stored in containers

```
#include <saga/saga.hpp>

void copy(saga::filesystem::file f, saga::url const& target)
{
    f.copy(target);
}

int main(int argc, char* argv[])
{
    saga::url source("ssh://tc17/etc/passwd");
    saga::url target(".");
    saga::filesystem::file f(source, saga::filesystem::Read);

    copy(f, target);
    return 0;
}
```

Sync/Async API's

- ▶ Almost all API objects expose 4 different sets of API functions:
 - ▶ Synchronous
 - ▶ Returns result synchronously
 - ▶ `saga::off_t file::get_size();`
 - ▶ Task based
 - ▶ Returns handle to deferred result (using a task object)
 - ▶ Asynchronous
 - `saga::task` not running yet
 - `saga::task file::get_size<saga::task::Async>();`
 - ▶ Task
 - `saga::task` already running
 - `saga::task file::get_size<saga::task::Task>();`
 - ▶ Synchronous
 - `saga::task` guaranteed to be finished
 - `saga::task file::get_size<saga::task::Sync>();`

Sync/Async API's

► Synchronous

```
saga::url source("ssh://tc17/etc/passwd");  
saga::filesystem::file f(source, saga::filesystem::Read);  
  
saga::off_t size = f.get_size();
```

► Asynchronous

```
saga::url source("ssh://tc17/etc/passwd");  
saga::filesystem::file f(source, saga::filesystem::Read);  
  
saga::task t = f.get_size<saga::task::ASync>();  
t.run();  
  
saga::off_t size = t.get_result<saga::off_t>();
```

Sync/Async API's

- ▶ Asynchronous creation of objects
- ▶ Factory functions

```
// create task
saga::task t =
    saga::filesystem::file::create<saga::task::ASync>("fileurl");

// ...
saga::filesystem::file f =
    t.get_result<saga::filesystem::file>();
```


Task container

- ▶ Special container allowing to handle many tasks as one

```
// create a list of tasks to run jobs
saga::job::service js("somehost");
saga::task_container tc;
for (int i = 0; i < num; ++i)
{
    tc.add_task(js.run_job<Task>(execs[i], hosts[i]));
}

// execute tasks
tc.run();

// wait for any of the tasks to finish
std::vector<saga::task> finished =
    tc.wait(saga::task_container::Any));
```

Error Handling

- ▶ Error handling is currently purely exception based

```
try {  
    saga::filesystem::file f("non-existing file");  
    // ...  
}  
catch (saga::exception const& e) {  
    std::cerr << e.what() << std::endl;  
}
```

- ▶ Also:

```
e.get_error();           // error code  
e.get_message();         // top level message  
e.get_object();          // get failing API object  
e.get_all_exceptions();  // get list of exceptions  
e.get_all_messages();    // get list of messages
```

A Simple SAGA Application

► Simple file copy example: copy.cpp

```
#include <saga/saga.hpp>

int main(int argc, char* argv[])
{
    saga::url source("ssh://tc17/etc/passwd");
    saga::url target(".");
    saga::filesystem::file f(source, saga::filesystem::Read);

    f.copy(target);
    return 0;
}
```

Compiling and Linking a SAGA Application

► Simple file copy example: Makefile

```
SAGA_SRC = $(wildcard *.cpp)
SAGA_ADD_BIN_OBJ = $(SAGA_SRC:%.cpp=%.o)
SAGA_BIN = copy

include $(SAGA_LOCATION)/share/saga/make/saga.application.mk
```

► Use saga-config

```
g++ -Wall `saga-config --cxxflags` `saga-config --lflags` copy.cpp
```

► Or do it the hard way:

```
g++ -Wall -I$SAGA_LOCATION/include -pthread \
-L$SAGA_LOCATION/lib \
-lsaga_engine -lsaga_package_file copy.cpp
```

Running a SAGA Application

- ▶ Make sure that the SAGA libraries can be found by the loader.
- ▶ Use: `/usr/local/saga/share/saga/saga-env.sh`
- ▶ If something goes wrong - use `SAGA_VERBOSE`:

```
SAGA_VERBOSE=6 ./copy
```

Programmers Guide

- ▶ Set of very small and easy examples, one for each package/paradigm
 - ▶ file_copy, file_copy (async)
 - ▶ Error handling
 - ▶ Attributes
 - ▶ Stream (server/client)

Hands on

- ▶ Try compiling and running other examples
 - ▶ Urls
 - ▶ Packages: file, job, replica
 - ▶ Modify stream server/client (see next slide)

Hands on

- ▶ Modify stream_server and stream_client to establish connection using advert service

- ▶ Server:

```
saga::stream::server s("tcp://tc17");  
saga::advert::entry e  
    ("advert://macpro01.cct.lsu.edu/NeSC2009/exercise_3/<uid>");  
e.store_object(s);  
...
```

- ▶ Client:

```
saga::advert::entry e  
    ("advert://macpro01.cct.lsu.edu/NeSC2009/exercise_3/<uid>");  
saga::stream::stream c (e.retrieve_object());  
c.connect();  
...
```


Example 1: hello_world

▶ Hello world

- ▶ Launch 3 jobs on different machines
 - ▶ Execute `"/bin/echo"`
- ▶ No job dependency
- ▶ Each job returns its passed input argument
 - ▶ "Hello"
 - ▶ "distributed"
 - ▶ "world!"
- ▶ Jobs are launched in parallel (in separate threads)
- ▶ As soon as result is collected it's printed on local console

Example 1: hello_world

- ▶ Hello world
 - ▶ Arbitrary sequence of results
 - ▶ Optimally: "Hello distributed world!"
 - ▶ Demonstrates
 - ▶ How to launch a remote job using SAGA job_service
 - ▶ Pass arguments using the command line
 - ▶ Collect result by output redirection
- ▶ The source code can be found here (see 'Example1'):
 - ▶ <http://faust.cct.lsu.edu/trac/saga/wiki/NeSC2009>
 - ▶ The example uses localhost to spawn childs
 - ▶ For remote execution change HOST1, HOST2, HOST3 from "localhost" to "[tc11, tc15, tc16, or tc17].nesc.ed.ac.uk"

Hands on

- ▶ Compile and run example locally
- ▶ Modify the code to run it remotely
- ▶ Compile and run example remotely
- ▶ Run other remote executables
- ▶ Modify hello_world to produce consistent results
 - ▶ “Hello distributed world!”
- ▶ ...

Example 2: chaining_jobs

- ▶ Launch 3 jobs on 3 different machines
- ▶ Output of previous job is needed to launch next job
- ▶ Simple sequential execution, but SAGA style
- ▶ Demonstrates
 - ▶ How to launch a job using SAGA job_service
 - ▶ How to feed input to launched job
 - ▶ How to collect output
- ▶ Launched job: `/usr/bin/bc`
- ▶ Increment the number passed as the argument
 - ▶ Pass returned incremented number to next job

Example 2: chaining_jobs

- ▶ Pass input
 - ▶ Command line (same as before)
 - ▶ Stage in of input files
 - ▶ Using job package API
 - ▶ Using file package API
 - ▶ Input redirection
 - ▶ Using Stream API
 - ▶ Using Advert service
- ▶ Collect output
 - ▶ Output redirection
 - ▶ Stage out of output files
 - ▶ Using job package API
 - ▶ Using file package API
 - ▶ Using Stream API
 - ▶ Using Advert service

Hands on

- ▶ **Coordinate execution sequence of multiple jobs**
 - ▶ 3 jobs are running, but wait for a flag in advert service to be set by previous job to continue

Example 3: depending_jobs

- ▶ Coordinating information from advert service
- ▶ Launch a single job sequentially on a set of remote resources
 - ▶ Simulating checkpointing/relaunching on different resource (migration)
- ▶ Maintain a single result value in advert service
 - ▶ Gets written by one job, and read by the next
- ▶ Demonstrates
 - ▶ How to launch remote job using SAGA job, while maintaining environment
 - ▶ Assembling argument lists
- ▶ Result is left in advert service, but accessed afterwards

Thanks! Questions?

<http://saga.cct.lsu.edu>