

# SAGA Tutorial - NeSC 2009

## Introduction to the SAGA API

# Agenda

- Background Recap
- SAGA API structure and scope
- walkthrough
- some implementation details
- SAGA extensions

- Middleware often targets legacy applications (Unicore, Globus, Condor, ...)
- some are distribution aware (MPICH-G, Ninf-G, ...)
- few **APIs** exist for Grid aware applications
  - GridFTP
  - GRAM
  - gLite
  - CoG
  - GAT
  - Cloud APIs

- diversity of Grid Middleware implies diversity of APIs
- some APIs try to generalize Grid programming concepts
- difficult to keep up with MW development, and to stay **simple**

- Open Grid Forum (GF, EGF, GGF) tries to standardize Grid MW
- e.g. single job description language (JSDL)
- focuses on interfaces, but also protocol and architecture

- OGF focuses on services
- some effort on higher level APIs
  - Distributed Resource Management Application API (DRMAA)
  - Remote Procedure Calls (GridRPC)
  - Checkpoint and Recovery (GridCPR)
  - Job Submission and Description Language (JSDL)
- numerous service interfaces, often WS-based (WSRF)

- implementable on all major resource management services
- simple means to define jobs, and to submit them
- basic job management features (status, kill)
- job templates for bulk job management

# DRMAA Example

## DRMAA Job Submit

```
drmaa_job_template_t * job_template;

if ( ! ( job_template = create_job_template (exe, 5, 0) ) )
{
    fprintf (stderr, "create_job_template failed\n");
    return 1;
}

while ( ( drmaa_errno = drmaa_run_job (job_id,
                                      sizeof (jobid)-1,
                                      job_template,
                                      diagnosis,
                                      sizeof (diagnosis)-1)
        ) == DRMAA_ERRNO_DRM_COMMUNICATION_FAILURE )
{
    fprintf(stderr, "drmaa_run_job failed: %s\n", diagnosis);
    sleep (1);
}
```



- 'standardizes' the three existing RPC implementations for Grids
- example of '*gridified API*'
- simple: get function handle, call function
- explicit support for async rpc calls

## GridRPC: Matrix Multiplication

```
double A[N*N], B[N*N], C[N*N];

initMatA (N, A);
initMatB (N, B);

grpc_initialize (argv[1]);

grpc_function_handle_t handle;
grpc_function_handle_default (&handle, "mat_mult");

if ( grpc_call (&handle, N, A, B, C) != GRPC_NO_ERROR)
{
    exit (1);
}

grpc_function_handle_destruct (&handle);
grpc_finalize ();
```

- Grids seem to favour application level checkpointing
- GridCPR allows to manage checkpoints
- defines an architecture, service interfaces, and, aehem, no API

- extensible XML based language for describing job requirements
- does not cover resource description (on purpose)
- does not cover workflows, or job dependencies etc (on purpose)
- JSDL is extensible (ParameterSweep, SPMD)

## JSDL: Simple Job

```
<jsd1:JobDefinition>
  <JobDescription>
    <Application>
      <jsd1-posix:POSIXApplication>
        <Executable>/bin/date</Executable>
      </jsdl-posix:POSIXApplication>
    </Application>
    <Resources ...>
      <OperatingSystem>
        <OperatingSystemType>
          <OperatingSystemName>LINUX</OperatingSystemName>
        </OperatingSystemType>
      </OperatingSystem>
    </Resources>
  </JobDescription>
</jsdl:JobDefinition>
```

- XML: embeddable into WSRF (WS-Agreement etc.)
- XML, but relatively flat
- maps well to existing JDLs, but is 'more complete'
- extensible (resource description, job dependencies, workflow)
- top down approach!

- some APIs exist in OGF, and are successful
- OGF APIs do not cover the complete OGF scope
- the various API standards are disjunct
- WSDL as service interface specification cannot replace an application level API (wrong level of abstraction)
- **SAGA tries to address these issues**

# OGF: top-down vs. bottom-up

- bottom-up often agrees on (semantic) LCD + backend specific extensions
- top-down usually focuses on semantics of application requirements
- bottom-up tends to be more powerful
- top-down tends to be simpler and more concise
- *we very much prefer top-down!*



# SAGA

## Simple API for Grid Applications

# SAGA Design Principles

- **SAGA: Simple API for Grid Applications**
- OGF approach to a uniform API layer (facade)
- **governing principle: 80:20 rule**  
simplicity versus control!
- **top-down approach:** use case driven!  
→ defines **application level** abstractions
- **extensible:** stable look & feel + API packages
- **influenced by:** DRMAA, GridRPC, OREP, JSDL, POSIX, GAT, CoG, LSF, Globus, ...
- API Specification is **Language Independent (IDL)**  
Renderings exist in C++, Python, Java  
Examples here are in C++

# SAGA Intro: Example 1

## SAGA: File Management

```
saga::filesystem::directory dir ("any://remote.host.net//data/");  
  
if ( dir.exists ("a") && ! dir.is_dir ("a") )  
{  
    dir.copy ("a", "b", Overwrite);  
}  
  
list <saga::url> names = dir.find ("*-{123}.txt");  
  
saga::filesystem::directory tmp  = dir.open_dir ("tmp/", Create);  
saga::filesystem::file      file = dir.open      ("tmp/data.txt");
```

# SAGA Intro: Example 1

- API is clearly POSIX (libc + shell) inspired
- where is my security??
- what is 'any: / /' ???

# SAGA Intro: Example 2

## SAGA: Job Submission

```
saga::job::description jd; // details left out
saga::job::service      js ("any://remote.host.net/");
saga::job::job          j = js.create_job (jd);

j.run ();

cout << "Job State: " << j.get_state () << endl;

j.wait ();

cout << "Retval " << j.get_attribute ("ExitCode") << endl;
```

# SAGA Intro: Example 2'

## SAGA: Job Submission

```
saga::job::service      js ("any://remote.host.net");  
saga::job::job          j = js.run_job ("touch /tmp/touch.me");  
  
cout << "Job State: " << j.get_state () << endl;  
  
j.wait ();  
  
cout << "Retval " << j.get_attribute ("ExitCode") << endl;
```

# SAGA Intro: Example 2

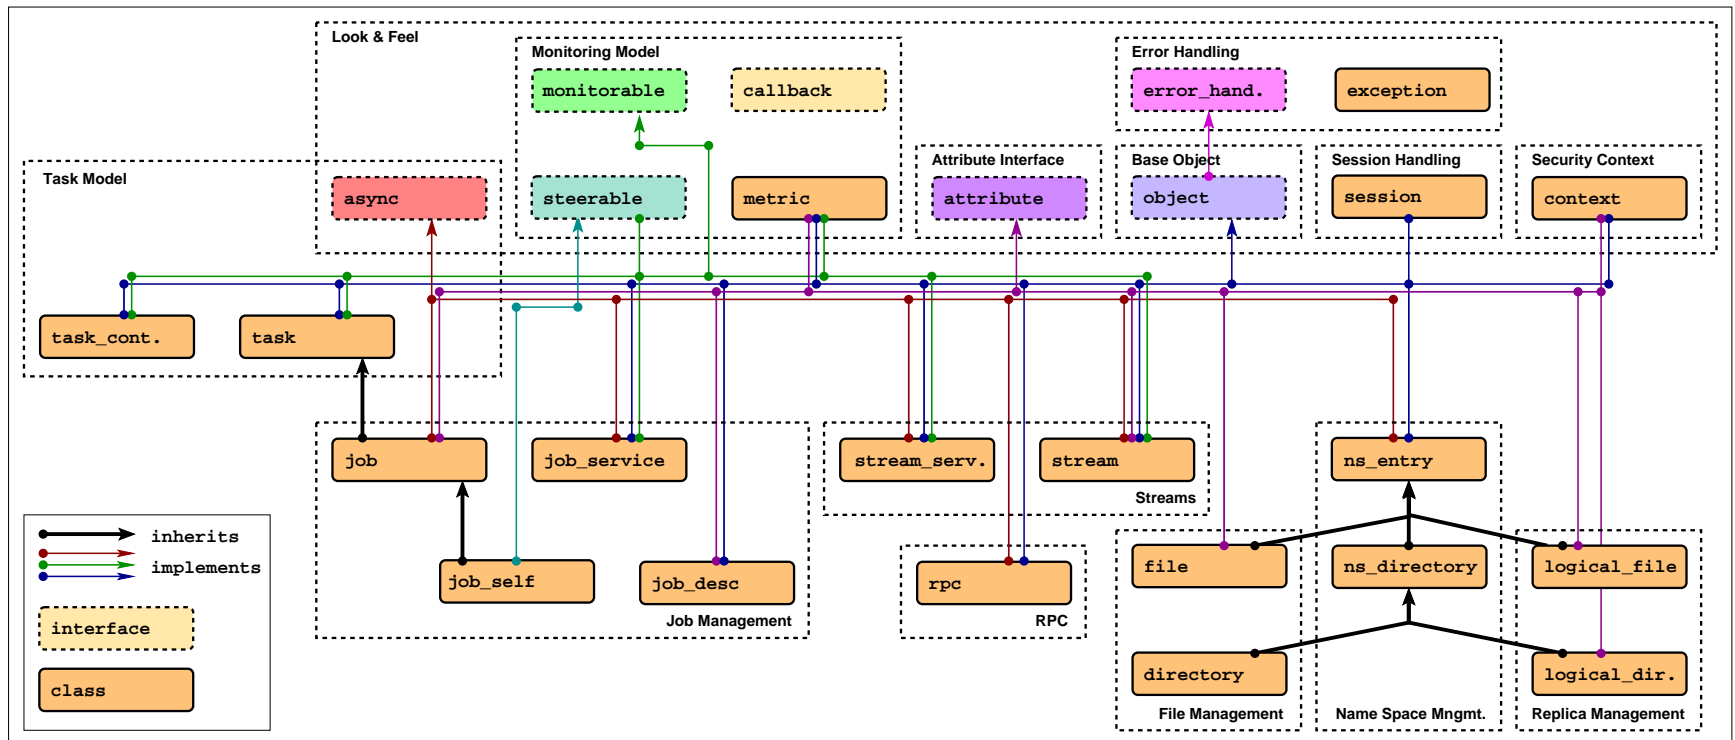
- stateful objects!
- yet another job description language? : – (
- many hidden/default parameters (to keep call signatures small)
- 'any : / /' again!
- TIMTOWTDI (there is more than one way to do it)

# SAGA Intro: 10.000 feet

- **object oriented:** inheritance, interfaces  
very moderate use of templates though!
- functional and non-functional elements strictly separated
  - *non-functional API:* look & feel- orthogonal to functional API  
often not mappable to remote operations
  - *functional API:* API 'Packages' - extensible  
typically mappable to remote operations
- few inter-package dependencies - allows for partial implementations




# SAGA: Class hierarchy

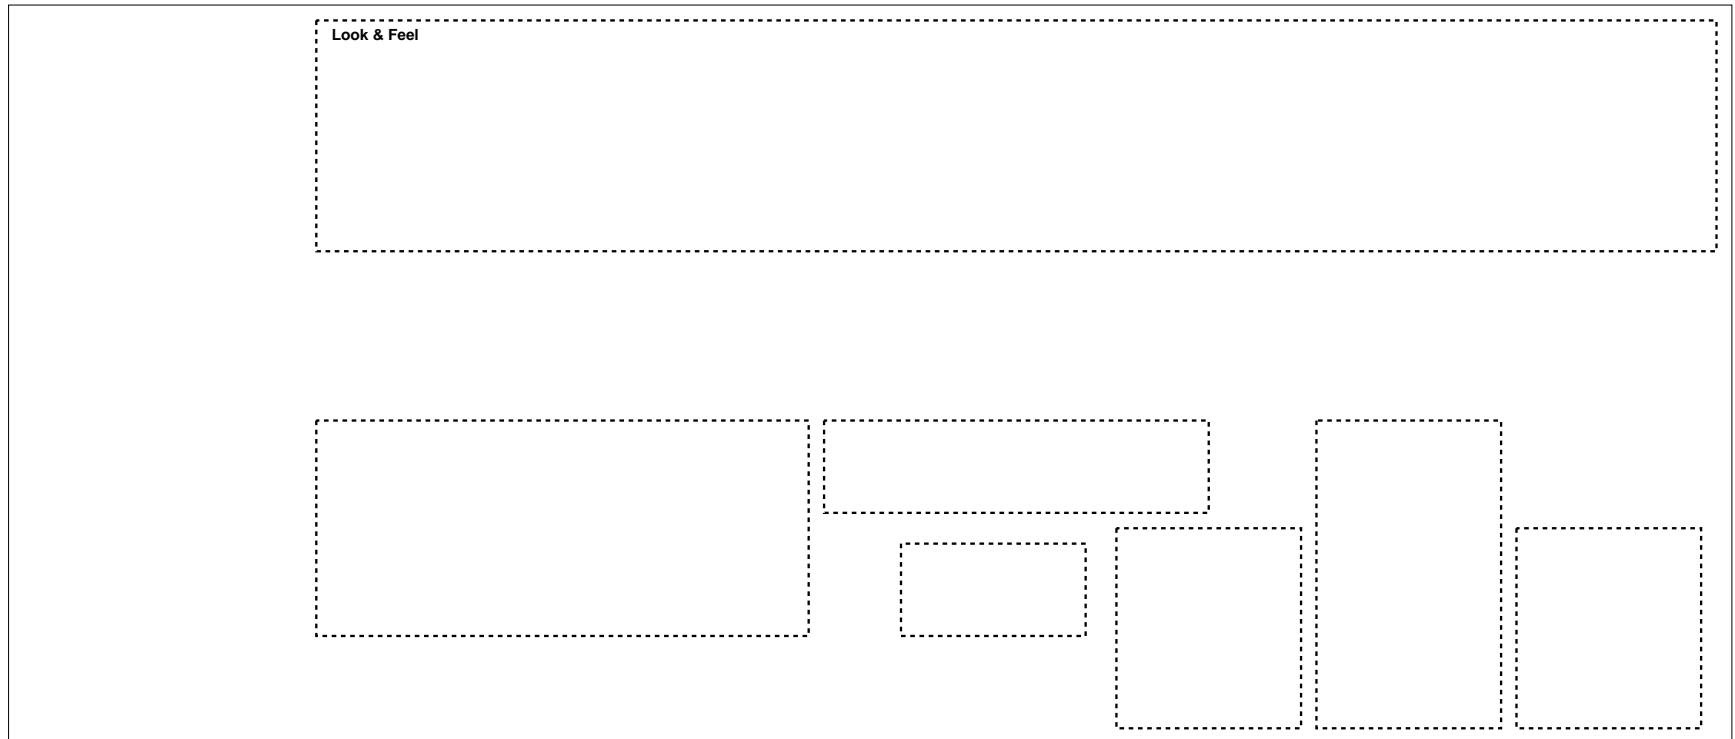


**SAGA API Structure:** look & feel (top) + API packages (bottom)

# Implementation



# SAGA: Class hierarchy



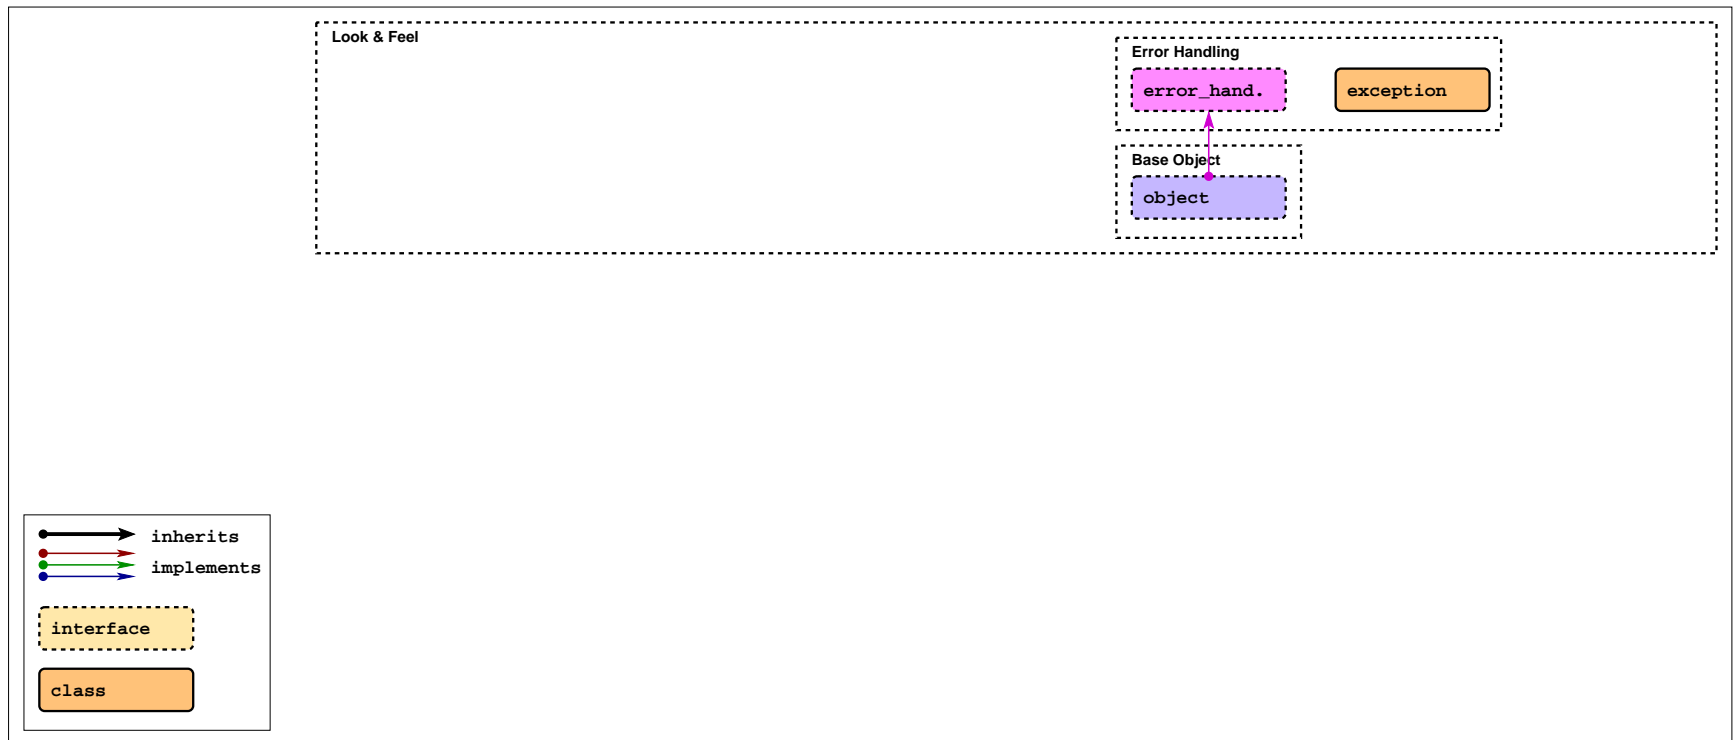
# SAGA: Class hierarchy



## SAGA Look & Feel:

`saga::object` allows for object uuids, `clone()` etc.

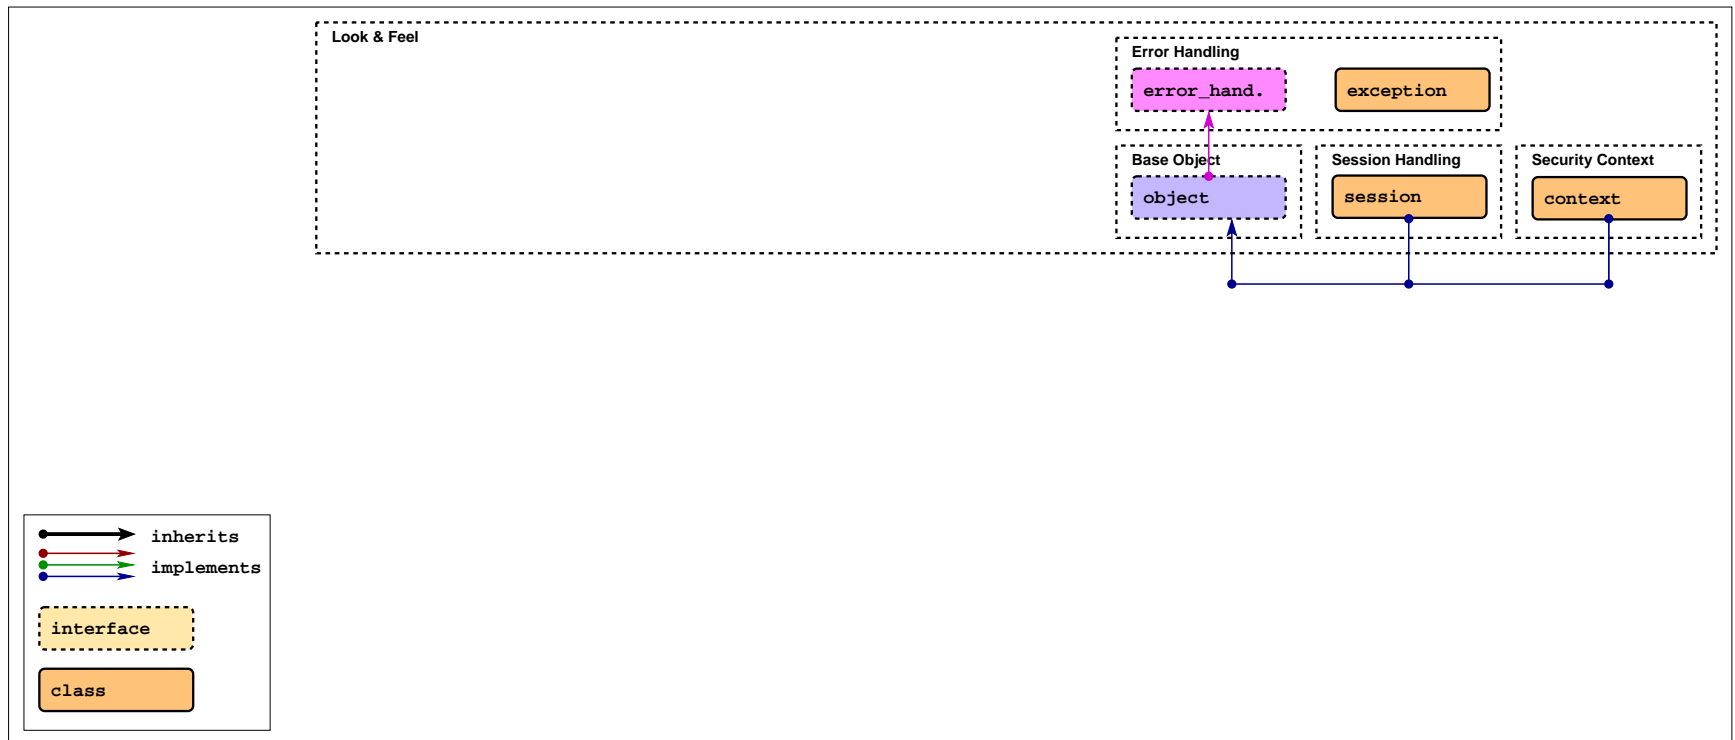
# SAGA: Class hierarchy



## SAGA Look & Feel:

errors are based on exceptions or error codes.

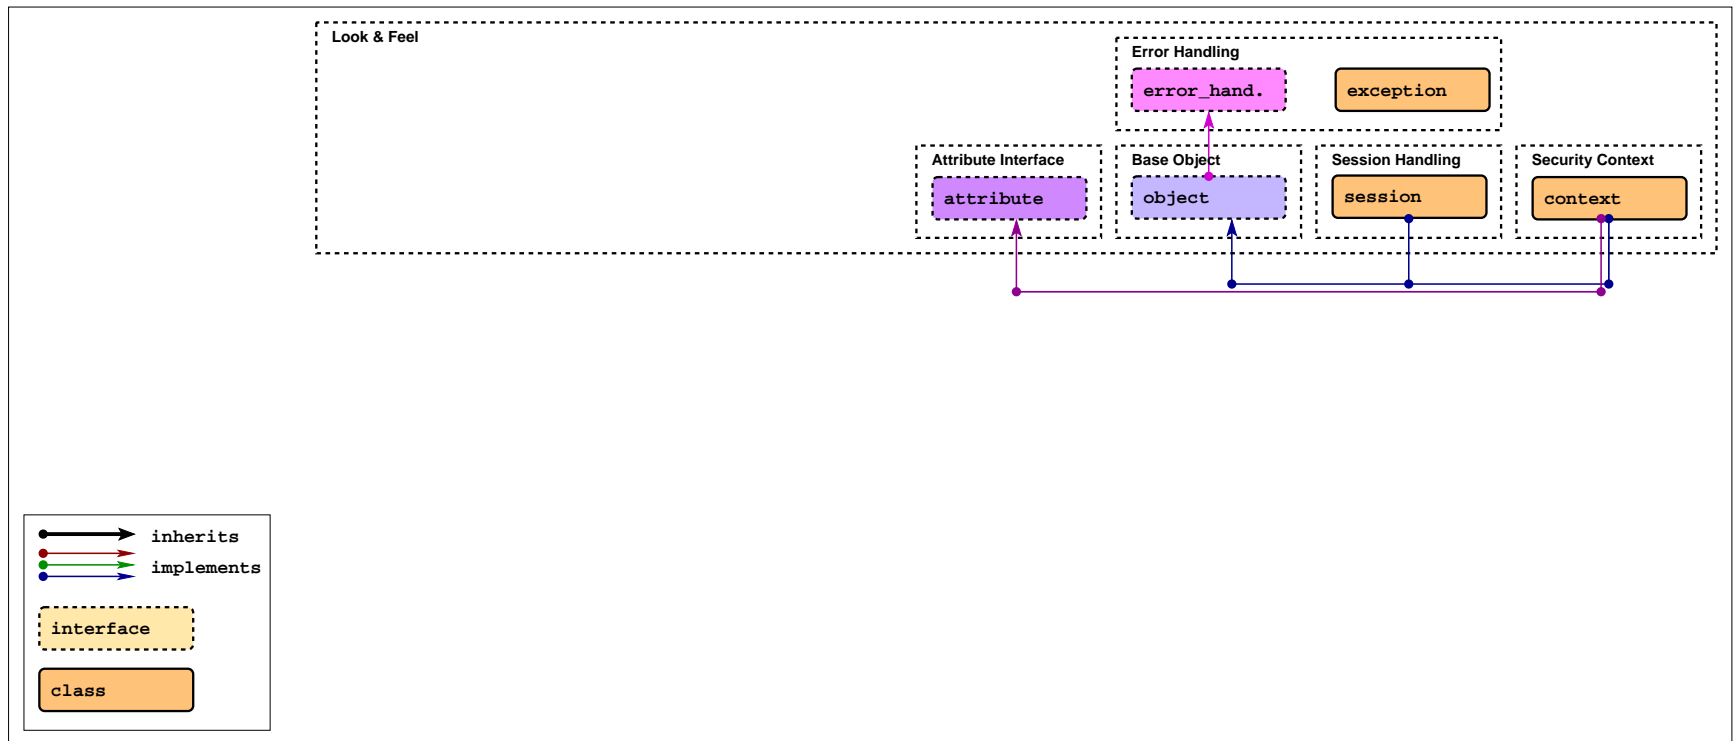
# SAGA: Class hierarchy



## SAGA Look & Feel:

session and credential management is hidden.

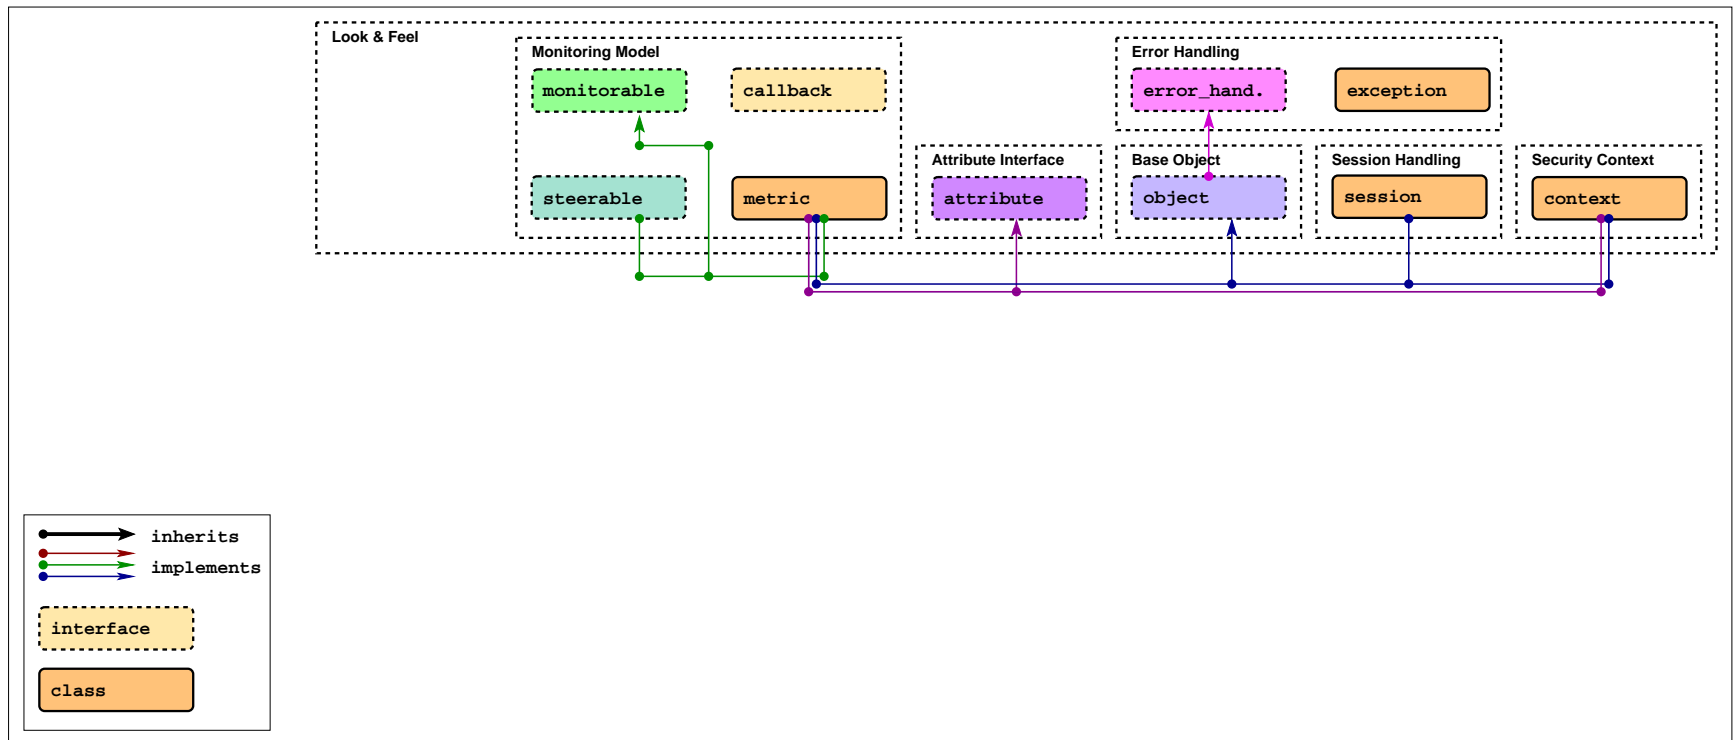
# SAGA: Class hierarchy



## SAGA Look & Feel:

Attribute interface for meta data.

# SAGA: Class hierarchy

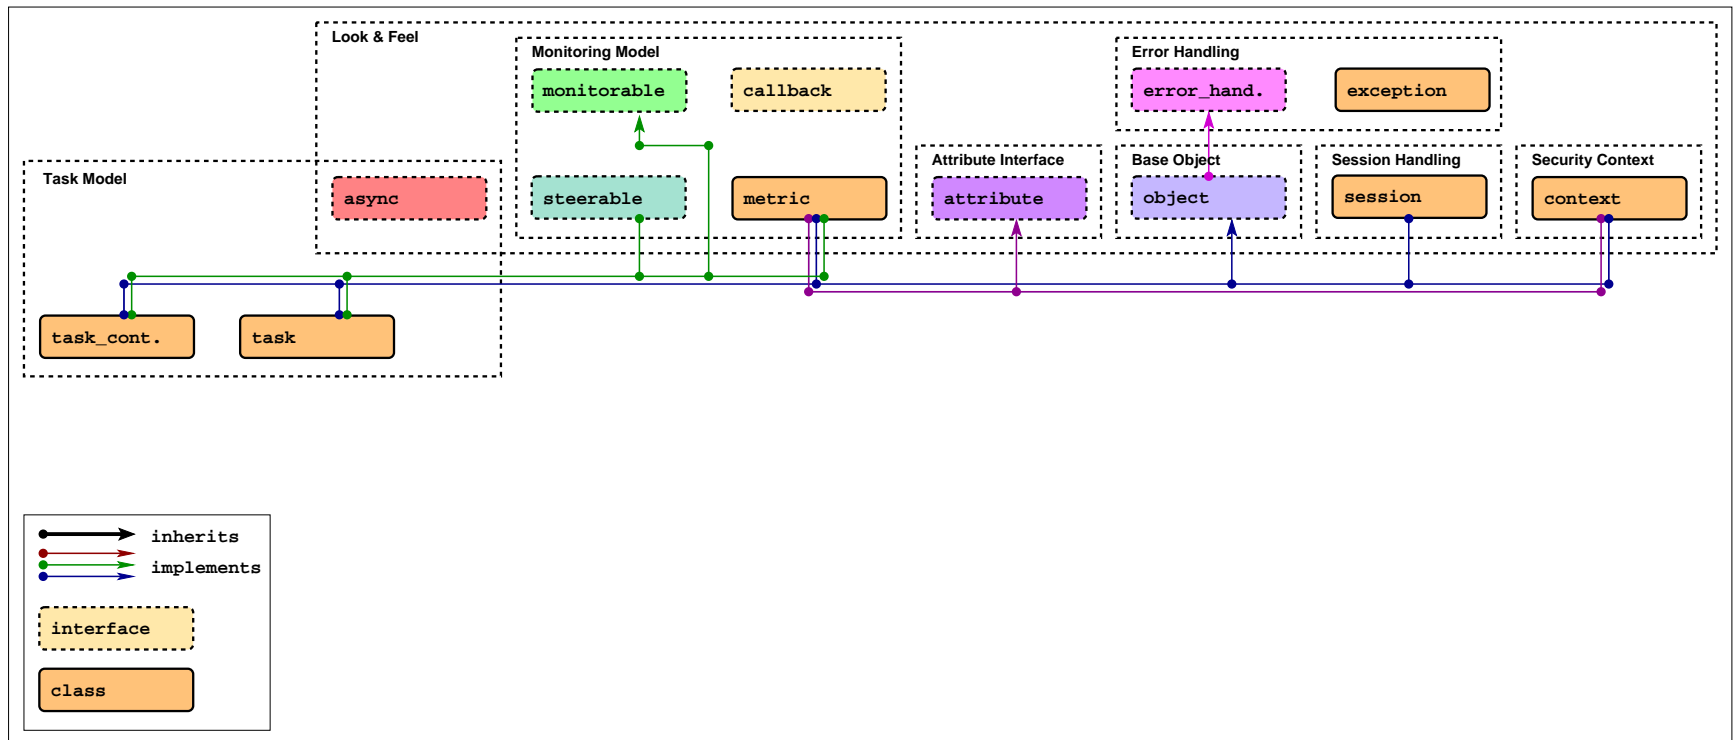


## SAGA Look & Feel:

Monitoring includes asynchronous notifications.



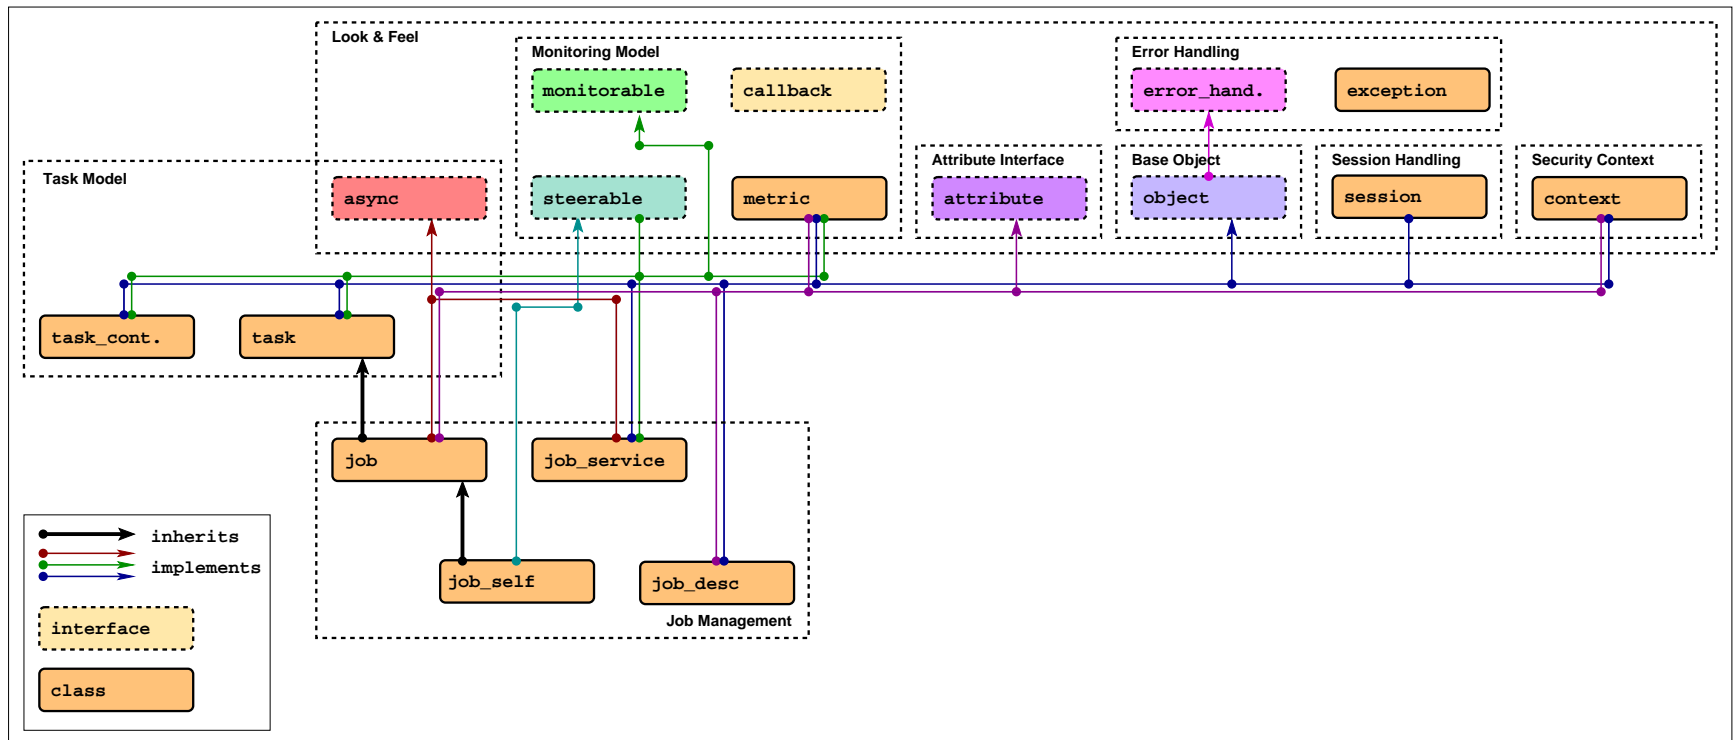
# SAGA: Class hierarchy



## SAGA Look & Feel:

the task model adds asynchronous operations.

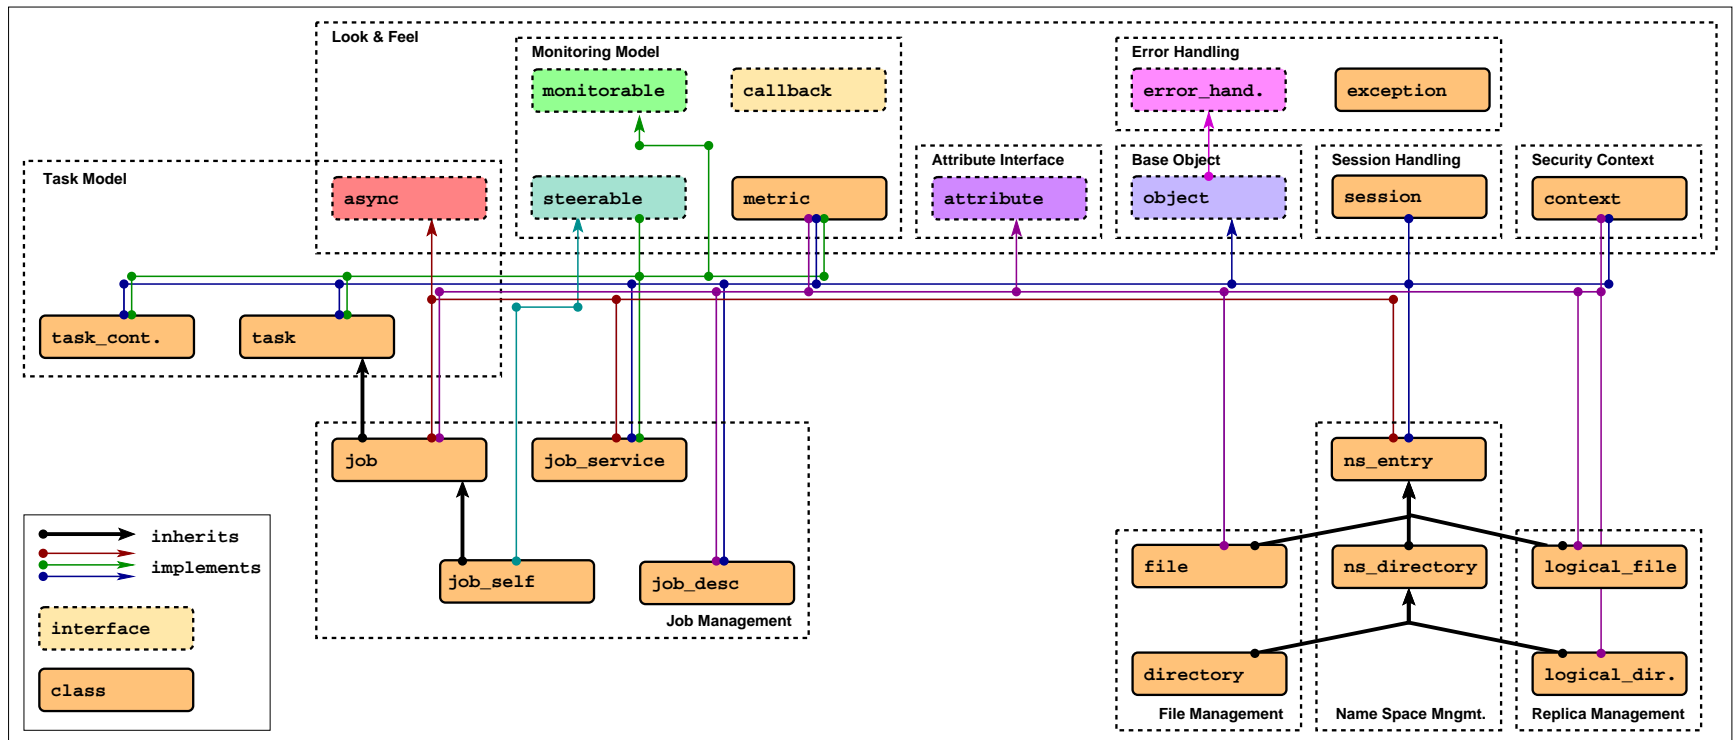
# SAGA: Class hierarchy



## SAGA API Package 'job':

create and manage remote processes.

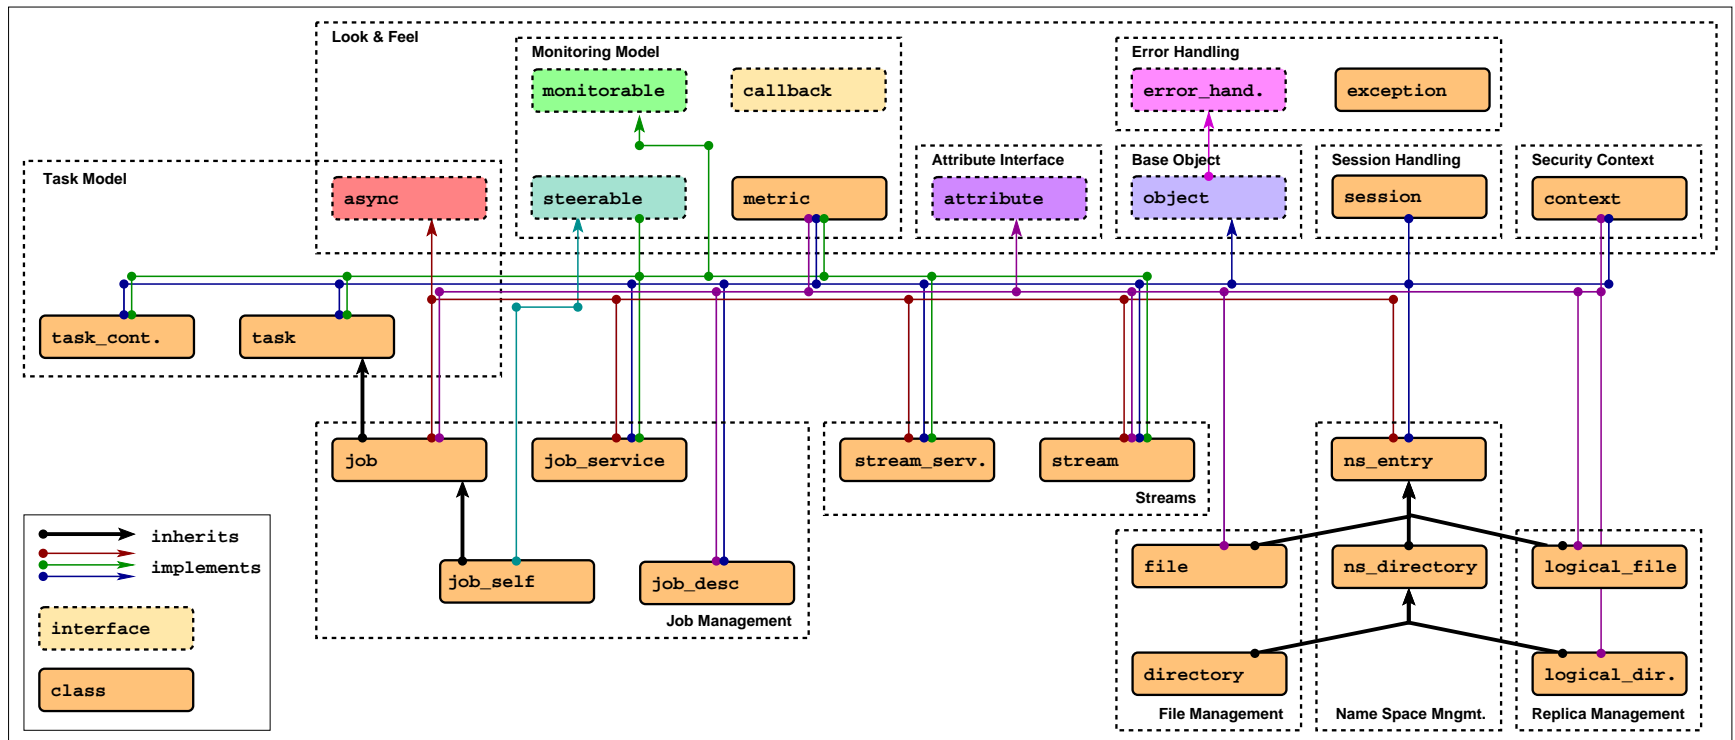
# SAGA: Class hierarchy



## SAGA API Package 'name\_spaces':

manage files, replicas, etc.

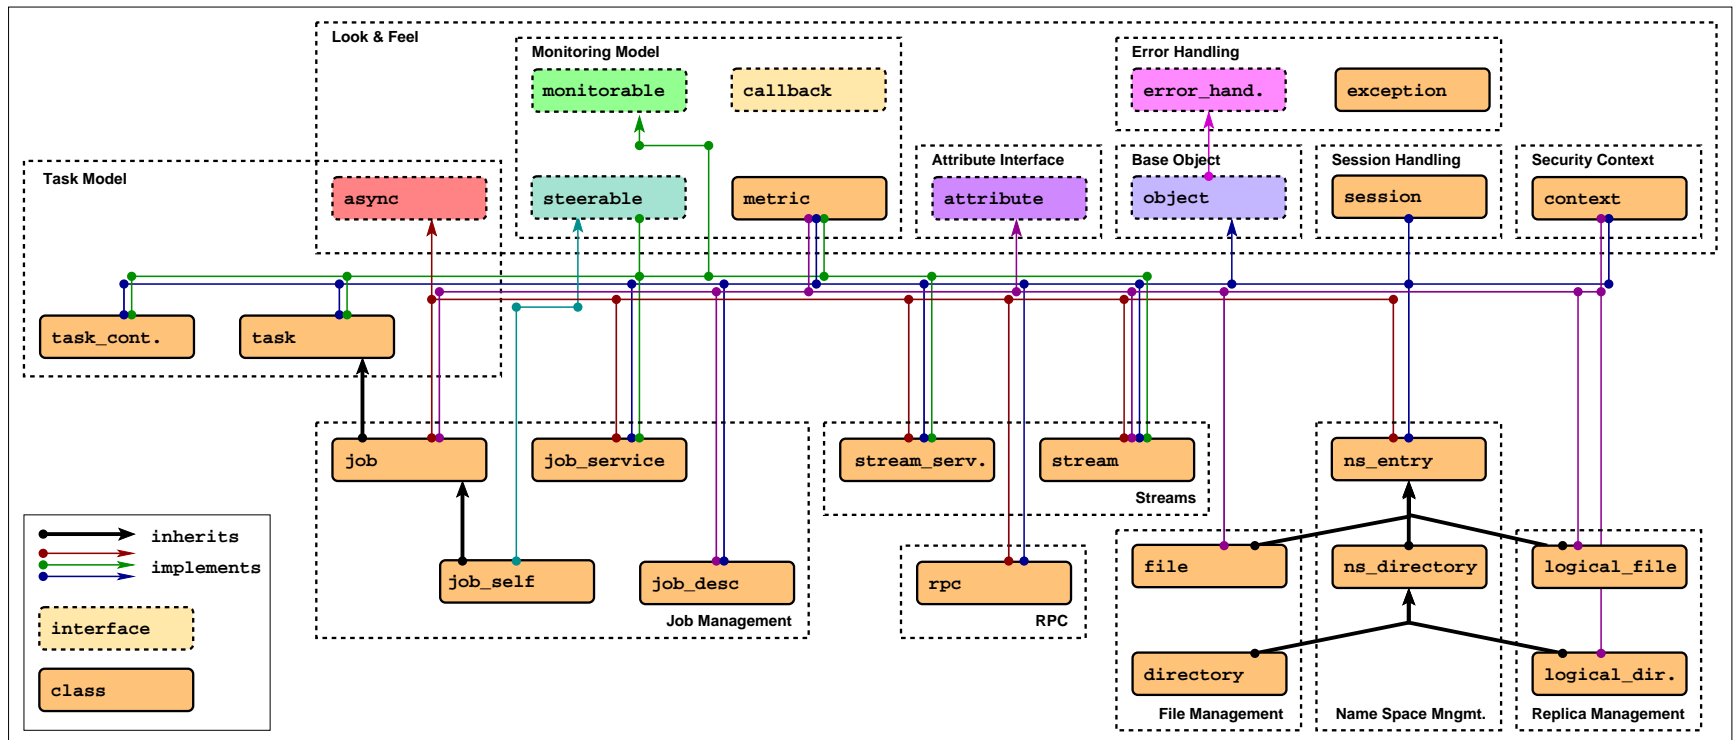
# SAGA: Class hierarchy



**SAGA API Package 'stream':**

SAGA rendering of BSD streams.

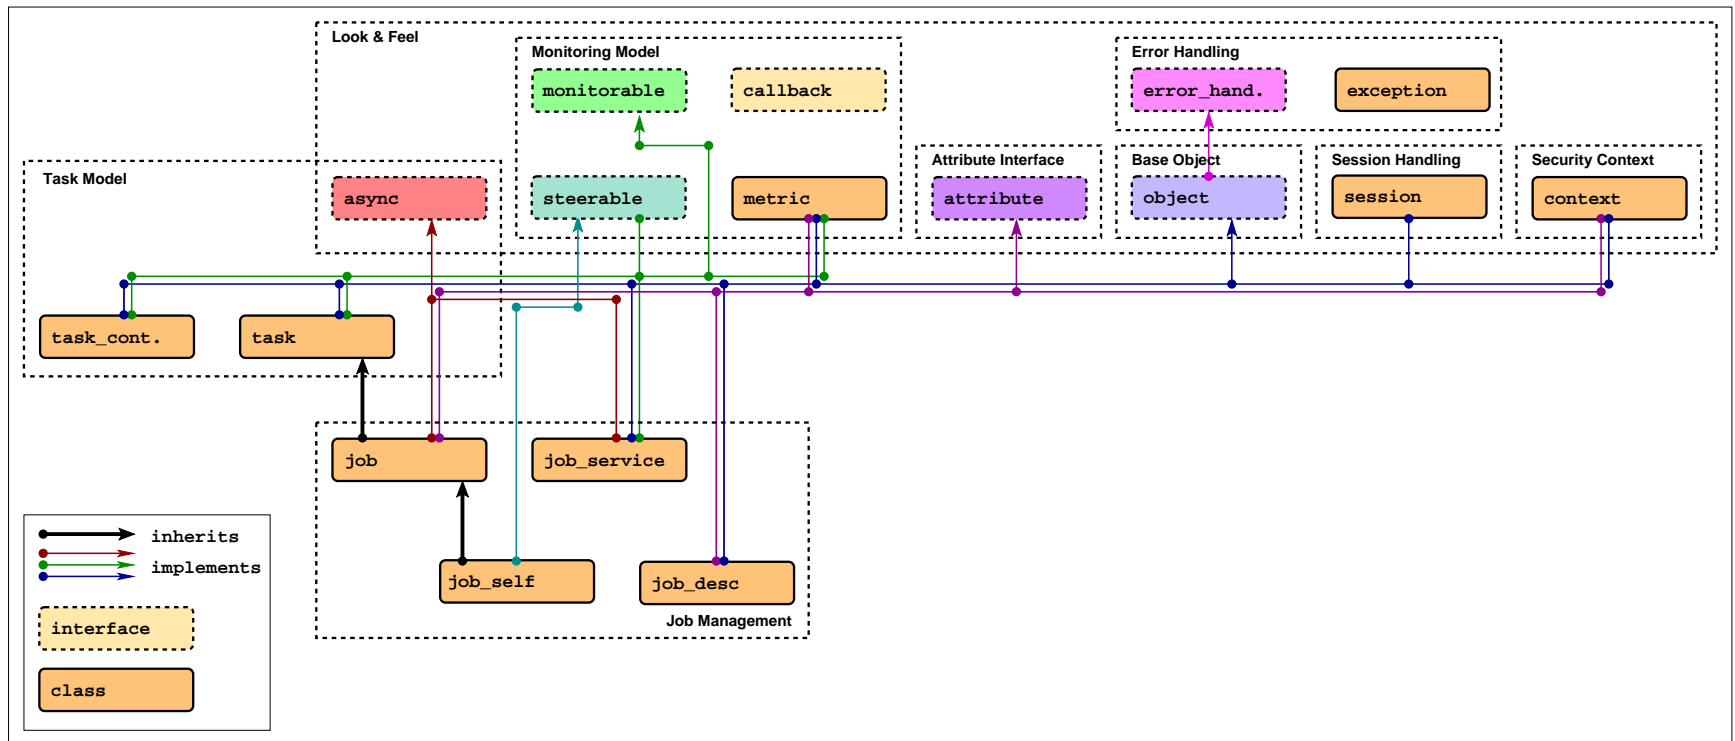
# SAGA: Class hierarchy



**SAGA API Package 'rpc':**  
remote procedure calls.

## Functional API Packages

# SAGA: Jobs

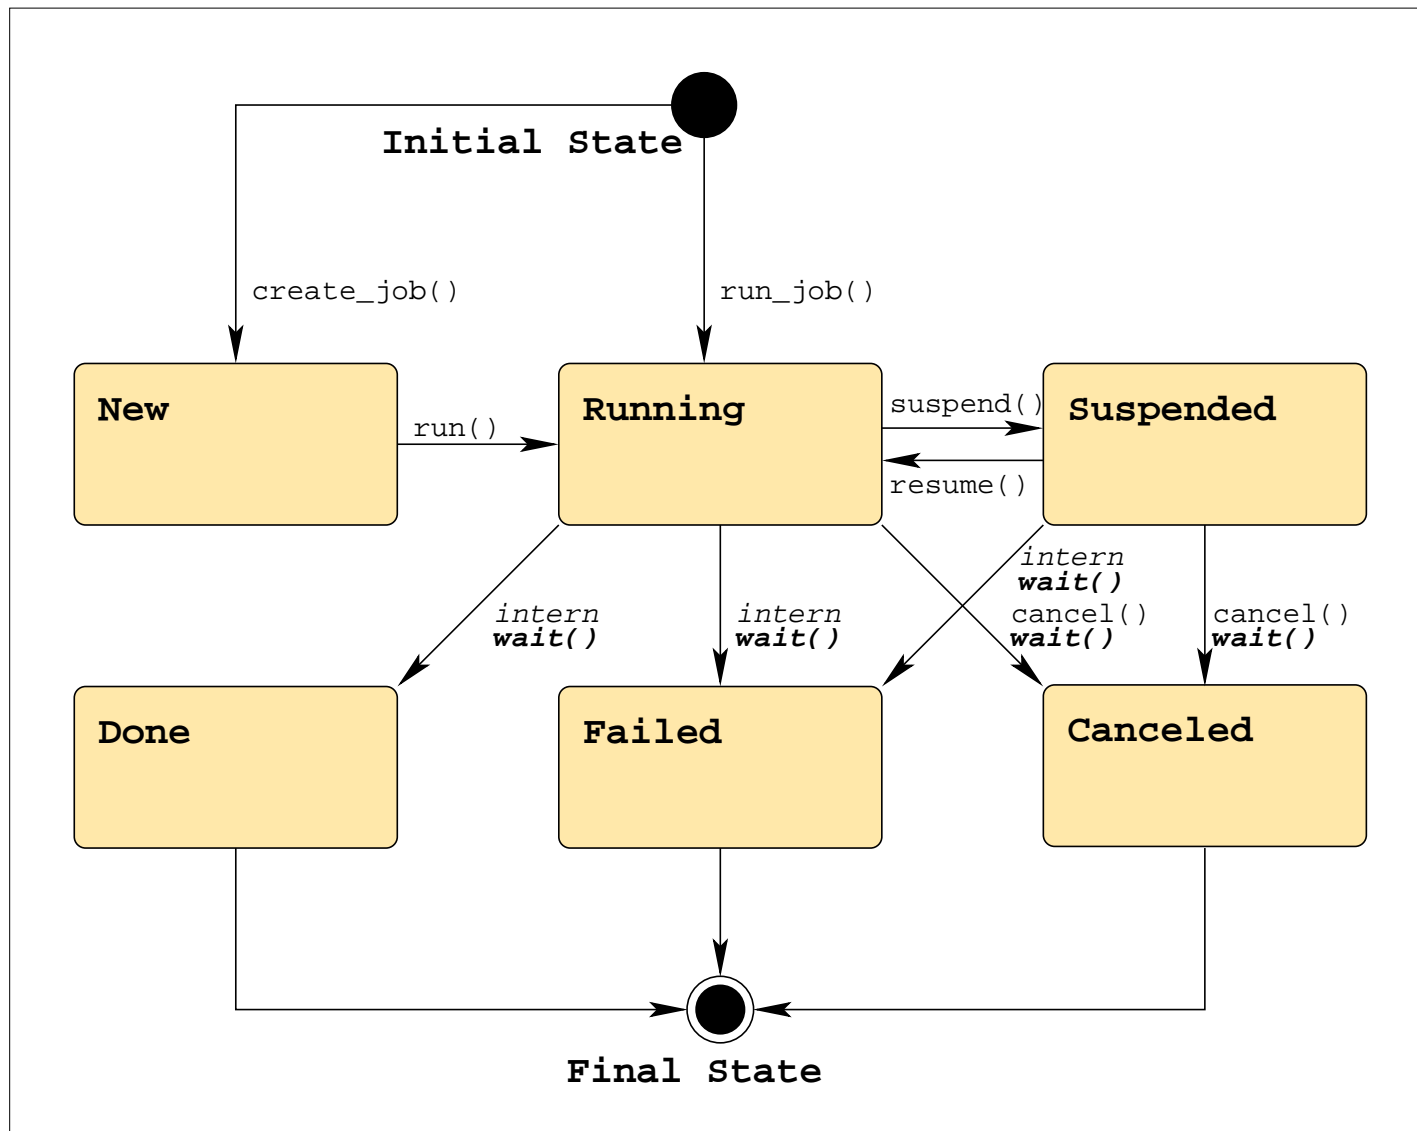


# SAGA: Jobs - Overview

- `job_service` uses `job_description` to create job instances
- `job_description` attributes are based on JSDL
- state model is based on / synced with BES
- `job_self` represents the SAGA application
- job submission and management, but no resource discovery, job dependencies, or workflows



# SAGA: Job States



# SAGA Examples: Jobs

\_\_\_\_\_ job submission \_\_\_\_\_

```
saga::job::service      js ("gram://headnode.gram.net");
saga::job::job          j = js.run_job ("/bin/sleep 10",
                                         "clusternode-2.gram.net");

cout << "Job State: " << j.get_state () << endl;

j.wait ();

cout << "Retval " << j.get_attribute ("ExitCode") << endl;
```

# SAGA Examples: Jobs

\_\_\_\_\_ job submission \_\_\_\_\_

```
saga::job::description jd;  
saga::job::service      js ("gram://remote.host.net");  
saga::job                j = js.create_job (jd);  
  
j.run ();  
  
cout << "Job State: " << j.get_state () << endl;  
  
j.wait ();  
  
cout << "Retval " << j.get_attribute ("ExitCode") << endl;
```

# SAGA Examples: Jobs

jobs (cont.)

```
j.run      ();  
j.wait     ();  
j.cancel   ();  
  
j.suspend  ();  
j.resume   ();  
  
j.signal    (SIGUSR1);  
j.checkpoint ();  
j.migrate   (jd);
```

# SAGA Examples: Job Descr.

\_\_\_\_\_ job description - JSDL based \_\_\_\_\_

```
saga::job::description jd;

jd.set_attribute ("Executable",      "/bin/tail");
jd.set_attribute ("WorkingDirectory", "data/");
jd.set_attribute ("Cleanup",         "False");

// pseudo code *blush*
jd.set_vector_attribute ("Arguments",  [ "-f", "my_log" ]);
jd.set_vector_attribute ("Environment", [ "TMPDIR=/tmp/" ]);
jd.set_vector_attribute ("FileTransfer", [ "my_log >> all_logs" ] );
```

# SAGA Job Description

SAGA JD attributes:

Executable	Arguments	Environment
CandidateHosts	SPMDVariation	TotalCPUCount
NumberOfProcesses	ProcessesPerHost	ThreadsPerProcess
WorkingDirectory	<i>Interactive</i>	Cleanup
Input	Output	Error
<i>JobStartTime</i>	WallTimeLimit	TotalCPUTime
TotalPhysicalMemory	CPUArchitecture	OperatingSystemType
<i>Queue</i>	JobProject	<i>JobContact</i>
FileTransfer		

# SAGA Job Description

- leaning heavily on **JSDL**, but flat
- borrowing from DRMAA
- mixes hardware, software and scheduling attributes!
- cannot be extended
- no support for 'native' job descriptions (RSL, JDL, ...)
- only 'Executable' is required
- backend MAY ignore unsupported keys!

---

```
cd /tmp/data && rm -rf *
```

# SAGA Example: job service

\_\_\_\_\_ job service \_\_\_\_\_

```
saga::job::service js ("gram://remote.host.net/");  
  
vector<string> ids = js.list (); // list known jobs  
  
while ( ids.size () )  
{  
    string id = ids.pop_back ();  
  
    saga::job j = js.get_job (id); // reconnect to job  
  
    cout << id << " : " << j.get_state () << endl;  
}
```



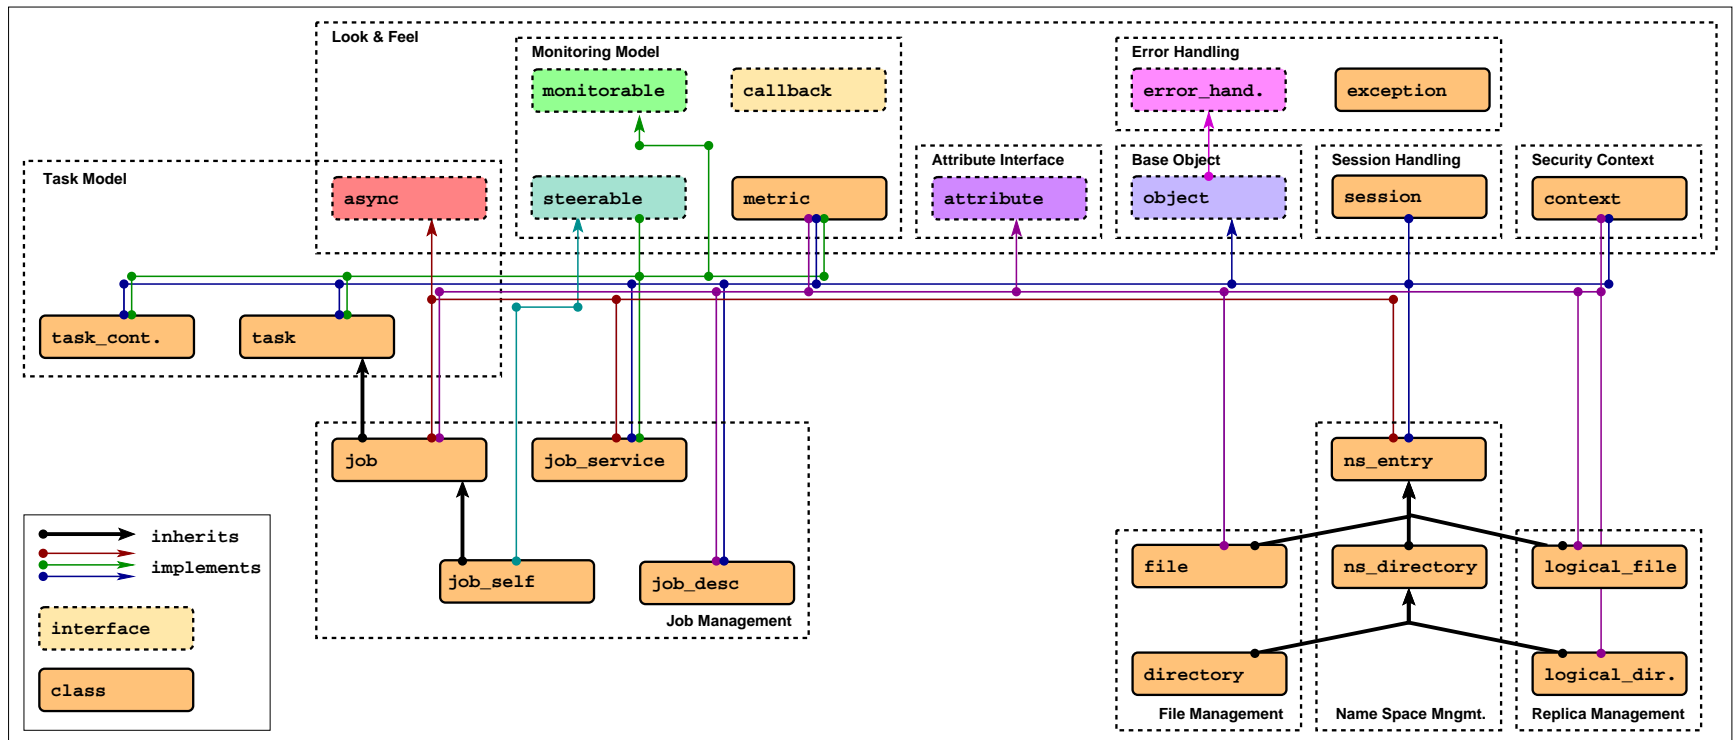
- represents a specific job submission endpoint
- job states are maintained on that endpoint (usually)
- full reconnect may not be possible (I/O streaming)
- lifetime of state up to backend
- reconnected jobs may have different job description (lossy translation)

# SAGA Examples: Job 'Self'

```
_____ jobs (cont.) _____  
  
saga::job::self self = js.get_self ();  
  
self.signal      (SIGUSR1);  
self.checkpoint  ();  
self.migrate     (jd);  
  
self.wait        ();    // blocks forever :-P  
self.cancel      ();
```

- represents the **calling application instance**
- `self.signal (SIGUSR1);`  
is different from  
`::kill (::getpid (), SIGUSR1);`  
as it goes over the job manager (accounting,  
logging, cleanup...)
- 'job::self' is the only SAGA object to which metrics  
can be added (details later...)

# SAGA: Name Spaces etc.



# SAGA: Name Spaces

- interfaces for managing entities in name spaces
- files, replicas, information, resources, steering parameter, checkpoints, ...
- manages hierarchy (mkdir, cd, ls, ...)
- manages NS entries as opaque (copy, move, delete, ...)

- implements name space interface, and adds access to content of NS entries (files)
- Posix oriented: read, write seek
- Grid optimizations: scattered I/O, pattern based I/O, extended I/O

# SAGA: Replicas

- implements name space interface, and adds access to properties of NS entries (logical files / replicas)
- O/REP oriented: list, add, remove replicas; manage meta data
- Grid optimizations are hidden (replica placement strategies, consistency and version management, ...)

- implements name space interface, and adds access to arbitrary key/value pairs on each entry.
- entries also allow to store a **serialized** SAGA Object!
- utterly useful for application bootstrapping, communication between different application modules, application persistency, etc etc.
- **not part of the SAGA Specification**, but is getting standardized as an extension



# SAGA Examples: NameSpaces

\_\_\_\_\_ name space management \_\_\_\_\_

```
saga::name_space::directory d ("ssh://remote.host.net//data/");

if ( d.is_entry ("a") && ! d.is_dir ("a") )
{
    d.copy ("a", "../b");
    d.link ("../b", "a", Overwrite);
}

list <saga::url> names = d.find ("*-{123}.text.");

saga::name_space::directory tmp  = d.open_dir ("tmp/data/1",
                                                saga::name_space::CreateParents);
saga::name_space::entry      data = tmp.open   ("data.txt");

data.copy ("data.bak", Overwrite);      // uses cwd
```

- name space entries are opaque: the name space package can never look inside
- directories are entries (inheritance)
- **inspection:**  
`get_cwd, get_url, get_name, exists, is_entry, is_dir, is_link, read_link`
- **manipulation:**  
`create (c'tor, open), copy, link, move, remove`
- **permissions:**  
`permissions_allow, permissions_deny`
- wildcards are supported (remember POSIX influence...)

# SAGA Examples: Files

```
_____ file access _____  
saga::filesystem::file f ("any://remote.host.net/data/data.bin");  
  
char mem[1024];  
saga::mutable_buffer buf (mem);  
  
if ( f.get_size () >= 1024 )  
{  
    buf.set_data (mem + 0, 512);  
    f.seek (512, saga::filesystem::Start);  
    f.read (buf);  
}  
  
if ( f.get_size () >= 512 )  
{  
    buf.set_data (mem + 512, 512);  
    f.seek (0, saga::filesystem::Start);  
    f.read (buf);  
}
```

- provides access to the **content** of filesystem entries (sequence of bytes)
- saga buffers are used to wrap raw memory buffers
- saga buffers can be allocated by the engine
- several incarnations of read/write: posix style, scattered, pattern based

# SAGA Name Spaces: Flags

```
enum flags {  
    None          = 0,  
    Overwrite     = 1,  
    Recursive     = 2,  
    Dereference   = 4,  
    Create        = 8,  
    Exclusive     = 16,  
    Lock          = 32,  
    CreateParents = 64,  
    Truncate      = 128,    // not on name_space  
    Append        = 256     // not on name_space  
    Read          = 512,  
    Write         = 1024,  
    ReadWrite     = 1536    // Read | Write  
    Binary        = 204     // only on filesystem  
}
```

# SAGA Examples: Replicas

\_\_\_\_\_ replica management \_\_\_\_\_

```
saga::replica::directory dir ("raptor://remote.host.net/data/");

if ( dir.is_entry ("a") || dir.is_link ("a") )
{
    dir.copy ("a", "../b");
    dir.link ("../b", "a");
}

saga::replica::file file = dir.open ("tmp/data.txt");
list <string>  locations = file.list_locations ();

file.replicate ("gridftp://other.host.net/tmp/a.dat");
```

- provides access to the **content** of replica system entries (list of physical locations, plus attributes)
- saga attribute interface is used for entry meta data. Meta data are maintained by application, and/or backend.
- `replicate( )` creates a new copy, and adds new location to list

# SAGA Examples: Replicas

\_\_\_\_\_ replica meta data \_\_\_\_\_

```
saga::replica::directory dir ("raptor://remote.host.net/data/");  
  
list <saga::url> files = dir.find ("*", "type=jpg");  
  
while ( file.size () )  
{  
    saga::logical_file lf (file.pop_front ());  
  
    lf.replicate ("file://localhost/data/images/",  
                 saga::replica::Overwrite);  
}
```



- persistent storage of application level information
- semantics of information defined by application
- allows storage of serialized SAGA objects (object persistency)

# SAGA Examples: Adverts

## Adverts

```
saga::advert::directory todo ("any//remote.host.net/my_tasks/");

// pseudo vector code
list <saga::url> urls = todo.find ("*", ["priority=urgent"]);

while ( urls.size () )
{
    saga::advert ad (urls.pop_front ());


    std::cout << ad.get_attribute ("description") << std::endl;
}
```

# SAGA Examples: Adverts

## Persistent SAGA Objects

```
saga::file    f    (url);  
saga::advert  ad  ("any//remote.host.net/files/my_file_ad", Create);  
  
ad.store_object (f);  
  
-----  
  
saga::advert  ad  ("any//remote.host.net/files/my_file_ad");  
saga::file    f  = ad.retrieve_object ();
```

# SAGA: Streams



# SAGA Examples: Streams

## stream server

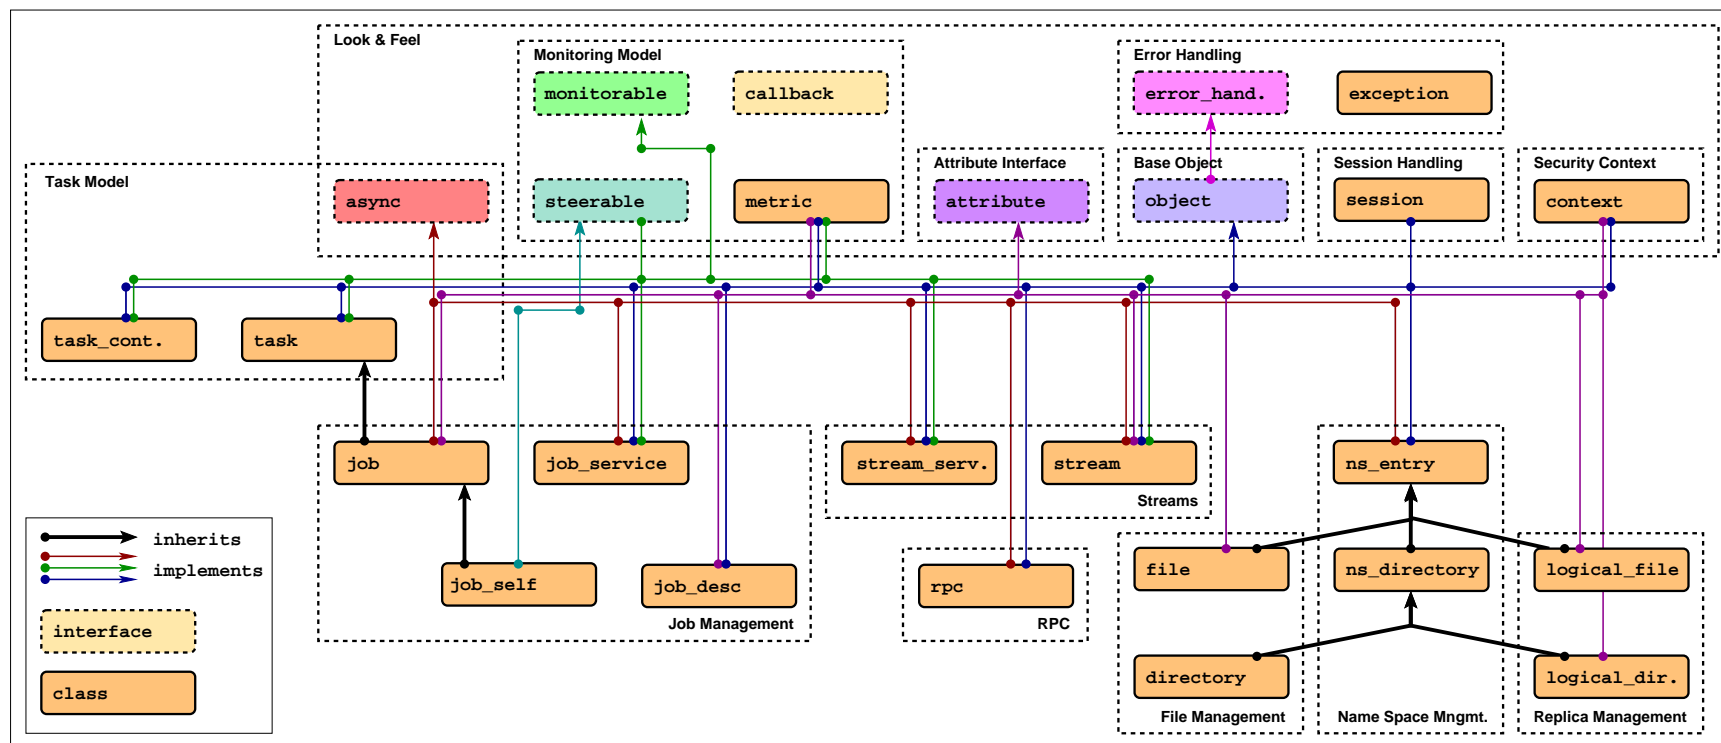
```
saga::stream_service ss ("tcp://localhost:1234");  
  
saga::stream_client sc = ss.serve ();  
  
sc.write ("Hello client", 13);
```

## stream client

```
char buf [13];  
saga::stream_client sc ("tcp://remote.host.net:1234");  
  
sc.connect ();  
sc.read      (buf, 13);  
  
cout << buf << endl;
```

- simple and BSD socket oriented
- not supposed to replace MPI etc, but allows for simple application level communication
- will be superceded by message package

# SAGA: RPC



# SAGA Examples: RPC

\_\_\_\_\_ remote procedure call \_\_\_\_\_

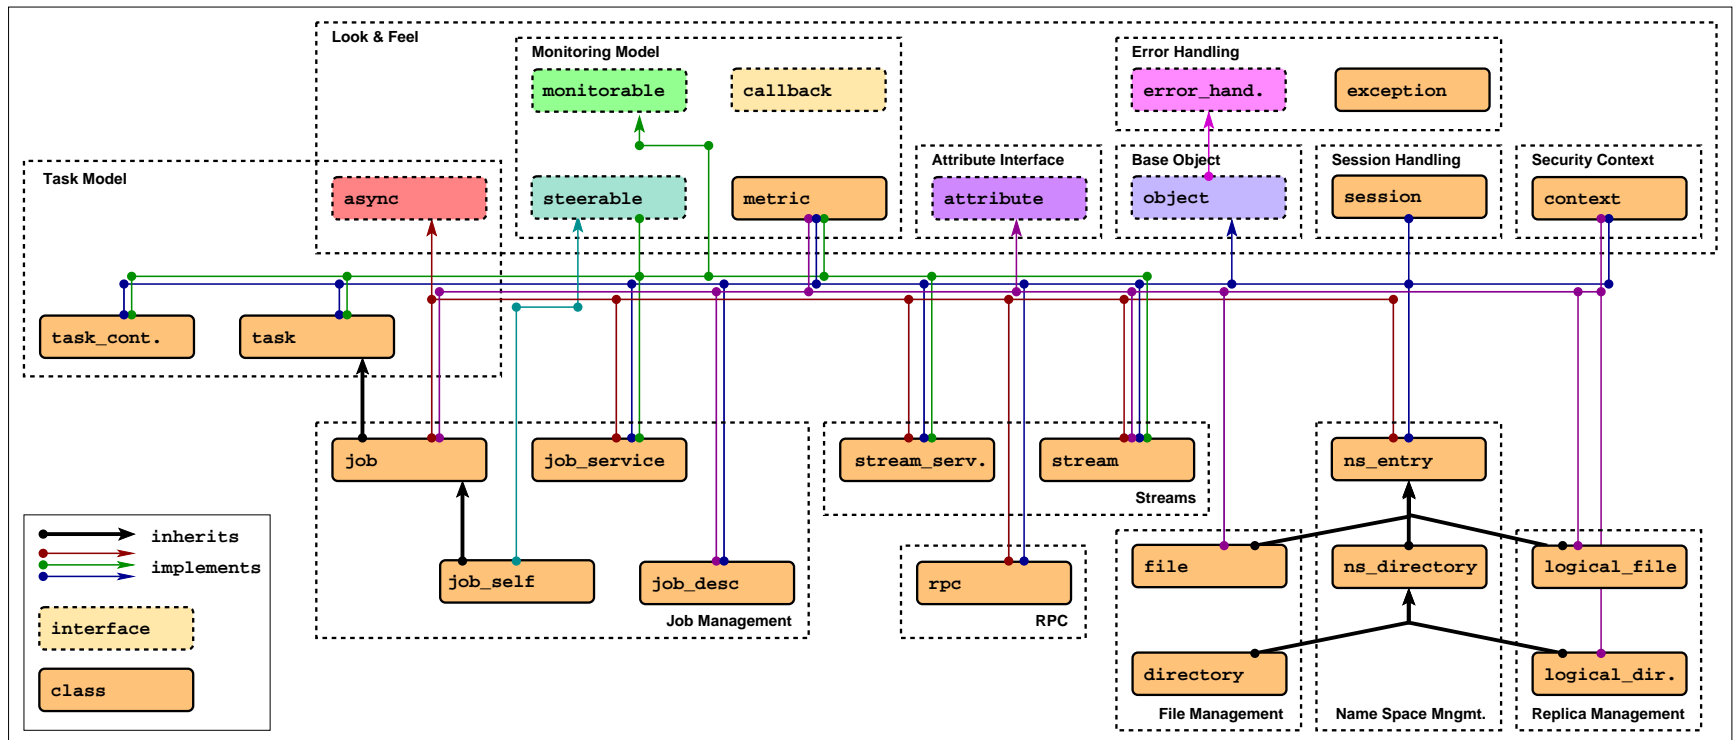
```
saga::rpc rpc ("ninfgr://remote.host.net:1234/random");  
  
list <saga::rpc::parameter> params;  
params.push_back (new saga::rpc::parameter (Out, 10));  
  
rpc.call (params);  
  
cout << "found random number: " << ::atoi (param.buffer) << endl;  
  
delete (params.pop_front ());
```



- maps GridRPC standard into the SAGA look & feel
- parameters are stack of structures (similar to scattered I/O)
- future revision will work on optimized data handling

## Non-Functional API Packages

# SAGA: Session and Context



# SAGA Examples: Session

\_\_\_\_\_ default sessions \_\_\_\_\_

```
saga::ns_dir dir ("any://remote.host.net//data/");

if ( dir.is_entry ("a") && ! dir.is_dir ("a") )
{
    dir.copy ("a", "../b");
    dir.link ("../b", "a", Overwrite);
}

list <saga::url> names = dir.find ("*-{123}.text.");

saga::name_space::directory tmp = dir.open_dir ("tmp/");
saga::name_space::entry entry = tmp.open ("data.txt");

entry.copy ("data.bak", Overwrite);
```

# SAGA Examples: Session

context management

```
saga::context c1 (saga::context::X509);  
saga::context c2 (saga::context::X509);  
  
c2.set_attribute ("UserProxy", "/tmp/x509up_u123.special");  
  
saga::session s;  
  
s.add_context (c1);  
s.add_context (c2);  
  
saga::name_space::dir dir (s, "any://remote.host.net/data/");
```

# SAGA: Session Management

- by default hidden (default session is used)
- session is identified by lifetime of security credentials and by objects in this session (jobs etc.)
- session is used on object creation (optional)
- `saga::context` is used to attach security tokens to a session
- the default session has default contexts

# SAGA Examples: Session

\_\_\_\_\_ session inheritance \_\_\_\_\_

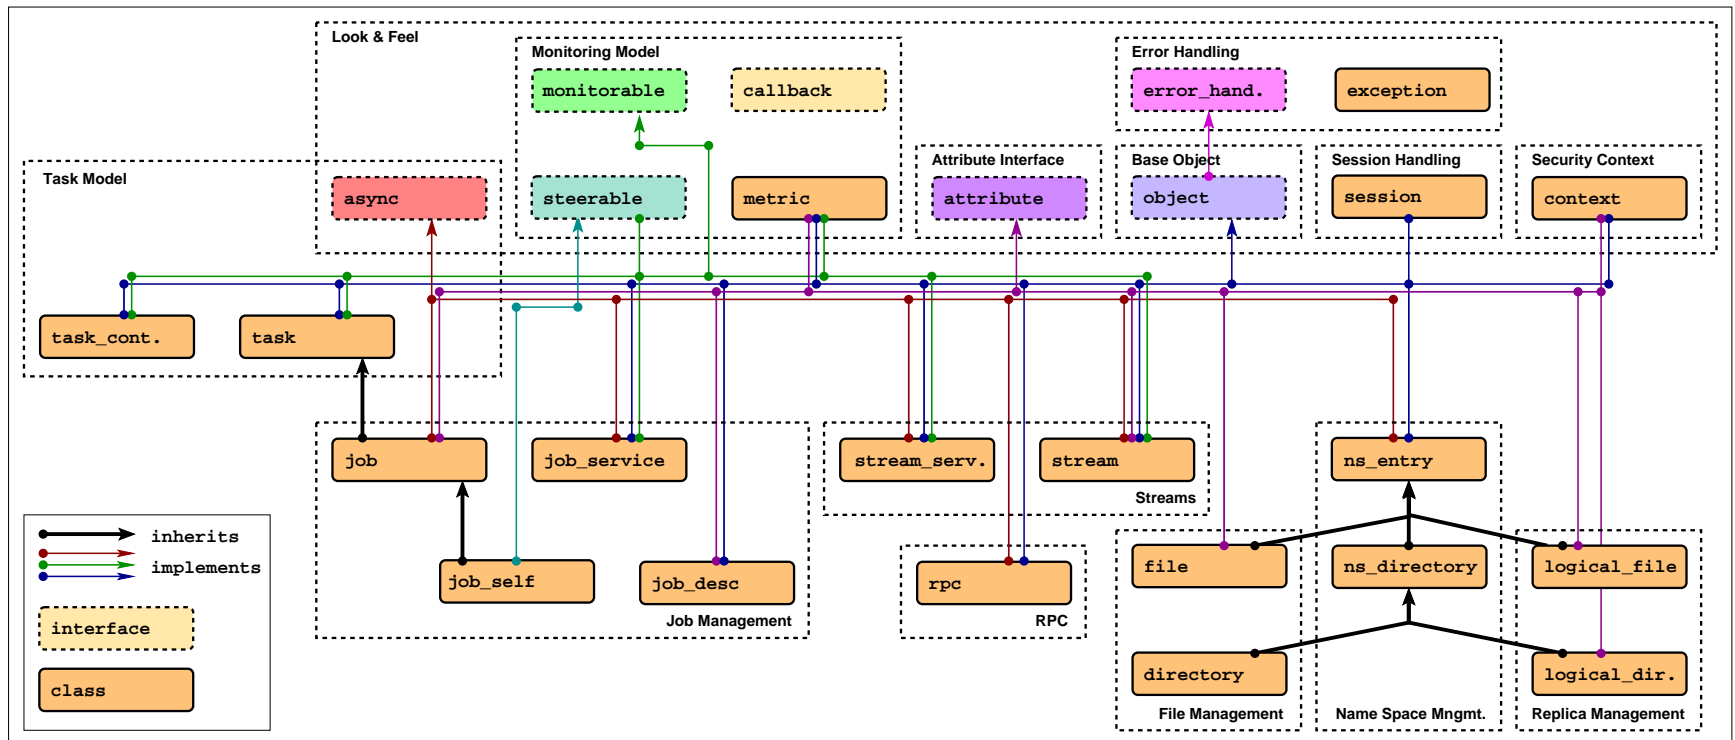
```
saga::dir  dir (s, "gridftp://remote.host.net/data/");  
  
saga::file file = dir.open ("data.bin");  
  
s.remove_context (c1);  
s.remove_context (c2);  
  
file.copy ("data.bin.bak");    // works - state is sticky!
```

# SAGA Examples: Session

```
_____ authorization _____  
  
// server side code  
saga::stream_service ss ("tcp://localhost:1234");  
  
saga::stream_client sc = ss.serve ();  
  
saga::context c = sc.get_context ();  
  
if ( c.get_type == Globus &&  
    c.attribute_equals ("RemoteID", "O=MyCA, O=MyOrg, CN=Joe" ) )  
{  
    sc.write ("welcome!", 9);  
}  
else  
{  
    sc.write ("bugger off!", 12);  
    sc.close ();  
}
```



# SAGA: Monitoring



- monitoring of Grid entities (jobs, files, . . .)
- monitoring of interactions (task state, notification, . . .)
- `monitorables` have metrics
- `metrics` can be pulled, or subscribed to (callbacks)
- some metrics can be written (basic steering)

# SAGA Examples: Monitoring

pull monitoring

```
saga::job job = js.create_job (jd);  
  
job.run ();  
  
saga::metric m = job.get_metric ("MemoryUsage");  
  
while ( 1 )  
{  
    cout << "Memory Usage: " << m.get_value () << endl;  
    sleep (1);  
}
```

# SAGA Examples: Monitoring

```
_____ callbacks _____  
class my_cb : public saga::callback  
{  
    public:  
        bool cb (saga::monitorable obj,  
                 saga::metric      m,  
                 saga::context      c)  
        {  
            cout << "Memory Usage: " << m.get_value () << endl;  
            return (true);  
        }  
};  
  
my_cb cb;  
saga::job job = js.create_job (jd);  
job.run ();  
  
saga::metric m = job.get_metric ("MemoryUsage");  
m.add_callback ("MemoryUsage", cb);
```


# SAGA Examples: Monitoring

```
_____ callbacks _____  
class my_cb : public saga::callback  
{  
    public:  
        bool cb (saga::monitorable obj,  
                 saga::metric          m,  
                 saga::context          c)  
        {  
            cout << "Memory Usage: " << m.get_value () << endl;  
            return (true);  
        }  
};  
  
my_cb cb;  
saga::job job = js.create_job (jd);  
job.run ();  
  
job.add_callback ("MemoryUsage", cb);
```

# SAGA Examples: Monitoring

```
_____ callbacks (cont.) _____  
class my_cb : public saga::callback  
{  
    public:  
        bool cb (saga::monitorable obj,  
                 saga::metric      m,  
                 saga::context      c)  
        {  
            cout << m.get_name () << " : " << m.get_value () << endl;  
            return (true);  
        }  
};  
  
list <string> metrics = job.list_metrics ();  
  
while ( metrics.size () )  
{  
    job.add_callback (metrics.pop_front (), cb);  
}
```

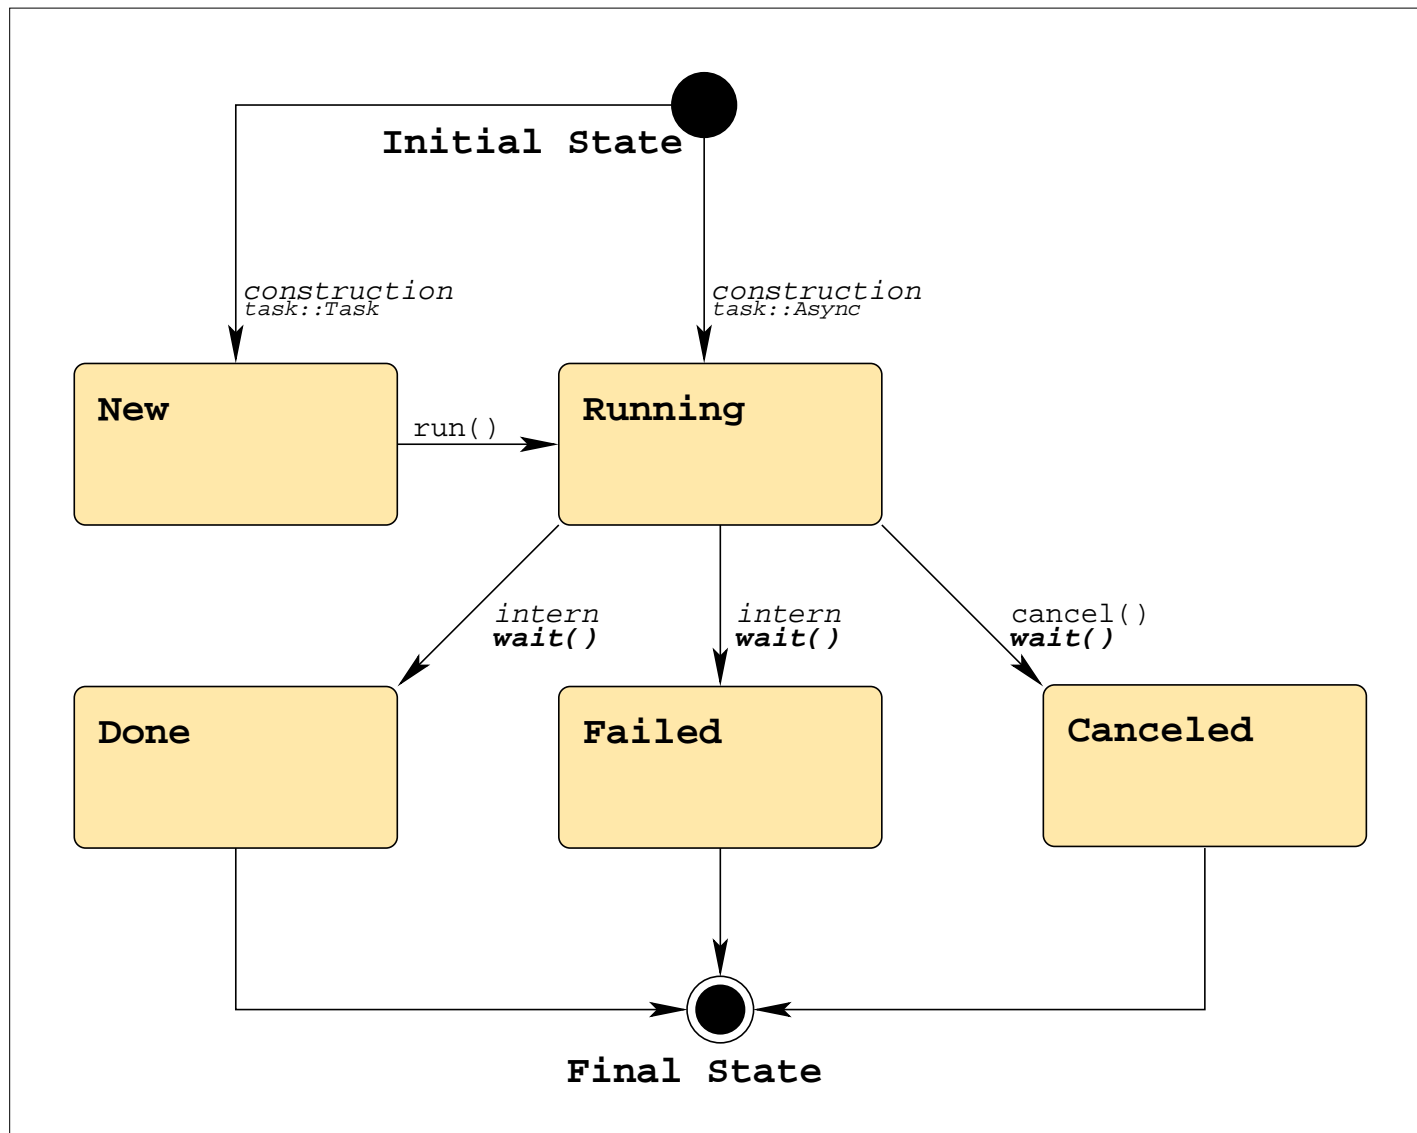
# SAGA: Tasks



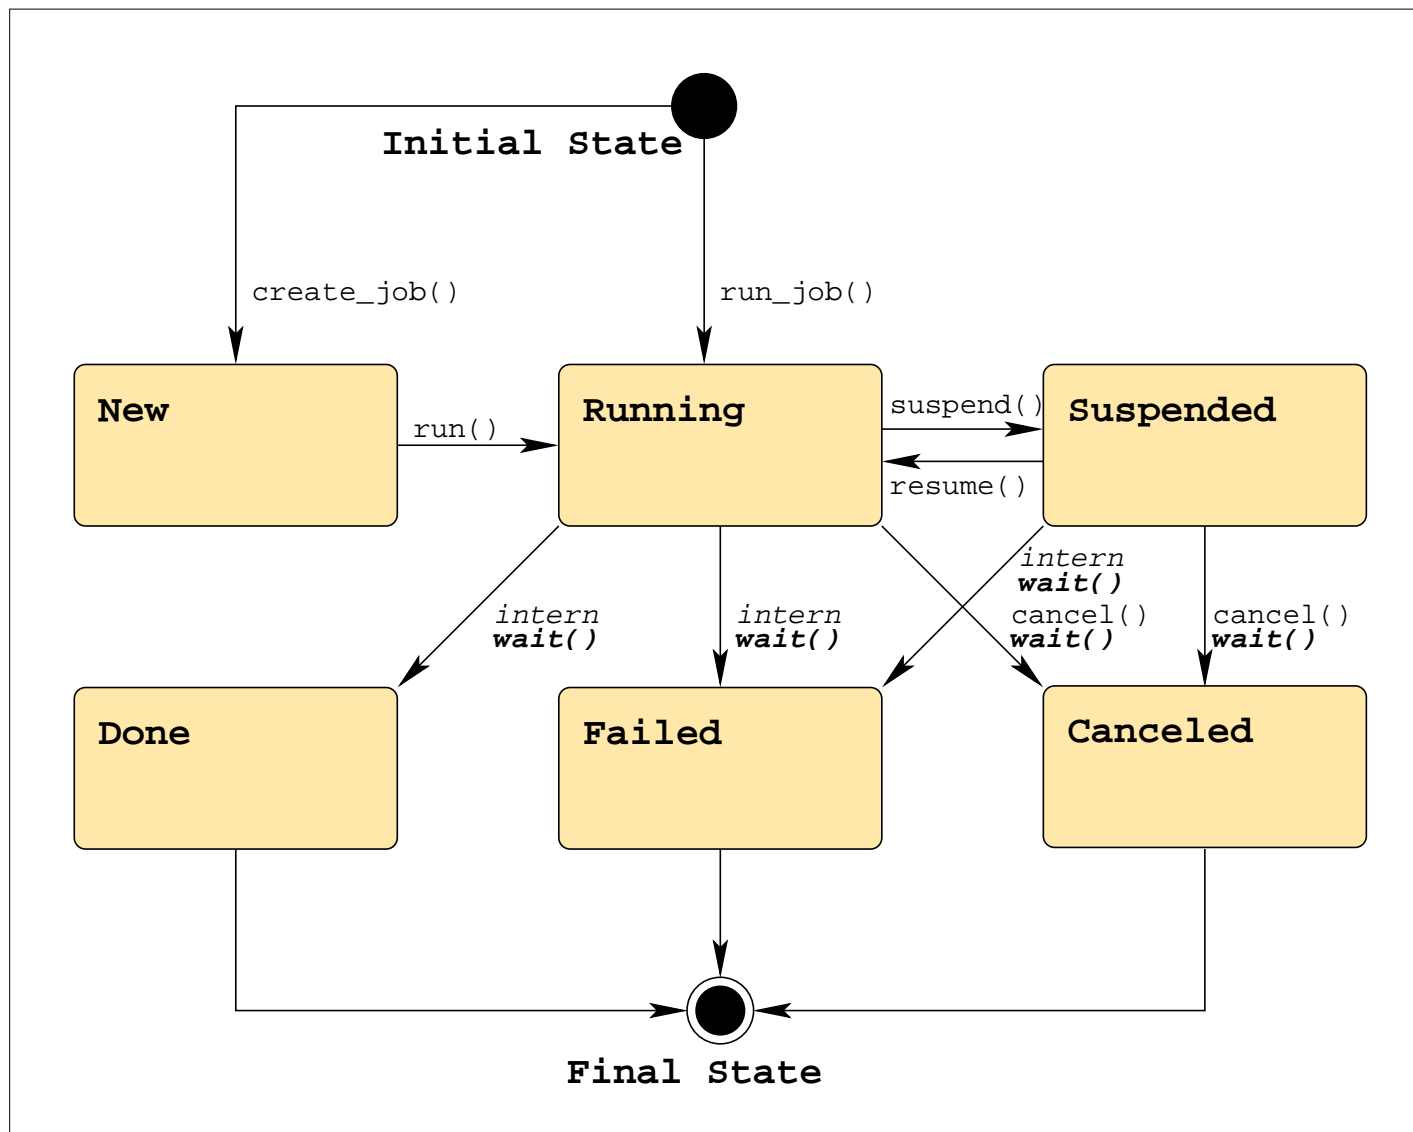
- asynchronous operations are a MUST in distributed systems, and Grids
- `saga::task` represents an synchronous operation (e.g. `file.copy ( )`)
- `saga::task_container` manages multiple tasks
- tasks are stateful (similar to jobs)



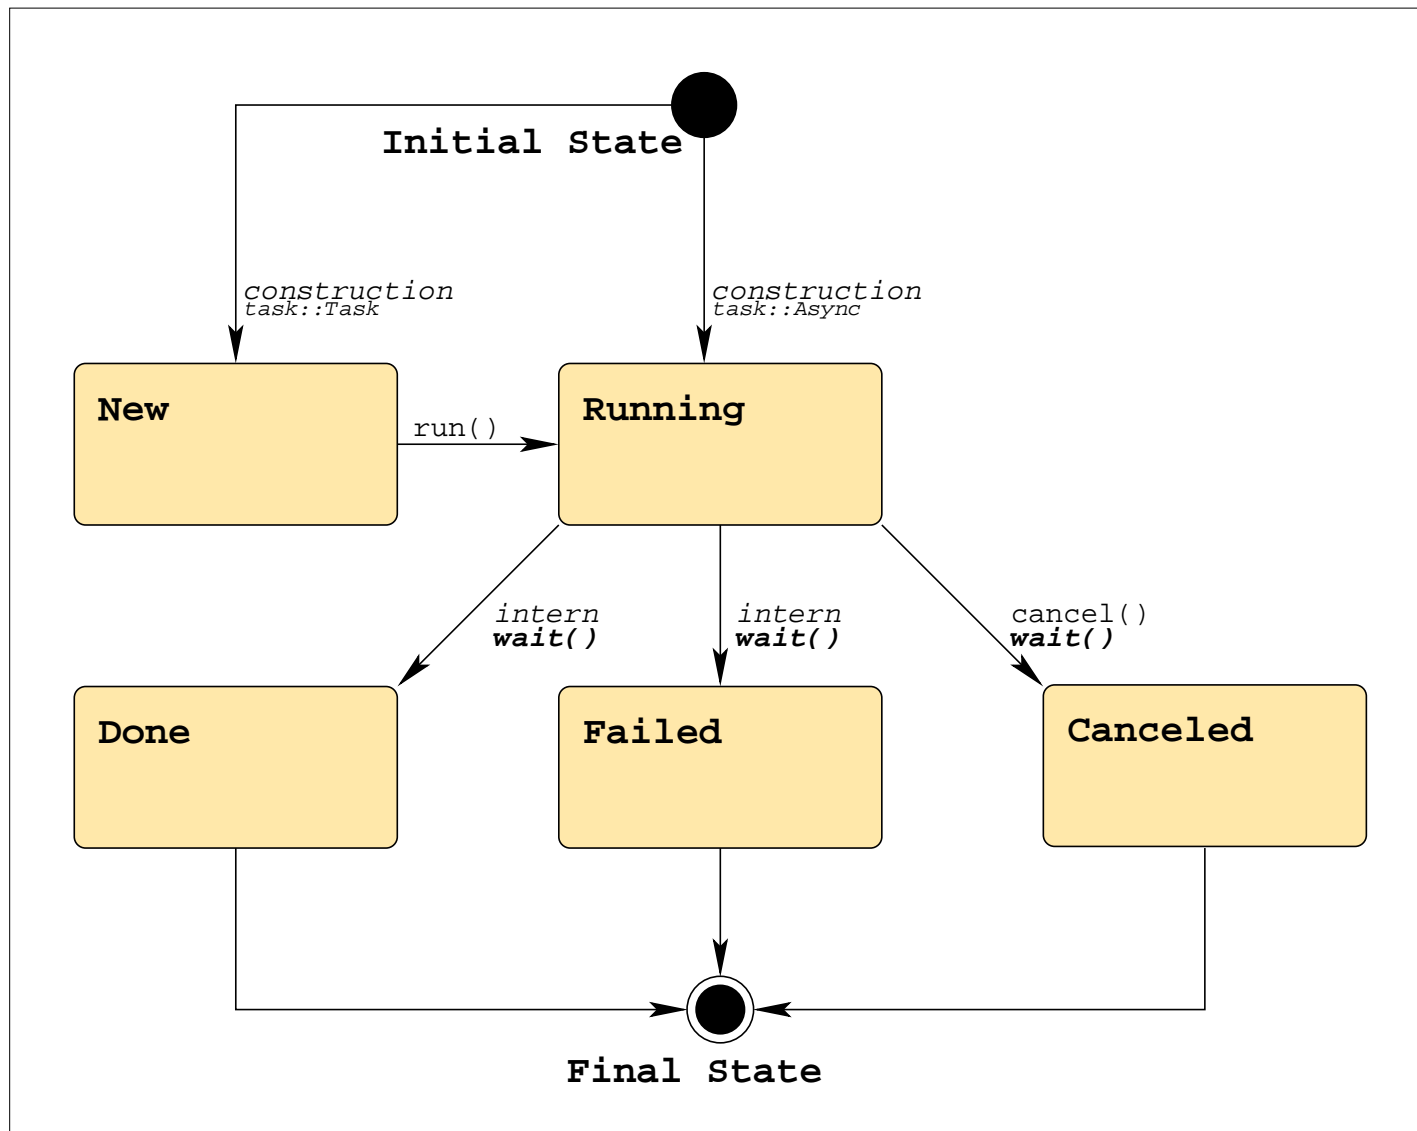
# SAGA: Task States



# SAGA: Job States



# SAGA: Task States



- different versions for each method call: sync, async, task
- signature *basically* the same
- differ in state of task representing that method

# SAGA Examples: Tasks

```
_____ tasks (i) _____

saga::file file ("gsiftp://remote.host.net/data/data.bin");

// normal, synchronous
file.copy ("data.bak");

// async versions
saga::task t1 = file.copy <saga::task::Sync> ("data.bak.1");
saga::task t2 = file.copy <saga::task::Async> ("data.bak.2");
saga::task t3 = file.copy <saga::task::Task> ("data.bak.3");

// t1: Done
// t2: Running
// t3: New
```

# SAGA Examples: Tasks

```
tasks (ii)

t3.run ();

cout << t3.get_state () << endl;  // Running

t2.wait ();
t3.wait ();

// t1, t2, t3: Done (or Failed...)
```

# SAGA Examples: Tasks

tasks container

```
saga::task_container tc;  
  
tc.add (t1);  
tc.add (t2);  
tc.add (t3);  
  
tc.run  ();  
  
saga::task done_task = tc.wait (Any);  
  
tc.wait (All);
```

# SAGA Examples: Tasks

tasks jobs and notification

```
saga::task task = file.copy <saga::task::Asyn> ("b");  
saga::job  job  = js.run_job ("remote.host.net", "/bin/date");  
  
task.add_callback ("State", my_cb);  
job.add_callback  ("State", my_cb);  
  
saga::task_container tc;  
  
tc.add (task);  
tc.add (job);  
  
tc.wait ();
```



# SAGA planned extensions

- service discovery
- message based communication
- information service (Advert Service)
- resource discovery and management
- checkpoint & recovery (GridCPR)

# SAGA v2: Messages

## \_\_\_\_\_ Messaging server \_\_\_\_\_

```
saga::sender snd ("tcp://localhost:1234");

saga::msg msg;

msg.set_size (100); // arbitrary size!
msg.set_data ("abcd...");

sc.send (msg);
```

## \_\_\_\_\_ Messaging client \_\_\_\_\_

```
char buf [13];
saga::receiver rec ("tcp://remote.host.net:1234");

// int size = rec.test ();

saga::msg = rec.receive (); // internal buffer allocation
```

# SAGA v2: Messages

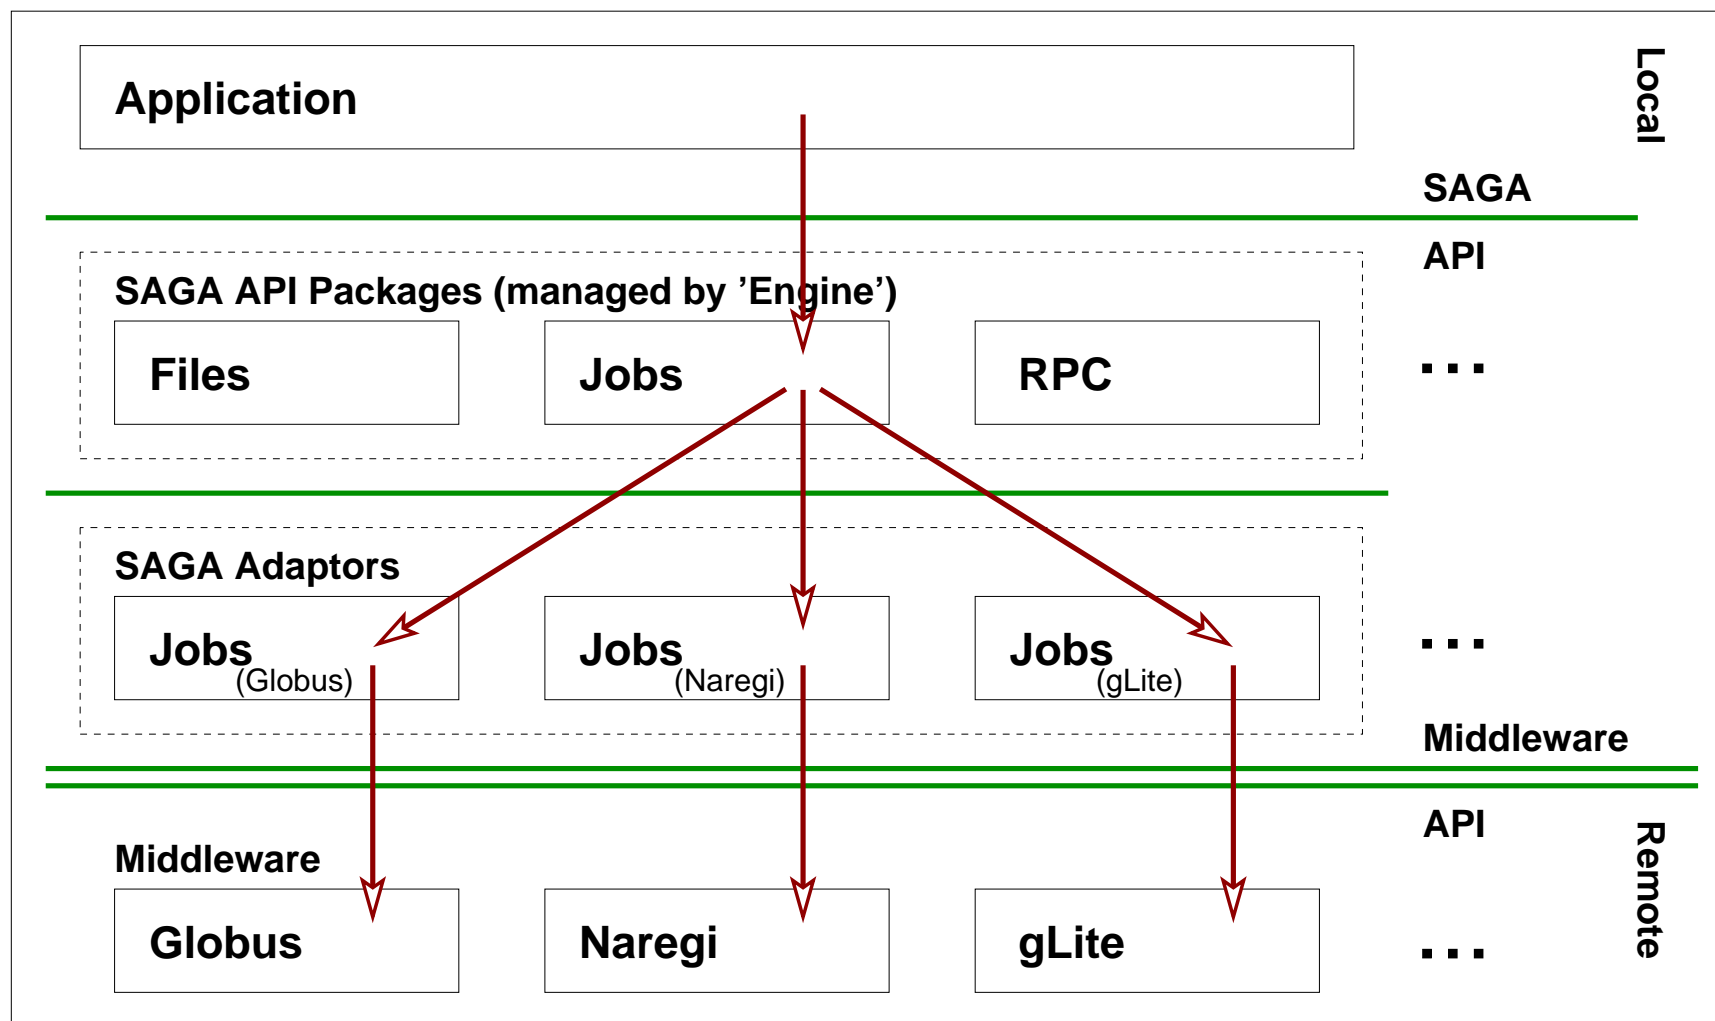
- messages are received intact or not at all
- implies protocol, but is silent about interop
- async zero copy implementation is possible

- no examples yet, API in flux (service spec in flux)
- allows to manage (find, move, stage, archive) checkpoints
- allows to trigger checkpointing of jobs
- probably name space based, with notification on CP creation

# Questions about API?

# Comments?

# Implementation



# Implementation Status

- C++
  - complete implementation by CCT/VU
  - in sync with spec, work in progress
  - other language bindings planned on top (C, Python, Perl, .Net)
- Java (out of sync with the spec)
  - partial implementation (jobs, files) by DEISA/EPCC
  - complete implementation by OMII-UK
  - possible complete implementation at VU
- MiddleWare Bindings
  - DEISA/EPCC binds to DEISA
  - OMII-UK binds to OMII stack, but is flexible
  - C++ adaptor based – GT4 and XtremOS planned

# Contact

`http://forge.ggf.org/sf/projects/saga-core-wg`

→ wiki, CVS details



# Questions?

# Backup Slides

# Recap: Grid Applications

## Types of Grid Applications

1. legacy applications
2. legacy distributed applications
3. Grid aware applications

# Legacy Applications

- no access to application code
- virtualization of heterogeneity
- use cases: remote resource utilization, high throughput
- **favourite technique:** sandboxing (virtualization)
- no need for Grid APIs (application exists, non-mutable)
- *not a topic for this tutorial*

# Legacy Distributed Applications

- aware of distribution (MPI, CORBA, ...)
- not aware of Grid properties (VO)
- usually not very dynamic or adaptive (bootstrapping!)
- use cases: scientific applications, bussiness applications
- **favourite technique:** emulation (GridMPI, etc.)
- no need for *new* Grid APIs (APIs non-mutable)
- *not a topic for this tutorial*

- aware of distribution, heterogeneity, VOs, elasticity etc.
- often dynamic and adaptive structure
- use cases: collaborative, adaptive, auto-optimized, scalable, ...
- **favourite technique:** depends on structure, middleware, ...
- need for Grid APIs (see Part 1)
- *topic for this tutorial :-)*

# Grid APIs: Globus (pre-WS)

- low level API for the Globus Grid Middleware
- scope reflects Globus services:
  - GridFTP
  - GRAM
  - MDS
  - Replicas
- some low level API abstractions (xio, gss-assist)
- **CoG** provides higher level API abstraction for Globus (Java)

# GridFTP Example

## GridFTP: Connection Setup

```
globus_module_activate (GLOBUS_FTP_CLIENT_MODULE);  
  
globus_ftp_client_handleattr_init      (&handle_attr);  
  
globus_ftp_client_handle_init          (&handle, &handle_attr);  
globus_ftp_client_handle_cache_url_state (&handle, server.c_str());
```



# GridFTP Example (ii)

\_\_\_\_\_ GridFTP: Get File Size \_\_\_\_\_

```
globus_ftp_client_operationattr_init      (&attr);
globus_ftp_client_operationattr_set_mode (&attr, ...);

globus_off_t      size      = GLOBUS_NULL;
globus_result_t success = globus_ftp_client_size
                           (&handle,
                            url.c_str(),
                            &attr,
                            &size,
                            GLOBUS_NULL, // done_callback,
                            GLOBUS_NULL);

if (success != GLOBUS_SUCCESS)
{ ... }
```

- API covers full scope of GridFTP protocol
- low level control over connection and operations
- allows synchronous and asynchronous calls (callbacks)

# GRAM Example

GRAM: Job Submit

```
globus_gram_client_callback_allow    (callback_func,  
                                      (void *) &Monitor,  
                                      &callback_contact);  
  
rc = globus_gram_client_job_request (rm_contact,  
                                     job_description,  
                                     job_state_mask,  
                                     callback_contact,  
                                     &job_contact);
```

- API provides full scope of GRAM protocol
- low level control over operations
- synchronous and asynchronous calls
- job details encapsulated in job description (RSL)

# GRAM Example (RSL)

GRAM: RSL example

```
+ ( &
  ( directory      = "/home/user/demo" )
  ( jobtype        = mpi )
  ( executable     = "/home/user/demo/mpi-application" )
  ( maxWallTime    = "10" )
  ( count          = "8" )
  ( architecture   = "i386" )
)
( &
  ( directory      = "/home/user/demo" )
  ( jobtype        = mpi )
  ( executable     = "/home/user/demo/mpi-application" )
  ( maxWallTime    = "10" )
  ( count          = "16" )
  ( architecture   = "i386" )
  ( resourceManagerContact = "fs2.das2.nikhef.nl" )
)
```

- GRAM comes with its own job / resource description language
- most middlewares invent their own languages
- requirements are interpreted in several places (resource broker, queue manager, ...)

# gLite Example

gLite: Job Submit

```
client.Delegate (delegID,  
                 "https://cream-ce-01:8443/.../CREAMDelegation",  
                 "/tmp/x509up_u202");  
client.Register ("https://cream-ce-01:8443/.../CREAM",  
                 "https://cream-ce-01:8443/.../CREAMDelegation",  
                 delegID,  
                 JobDescriptionBuffer,  
                 "/tmp/x509up_u202",  
                 uploadURL_and_jobID,  
                 0, false);  
client.Start    ("https://cream-ce-01:8443/.../CREAM",  
                 uploadURL_and_jobID[1]);
```

- moves security details to API level
- in some sense, is a customized globus like environment
- shows its Globus foundations
- faithful to the web service paradigm (app-level WSDL)  
→ API reflects gLite service interface
- JDL vs. RSL



# CoG Example

CoG: Job Submit

```
String gramContact = "pitcairn.mcs.anl.gov:6722:...";
String rsl          = "&(executable=...)(...)(...)";

GramJob job = null;
try {
    job = new GramJob (rsl);
    Gram.request (gramContact, job);
}
catch (GramException e) {
    ...
}
```

- covers same scope as Globus API
- hides complexity and API evolution
- separates functional and non functional API parts
- new versions provide additional functionality (workflow, GUI, ...) and cover non-globus middleware

# GAT Example

GAT: Job Submit

```
ResourceBroker      rb ();  
SoftwareDescription sd (("location",  "/bin/ls")  
                        ("arguments", "-l"));  
  
HardwareDescription hd (("memory.size", 1024.f)  
                        ("disk.size",    10.f)  
                        ("machine.type", "i686"));  
  
Job job = broker.SubmitJob (JobDescription (sd, hd));
```

- tries to abstract Grid Middleware functionality
- tries to hide middleware details
- implementable on multiple middleware systems
- usability limited by scope of its use cases (historical reasons)